
Pyro Documentation

Uber AI Labs

Jul 13, 2019

1	Installation	1
1.1	Install from Source	1
2	Getting Started	3
3	Primitives	5
4	Inference	9
4.1	SVI	9
4.2	ELBO	10
4.3	Importance	16
4.4	Discrete Inference	17
4.5	Inference Utilities	18
4.6	MCMC	20
5	Distributions	27
5.1	PyTorch Distributions	27
5.2	Pyro Distributions	31
5.3	Transformed Distributions	47
6	Parameters	61
6.1	ParamStore	61
7	Neural Network	65
7.1	AutoRegressiveNN	65
8	Optimization	69
8.1	Pyro Optimizers	69
8.2	PyTorch Optimizers	71
8.3	Higher-Order Optimizers	72
9	Poutine (Effect handlers)	75
9.1	Handlers	75
9.2	Trace	82
9.3	Messengers	84
9.4	Runtime	88
9.5	Utilities	89

10 Miscellaneous Ops	91
10.1 Utilities for HMC	91
10.2 Newton Optimizers	93
10.3 Tensor Indexing	95
10.4 Tensor Contraction	96
10.5 Statistical Utilities	98
11 Generic Interface	101
12 Automatic Guide Generation	103
12.1 AutoGuide	103
12.2 AutoGuideList	104
12.3 AutoCallable	104
12.4 AutoDelta	105
12.5 AutoContinuous	106
12.6 AutoMultivariateNormal	107
12.7 AutoDiagonalNormal	107
12.8 AutoLowRankMultivariateNormal	108
12.9 AutoIAFNormal	108
12.10 AutoLaplaceApproximation	109
12.11 AutoDiscreteParallel	109
12.12 Initialization	110
13 Automatic Name Generation	111
13.1 Named Data Structures	113
13.2 Scoping	115
14 Bayesian Neural Networks	119
14.1 HiddenLayer	119
15 Easy Custom Guides	121
15.1 EasyGuide	121
15.2 easy_guide	122
15.3 Group	122
16 Generalised Linear Mixed Models	125
17 Gaussian Processes	127
17.1 Models	127
17.2 Kernels	137
17.3 Likelihoods	144
17.4 Parameterized	147
17.5 Util	148
18 Mini Pyro	151
19 Optimal Experiment Design	153
19.1 Expected Information Gain	153
20 Tracking	161
20.1 Data Association	161
20.2 Distributions	164
20.3 Dynamic Models	164
20.4 Extended Kalman Filter	168
20.5 Hashing	170
20.6 Measurements	172

21 Indices and tables	175
Python Module Index	177
Index	179

1.1 Install from Source

Pyro supports Python 2.7.* and Python 3. To setup, install [PyTorch](#) then run:

```
pip install pyro-ppl
```

or install from source:

```
git clone https://github.com/uber/pyro.git
cd pyro
python setup.py install
```


CHAPTER 2

Getting Started

- [Install Pyro.](#)
- Learn the basic concepts of Pyro: [models](#) and [inference](#).
- Dive in to other [tutorials](#) and [examples](#).

sample (*name*, *fn*, **args*, ***kwargs*)

Calls the stochastic function *fn* with additional side-effects depending on *name* and the enclosing context (e.g. an inference algorithm). See [Intro I](#) and [Intro II](#) for a discussion.

Parameters

- **name** – name of sample
- **fn** – distribution class or function
- **obs** – observed datum (optional; should only be used in context of inference) optionally specified in *kwargs*
- **infer** (*dict*) – Optional dictionary of inference parameters specified in *kwargs*. See inference documentation for details.

Returns sample

param (*name*, **args*, ***kwargs*)

Saves the variable as a parameter in the param store. To interact with the param store or write to disk, see [Parameters](#).

Parameters

- **name** (*str*) – name of parameter
- **init_tensor** (*torch.Tensor* or *callable*) – initial tensor or lazy callable that returns a tensor. For large tensors, it may be cheaper to write e.g. `lambda: torch.randn(100000)`, which will only be evaluated on the initial statement.
- **constraint** (*torch.distributions.constraints.Constraint*) – torch constraint, defaults to `constraints.real`.
- **event_dim** (*int*) – (optional) number of rightmost dimensions unrelated to batching. Dimension to the left of this will be considered batch dimensions; if the param statement is inside a subsampled plate, then corresponding batch dimensions of the parameter will be correspondingly subsampled. If unspecified, all dimensions will be considered event dims and no subsampling will be performed.

Returns parameter

Return type `torch.Tensor`

module (*name*, *nn_module*, *update_module_params=False*)

Takes a `torch.nn.Module` and registers its parameters with the `ParamStore`. In conjunction with the `ParamStore` `save()` and `load()` functionality, this allows the user to save and load modules.

Parameters

- **name** (*str*) – name of module
- **nn_module** (*torch.nn.Module*) – the module to be registered with Pyro
- **update_module_params** – determines whether Parameters in the PyTorch module get overridden with the values found in the `ParamStore` (if any). Defaults to *False*

Returns `torch.nn.Module`

random_module (*name*, *nn_module*, *prior*, **args*, ***kwargs*)

Places a prior over the parameters of the module *nn_module*. Returns a distribution (callable) over *nn.Modules*, which upon calling returns a sampled *nn.Module*.

See the [Bayesian Regression tutorial](#) for an example.

Parameters

- **name** (*str*) – name of pyro module
- **nn_module** (*torch.nn.Module*) – the module to be registered with pyro
- **prior** – pyro distribution, stochastic function, or python dict with parameter names as keys and respective distributions/stochastic functions as values.

Returns a callable which returns a sampled module

class plate (*name*, *size=None*, *subsample_size=None*, *subsample=None*, *dim=None*, *use_cuda=None*, *device=None*)

Construct for conditionally independent sequences of variables.

`plate` can be used either sequentially as a generator or in parallel as a context manager (formerly `irange` and `iarange`, respectively).

Sequential `plate` is similar to `range()` in that it generates a sequence of values.

Vectorized `plate` is similar to `torch.arange()` in that it yields an array of indices by which other tensors can be indexed. `plate` differs from `torch.arange()` in that it also informs inference algorithms that the variables being indexed are conditionally independent. To do this, `plate` is provided as context manager rather than a function, and users must guarantee that all computation within an `plate` context is conditionally independent:

```
with plate("name", size) as ind:
    # ...do conditionally independent stuff with ind...
```

Additionally, `plate` can take advantage of the conditional independence assumptions by subsampling the indices and informing inference algorithms to scale various computed values. This is typically used to subsample minibatches of data:

```
with plate("data", len(data), subsample_size=100) as ind:
    batch = data[ind]
    assert len(batch) == 100
```

By default `subsample_size=False` and this simply yields a `torch.arange(0, size)`. If $0 < \text{subsample_size} \leq \text{size}$ this yields a single random batch of indices of size `subsample_size` and scales all log likelihood terms by `size/batch_size`, within this context.

Warning: This is only correct if all computation is conditionally independent within the context.

Parameters

- **name** (*str*) – A unique name to help inference algorithms match *plate* sites between models and guides.
- **size** (*int*) – Optional size of the collection being subsampled (like *stop* in builtin *range*).
- **subsample_size** (*int*) – Size of minibatches used in subsampling. Defaults to *size*.
- **subsample** (Anything supporting *len()*) – Optional custom subsample for user-defined subsampling schemes. If specified, then *subsample_size* will be set to *len(subsample)*.
- **dim** (*int*) – An optional dimension to use for this independence index. If specified, *dim* should be negative, i.e. should index from the right. If not specified, *dim* is set to the rightmost *dim* that is left of all enclosing *plate* contexts.
- **use_cuda** (*bool*) – DEPRECATED, use the *device* arg instead. Optional *bool* specifying whether to use cuda tensors for *subsample* and *log_prob*. Defaults to `torch.Tensor.is_cuda`.
- **device** (*str*) – Optional keyword specifying which device to place the results of *subsample* and *log_prob* on. By default, results are placed on the same device as the default tensor.

Returns A reusable context manager yielding a single 1-dimensional `torch.Tensor` of indices.

Examples:

```
>>> # This version declares sequential independence and subsamples data:
>>> for i in plate('data', 100, subsample_size=10):
...     if z[i]: # Control flow in this example prevents vectorization.
...         obs = sample('obs_{}'.format(i), dist.Normal(loc, scale),
... obs=data[i])
```

```
>>> # This version declares vectorized independence:
>>> with plate('data'):
...     obs = sample('obs', dist.Normal(loc, scale), obs=data)
```

```
>>> # This version subsamples data in vectorized way:
>>> with plate('data', 100, subsample_size=10) as ind:
...     obs = sample('obs', dist.Normal(loc, scale), obs=data[ind])
```

```
>>> # This wraps a user-defined subsampling method for use in pyro:
>>> ind = torch.randint(0, 100, (10,)).long() # custom subsample
>>> with plate('data', 100, subsample=ind):
...     obs = sample('obs', dist.Normal(loc, scale), obs=data[ind])
```

```
>>> # This reuses two different independence contexts.
>>> x_axis = plate('outer', 320, dim=-1)
>>> y_axis = plate('inner', 200, dim=-2)
```

(continues on next page)

(continued from previous page)

```

>>> with x_axis:
...     x_noise = sample("x_noise", dist.Normal(loc, scale))
...     assert x_noise.shape == (320,)
>>> with y_axis:
...     y_noise = sample("y_noise", dist.Normal(loc, scale))
...     assert y_noise.shape == (200, 1)
>>> with x_axis, y_axis:
...     xy_noise = sample("xy_noise", dist.Normal(loc, scale))
...     assert xy_noise.shape == (200, 320)

```

See [SVI Part II](#) for an extended discussion.

get_param_store()

Returns the ParamStore

clear_param_store()

Clears the ParamStore. This is especially useful if you're working in a REPL.

validation_enabled(*args, **kws)

Context manager that is useful when temporarily enabling/disabling validation checks.

Parameters **is_validate** (*bool*) – (optional; defaults to True) temporary validation check override.

enable_validation(is_validate=True)

Enable or disable validation checks in Pyro. Validation checks provide useful warnings and errors, e.g. NaN checks, validating distribution arguments and support values, etc. which is useful for debugging. Since some of these checks may be expensive, we recommend turning this off for mature models.

Parameters **is_validate** (*bool*) – (optional; defaults to True) whether to enable validation checks.

trace(fn=None, ignore_warnings=False, jit_options=None)

Lazy replacement for `torch.jit.trace()` that works with Pyro functions that call `pyro.param()`.

The actual compilation artifact is stored in the `compiled` attribute of the output. Call diagnostic methods on this attribute.

Example:

```

def model(x):
    scale = pyro.param("scale", torch.tensor(0.5), constraint=constraints.
    ↳positive)
    return pyro.sample("y", dist.Normal(x, scale))

@pyro.ops.jit.trace
def model_log_prob_fn(x, y):
    cond_model = pyro.condition(model, data={"y": y})
    tr = pyro.poutine.trace(cond_model).get_trace(x)
    return tr.log_prob_sum()

```

Parameters

- **fn** (*callable*) – The function to be traced.
- **ignore_warnins** (*bool*) – Whether to ignore jit warnings.
- **jit_options** (*dict*) – Optional dict of options to pass to `torch.jit.trace()`, e.g. `{"optimize": False}`.

In the context of probabilistic modeling, learning is usually called inference. In the particular case of Bayesian inference, this often involves computing (approximate) posterior distributions. In the case of parameterized models, this usually involves some sort of optimization. Pyro supports multiple inference algorithms, with support for stochastic variational inference (SVI) being the most extensive. Look here for more inference algorithms in future versions of Pyro.

See [Intro II](#) for a discussion of inference in Pyro.

4.1 SVI

class SVI (*model, guide, optim, loss, loss_and_grads=None, num_samples=10, num_steps=0, **kwargs*)
 Bases: `pyro.infer.abstract_infer.TracePosterior`

Parameters

- **model** – the model (callable containing Pyro primitives)
- **guide** – the guide (callable containing Pyro primitives)
- **optim** (`pyro.optim.PyroOptim`) – a wrapper for a PyTorch optimizer
- **loss** (`pyro.infer.elbo.ELBO`) – an instance of a subclass of `ELBO`. Pyro provides three built-in losses: `Trace_ELBO`, `TraceGraph_ELBO`, and `TraceEnum_ELBO`. See the `ELBO` docs to learn how to implement a custom loss.
- **num_samples** – the number of samples for Monte Carlo posterior approximation
- **num_steps** – the number of optimization steps to take in `run()`

A unified interface for stochastic variational inference in Pyro. The most commonly used loss is `loss=Trace_ELBO()`. See the tutorial [SVI Part I](#) for a discussion.

evaluate_loss (**args, **kwargs*)

Returns estimate of the loss

Return type `float`

Evaluate the loss function. Any args or kwargs are passed to the model and guide.

run (**args*, ***kwargs*)

step (**args*, ***kwargs*)

Returns estimate of the loss

Return type `float`

Take a gradient step on the loss function (and any auxiliary loss functions generated under the hood by *loss_and_grads*). Any args or kwargs are passed to the model and guide

4.2 ELBO

```
class ELBO(num_particles=1, max_plate_nesting=inf, max_iarange_nesting=None, vector-
           ize_particles=False, strict_enumeration_warning=True, ignore_jit_warnings=False,
           jit_options=None, retain_graph=None, tail_adaptive_beta=-1.0)
```

Bases: `object`

ELBO is the top-level interface for stochastic variational inference via optimization of the evidence lower bound.

Most users will not interact with this base class *ELBO* directly; instead they will create instances of derived classes: *Trace_ELBO*, *TraceGraph_ELBO*, or *TraceEnum_ELBO*.

Parameters

- **num_particles** – The number of particles/samples used to form the ELBO (gradient) estimators.
- **max_plate_nesting** (*int*) – Optional bound on max number of nested *pyro.plate()* contexts. This is only required when enumerating over sample sites in parallel, e.g. if a site sets `infer={"enumerate": "parallel"}`. If omitted, ELBO may guess a valid value by running the (model,guide) pair once, however this guess may be incorrect if model or guide structure is dynamic.
- **vectorize_particles** (*bool*) – Whether to vectorize the ELBO computation over *num_particles*. Defaults to False. This requires static structure in model and guide.
- **strict_enumeration_warning** (*bool*) – Whether to warn about possible misuse of enumeration, i.e. that *pyro.infer.traceenum_elbo.TraceEnum_ELBO* is used iff there are enumerated sample sites.
- **ignore_jit_warnings** (*bool*) – Flag to ignore warnings from the JIT tracer. When this is True, all `torch.jit.TracerWarning` will be ignored. Defaults to False.
- **jit_options** (*bool*) – Optional dict of options to pass to `torch.jit.trace()`, e.g. `{"optimize": False}`.
- **retain_graph** (*bool*) – Whether to retain autograd graph during an SVI step. Defaults to None (False).
- **tail_adaptive_beta** (*float*) – Exponent beta with $-1.0 \leq \text{beta} < 0.0$ for use with *TraceTailAdaptive_ELBO*.

References

- [1] *Automated Variational Inference in Probabilistic Programming* David Wingate, Theo Weber
- [2] *Black Box Variational Inference*, Rajesh Ranganath, Sean Gerrish, David M. Blei


```
class Trace_ELBO (num_particles=1, max_plate_nesting=inf, max_iarange_nesting=None, vectorize_particles=False, strict_enumeration_warning=True, ignore_jit_warnings=False, jit_options=None, retain_graph=None, tail_adaptive_beta=-1.0)
```

Bases: `pyro.infer.elbo.ELBO`

A trace implementation of ELBO-based SVI. The estimator is constructed along the lines of references [1] and [2]. There are no restrictions on the dependency structure of the model or the guide. The gradient estimator includes partial Rao-Blackwellization for reducing the variance of the estimator when non-reparameterizable random variables are present. The Rao-Blackwellization is partial in that it only uses conditional independence information that is marked by `plate` contexts. For more fine-grained Rao-Blackwellization, see `TraceGraph_ELBO`.

References

[1] **Automated Variational Inference in Probabilistic Programming**, David Wingate, Theo Weber

[2] **Black Box Variational Inference**, Rajesh Ranganath, Sean Gerrish, David M. Blei

```
loss (model, guide, *args, **kwargs)
```

Returns returns an estimate of the ELBO

Return type `float`

Evaluates the ELBO with an estimator that uses `num_particles` many samples/particles.

```
differentiable_loss (model, guide, *args, **kwargs)
```

Computes the surrogate loss that can be differentiated with autograd to produce gradient estimates for the model and guide parameters

```
loss_and_grads (model, guide, *args, **kwargs)
```

Returns returns an estimate of the ELBO

Return type `float`

Computes the ELBO as well as the surrogate ELBO that is used to form the gradient estimator. Performs backward on the latter. Num_particle many samples are used to form the estimators.

```
class JitTrace_ELBO (num_particles=1, max_plate_nesting=inf, max_iarange_nesting=None, vectorize_particles=False, strict_enumeration_warning=True, ignore_jit_warnings=False, jit_options=None, retain_graph=None, tail_adaptive_beta=-1.0)
```

Bases: `pyro.infer.trace_elbo.Trace_ELBO`

Like `Trace_ELBO` but uses `pyro.ops.jit.compile()` to compile `loss_and_grads()`.

This works only for a limited set of models:

- Models must have static structure.
- Models must not depend on any global data (except the param store).
- All model inputs that are tensors must be passed in via `*args`.
- All model inputs that are *not* tensors must be passed in via `**kwargs`, and compilation will be triggered once per unique `**kwargs`.

```
loss_and_surrogate_loss (model, guide, *args, **kwargs)
```

```
differentiable_loss (model, guide, *args, **kwargs)
```

```
loss_and_grads (model, guide, *args, **kwargs)
```

```
class TraceGraph_ELBO (num_particles=1, max_plate_nesting=inf, max_iarange_nesting=None,
                        vectorize_particles=False, strict_enumeration_warning=True, ignore_jit_warnings=False,
                        jit_options=None, retain_graph=None, tail_adaptive_beta=-1.0)
```

Bases: `pyro.infer.elbo.ELBO`

A `TraceGraph` implementation of ELBO-based SVI. The gradient estimator is constructed along the lines of reference [1] specialized to the case of the ELBO. It supports arbitrary dependency structure for the model and guide as well as baselines for non-reparameterizable random variables. Where possible, conditional dependency information as recorded in the `Trace` is used to reduce the variance of the gradient estimator. In particular two kinds of conditional dependency information are used to reduce variance:

- the sequential order of samples (z is sampled after y => y does not depend on z)
- `plate` generators

References

[1] *Gradient Estimation Using Stochastic Computation Graphs*, John Schulman, Nicolas Heess, Theophane Weber, Pieter Abbeel

[2] *Neural Variational Inference and Learning in Belief Networks* Andriy Mnih, Karol Gregor

```
loss (model, guide, *args, **kwargs)
```

Returns returns an estimate of the ELBO

Return type `float`

Evaluates the ELBO with an estimator that uses `num_particles` many samples/particles.

```
loss_and_grads (model, guide, *args, **kwargs)
```

Returns returns an estimate of the ELBO

Return type `float`

Computes the ELBO as well as the surrogate ELBO that is used to form the gradient estimator. Performs backward on the latter. `Num_particle` many samples are used to form the estimators. If baselines are present, a baseline loss is also constructed and differentiated.

```
class JitTraceGraph_ELBO (num_particles=1, max_plate_nesting=inf, max_iarange_nesting=None,
                           vectorize_particles=False, strict_enumeration_warning=True, ignore_jit_warnings=False,
                           jit_options=None, retain_graph=None, tail_adaptive_beta=-1.0)
```

Bases: `pyro.infer.tracegraph_elbo.TraceGraph_ELBO`

Like `TraceGraph_ELBO` but uses `torch.jit.trace()` to compile `loss_and_grads()`.

This works only for a limited set of models:

- Models must have static structure.
- Models must not depend on any global data (except the param store).
- All model inputs that are tensors must be passed in via `*args`.
- All model inputs that are *not* tensors must be passed in via `**kwargs`, and compilation will be triggered once per unique `**kwargs`.

```
loss_and_grads (model, guide, *args, **kwargs)
```

```
class BackwardSampleMessenger (enum_trace, guide_trace)
```

Bases: `pyro.poutine.messenger.Messenger`

Implements forward filtering / backward sampling for sampling from the joint posterior distribution

```
class TraceEnum_ELBO (num_particles=1,    max_plate_nesting=inf,    max_iarange_nesting=None,
                      vectorize_particles=False,    strict_enumeration_warning=True,    ignore_jit_warnings=False,
                      jit_options=None,    retain_graph=None,
                      tail_adaptive_beta=-1.0)
```

Bases: `pyro.infer.elbo.ELBO`

A trace implementation of ELBO-based SVI that supports - exhaustive enumeration over discrete sample sites, and - local parallel sampling over any sample site.

To enumerate over a sample site in the guide, mark the site with either `infer={'enumerate': 'sequential'}` or `infer={'enumerate': 'parallel'}`. To configure all guide sites at once, use `config_enumerate()`. To enumerate over a sample site in the model, mark the site `infer={'enumerate': 'parallel'}` and ensure the site does not appear in the guide.

This assumes restricted dependency structure on the model and guide: variables outside of an `plate` can never depend on variables inside that `plate`.

loss (model, guide, *args, **kwargs)

Returns an estimate of the ELBO

Return type `float`

Estimates the ELBO using `num_particles` many samples (particles).

differentiable_loss (model, guide, *args, **kwargs)

Returns a differentiable estimate of the ELBO

Return type `torch.Tensor`

Raises `ValueError` – if the ELBO is not differentiable (e.g. is identically zero)

Estimates a differentiable ELBO using `num_particles` many samples (particles). The result should be infinitely differentiable (as long as underlying derivatives have been implemented).

loss_and_grads (model, guide, *args, **kwargs)

Returns an estimate of the ELBO

Return type `float`

Estimates the ELBO using `num_particles` many samples (particles). Performs backward on the ELBO of each particle.

compute_marginals (model, guide, *args, **kwargs)

Computes marginal distributions at each model-enumerated sample site.

Returns a dict mapping site name to marginal `Distribution` object

Return type `OrderedDict`

sample_posterior (model, guide, *args, **kwargs)

Sample from the joint posterior distribution of all model-enumerated sites given all observations

```
class JitTraceEnum_ELBO (num_particles=1,    max_plate_nesting=inf,    max_iarange_nesting=None,
                        vectorize_particles=False,    strict_enumeration_warning=True,    ignore_jit_warnings=False,
                        jit_options=None,    retain_graph=None,
                        tail_adaptive_beta=-1.0)
```

Bases: `pyro.infer.traceenum_elbo.TraceEnum_ELBO`

Like `TraceEnum_ELBO` but uses `pyro.ops.jit.compile()` to compile `loss_and_grads()`.

This works only for a limited set of models:

- Models must have static structure.

- Models must not depend on any global data (except the param store).
- All model inputs that are tensors must be passed in via `*args`.
- All model inputs that are *not* tensors must be passed in via `**kwargs`, and compilation will be triggered once per unique `**kwargs`.

`differentiable_loss` (*model*, *guide*, **args*, ***kwargs*)

`loss_and_grads` (*model*, *guide*, **args*, ***kwargs*)

```
class TraceMeanField_ELBO (num_particles=1, max_plate_nesting=inf, max_iarange_nesting=None,
                           vectorize_particles=False, strict_enumeration_warning=True, ignore_jit_warnings=False,
                           jit_options=None, retain_graph=None, tail_adaptive_beta=-1.0)
```

Bases: `pyro.infer.trace_elbo.Trace_ELBO`

A trace implementation of ELBO-based SVI. This is currently the only ELBO estimator in Pyro that uses analytic KL divergences when those are available.

In contrast to, e.g., `TraceGraph_ELBO` and `Trace_ELBO` this estimator places restrictions on the dependency structure of the model and guide. In particular it assumes that the guide has a mean-field structure, i.e. that it factorizes across the different latent variables present in the guide. It also assumes that all of the latent variables in the guide are reparameterized. This latter condition is satisfied for, e.g., the Normal distribution but is not satisfied for, e.g., the Categorical distribution.

Warning: This estimator may give incorrect results if the mean-field condition is not satisfied.

Note for advanced users:

The mean field condition is a sufficient but not necessary condition for this estimator to be correct. The precise condition is that for every latent variable z in the guide, its parents in the model must not include any latent variables that are descendants of z in the guide. Here ‘parents in the model’ and ‘descendants in the guide’ is with respect to the corresponding (statistical) dependency structure. For example, this condition is always satisfied if the model and guide have identical dependency structures.

`loss` (*model*, *guide*, **args*, ***kwargs*)

Returns returns an estimate of the ELBO

Return type `float`

Evaluates the ELBO with an estimator that uses `num_particles` many samples/particles.

```
class JitTraceMeanField_ELBO (num_particles=1, max_plate_nesting=inf,
                              max_iarange_nesting=None, vectorize_particles=False,
                              strict_enumeration_warning=True, ignore_jit_warnings=False,
                              jit_options=None, retain_graph=None, tail_adaptive_beta=-1.0)
```

Bases: `pyro.infer.trace_mean_field_elbo.TraceMeanField_ELBO`

Like `TraceMeanField_ELBO` but uses `pyro.ops.jit.trace()` to compile `loss_and_grads()`.

This works only for a limited set of models:

- Models must have static structure.
- Models must not depend on any global data (except the param store).
- All model inputs that are tensors must be passed in via `*args`.
- All model inputs that are *not* tensors must be passed in via `**kwargs`, and compilation will be triggered once per unique `**kwargs`.

`differentiable_loss (model, guide, *args, **kwargs)`

`loss_and_grads (model, guide, *args, **kwargs)`

```
class TraceTailAdaptive_ELBO (num_particles=1, max_plate_nesting=inf,
                              max_iarange_nesting=None, vectorize_particles=False,
                              strict_enumeration_warning=True, ignore_jit_warnings=False,
                              jit_options=None, retain_graph=None, tail_adaptive_beta=-1.0)
```

Bases: `pyro.infer.trace_elbo.Trace_ELBO`

Interface for Stochastic Variational Inference with an adaptive f-divergence as described in ref. [1]. Users should specify `num_particles > 1` and `vectorize_particles==True`. The argument `tail_adaptive_beta` can be specified to modify how the adaptive f-divergence is constructed. See reference for details.

Note that this interface does not support computing the variational objective itself; rather it only supports computing gradients of the variational objective. Consequently, one might want to use another SVI interface (e.g. `RenyiELBO`) in order to monitor convergence.

Note that this interface only supports models in which all the latent variables are fully reparameterized. It also does not support data subsampling.

References [1] “Variational Inference with Tail-adaptive f-Divergence”, Dilin Wang, Hao Liu, Qiang Liu, NeurIPS 2018 <https://papers.nips.cc/paper/7816-variational-inference-with-tail-adaptive-f-divergence>

`loss (model, guide, *args, **kwargs)`

It is not necessary to estimate the tail-adaptive f-divergence itself in order to compute the corresponding gradients. Consequently the loss method is left unimplemented.

```
class RenyiELBO (alpha=0, num_particles=2, max_plate_nesting=inf, max_iarange_nesting=None,
                 vectorize_particles=False, strict_enumeration_warning=True)
```

Bases: `pyro.infer.elbo.ELBO`

An implementation of Renyi’s α -divergence variational inference following reference [1].

In order for the objective to be a strict lower bound, we require $\alpha \geq 0$. Note, however, that according to reference [1], depending on the dataset $\alpha < 0$ might give better results. In the special case $\alpha = 0$, the objective function is that of the important weighted autoencoder derived in reference [2].

Note: Setting $\alpha < 1$ gives a better bound than the usual ELBO. For $\alpha = 1$, it is better to use `Trace_ELBO` class because it helps reduce variances of gradient estimations.

Warning: Mini-batch training is not supported yet.

Parameters

- **alpha** (*float*) – The order of α -divergence. Here $\alpha \neq 1$. Default is 0.
- **num_particles** – The number of particles/samples used to form the objective (gradient) estimator. Default is 2.
- **max_plate_nesting** (*int*) – Bound on max number of nested `pyro.plate()` contexts. Default is infinity.
- **strict_enumeration_warning** (*bool*) – Whether to warn about possible misuse of enumeration, i.e. that `TraceEnum_ELBO` is used iff there are enumerated sample sites.

References:

[1] *Renyi Divergence Variational Inference*, Yingzhen Li, Richard E. Turner

[2] *Importance Weighted Autoencoders*, Yuri Burda, Roger Grosse, Ruslan Salakhutdinov

loss (*model*, *guide*, **args*, ***kwargs*)

Returns returns an estimate of the ELBO

Return type `float`

Evaluates the ELBO with an estimator that uses `num_particles` many samples/particles.

loss_and_grads (*model*, *guide*, **args*, ***kwargs*)

Returns returns an estimate of the ELBO

Return type `float`

Computes the ELBO as well as the surrogate ELBO that is used to form the gradient estimator. Performs backward on the latter. `Num_particle` many samples are used to form the estimators.

4.3 Importance

class Importance (*model*, *guide*=None, *num_samples*=None)

Bases: `pyro.infer.abstract_infer.TracePosterior`

Parameters

- **model** – probabilistic model defined as a function
- **guide** – guide used for sampling defined as a function
- **num_samples** – number of samples to draw from the guide (default 10)

This method performs posterior inference by importance sampling using the guide as the proposal distribution. If no guide is provided, it defaults to proposing from the model’s prior.

get_ESS ()

Compute (Importance Sampling) Effective Sample Size (ESS).

get_log_normalizer ()

Estimator of the normalizing constant of the target distribution. (mean of the unnormalized weights)

get_normalized_weights (*log_scale*=False)

Compute the normalized importance weights.

psis_diagnostic (**args*, ***kwargs*)

Computes the Pareto tail index k for a model/guide pair using the technique described in [1], which builds on previous work in [2]. If $0 < k < 0.5$ the guide is a good approximation to the model posterior, in the sense described in [1]. If $0.5 \leq k \leq 0.7$, the guide provides a suboptimal approximation to the posterior, but may still be useful in practice. If $k > 0.7$ the guide program provides a poor approximation to the full posterior, and caution should be used when using the guide. Note, however, that a guide may be a poor fit to the full posterior while still yielding reasonable model predictions. If $k < 0.0$ the importance weights corresponding to the model and guide appear to be bounded from above; this would be a bizarre outcome for a guide trained via ELBO maximization. Please see [1] for a more complete discussion of how the tail index k should be interpreted.

Please be advised that a large number of samples may be required for an accurate estimate of k .

Note that we assume that the model and guide are both vectorized and have static structure. As is canonical in Pyro, the args and kwargs are passed to the model and guide.

References [1] ‘Yes, but Did It Work?: Evaluating Variational Inference.’ Yuling Yao, Aki Vehtari, Daniel Simpson, Andrew Gelman [2] ‘Pareto Smoothed Importance Sampling.’ Aki Vehtari, Andrew Gelman, Jonah Gabry

Parameters

- **model** (*callable*) – the model program.
- **guide** (*callable*) – the guide program.
- **num_particles** (*int*) – the total number of times we run the model and guide in order to compute the diagnostic. defaults to 1000.
- **max_simultaneous_particles** – the maximum number of simultaneous samples drawn from the model and guide. defaults to *num_particles*. *num_particles* must be divisible by *max_simultaneous_particles*. compute the diagnostic. defaults to 1000.
- **max_plate_nesting** (*int*) – optional bound on max number of nested `pyro.plate()` contexts in the model/guide. defaults to 7.

Returns float the PSIS diagnostic k

vectorized_importance_weights (*model*, *guide*, **args*, ***kwargs*)

Parameters

- **model** – probabilistic model defined as a function
- **guide** – guide used for sampling defined as a function
- **num_samples** – number of samples to draw from the guide (default 1)
- **max_plate_nesting** (*int*) – Bound on max number of nested `pyro.plate()` contexts.
- **normalized** (*bool*) – set to True to return self-normalized importance weights

Returns returns a $(\text{num_samples},)$ -shaped tensor of importance weights and the model and guide traces that produced them

Vectorized computation of importance weights for models with static structure:

```
log_weights, model_trace, guide_trace = \
    vectorized_importance_weights(model, guide, *args,
                                  num_samples=1000,
                                  max_plate_nesting=4,
                                  normalized=False)
```

4.4 Discrete Inference

infer_discrete (*fn=None*, *first_available_dim=None*, *temperature=1*)

A routine that samples discrete sites marked with `site["infer"]["enumerate"] = "parallel"` from the posterior, conditioned on observations.

Example:

```
@infer_discrete(first_available_dim=-1, temperature=0)
@config_enumerate
def viterbi_decoder(data, hidden_dim=10):
    transition = 0.3 / hidden_dim + 0.7 * torch.eye(hidden_dim)
    means = torch.arange(float(hidden_dim))
    states = [0]
    for t in pyro.markov(range(len(data))):
        states.append(pyro.sample("states_{}".format(t),
```

(continues on next page)

(continued from previous page)

```

                                dist.Categorical(transition[states[-1]]))
    pyro.sample("obs_{}".format(t),
                dist.Normal(means[states[-1]], 1.),
                obs=data[t])
    return states # returns maximum likelihood states

```

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **first_available_dim** (*int*) – The first tensor dimension (counting from the right) that is available for parallel enumeration. This dimension and all dimensions left may be used internally by Pyro. This should be a negative integer.
- **temperature** (*int*) – Either 1 (sample via forward-filter backward-sample) or 0 (optimize via Viterbi-like MAP inference). Defaults to 1 (sample).

```

class TraceEnumSample_ELBO(num_particles=1,                                max_plate_nesting=inf,
                           max_iarange_nesting=None,                    vectorize_particles=False,
                           strict_enumeration_warning=True,              ignore_jit_warnings=False,
                           jit_options=None, retain_graph=None, tail_adaptive_beta=-1.0)

```

Bases: `pyro.infer.traceenum_elbo.TraceEnum_ELBO`

This extends `TraceEnum_ELBO` to make it cheaper to sample from discrete latent states during SVI.

The following are equivalent but the first is cheaper, sharing work between the computations of `loss` and `z`:

```

# Version 1.
elbo = TraceEnumSample_ELBO(max_plate_nesting=1)
loss = elbo.loss(*args, **kwargs)
z = elbo.sample_saved()

# Version 2.
elbo = TraceEnum_ELBO(max_plate_nesting=1)
loss = elbo.loss(*args, **kwargs)
guide_trace = poutine.trace(guide).get_trace(*args, **kwargs)
z = infer_discrete(poutine.replay(model, guide_trace),
                  first_available_dim=-2)(*args, **kwargs)

```

sample_saved()

Generate latent samples while reusing work from `SVI.step()`.

4.5 Inference Utilities

```

class EmpiricalMarginal(trace_posterior, sites=None, validate_args=None)

```

Bases: `pyro.distributions.empirical.Empirical`

Marginal distribution over a single site (or multiple, provided they have the same shape) from the `TracePosterior`'s model.

Note: If multiple sites are specified, they must have the same tensor shape. Samples from each site will be stacked and stored within a single tensor. See [Empirical](#). To hold the marginal distribution of sites having different shapes, use [Marginals](#) instead.

Parameters

- **trace_posterior** (`TracePosterior`) – a `TracePosterior` instance representing a Monte Carlo posterior.
- **sites** (`list`) – optional list of sites for which we need to generate the marginal distribution.

class `Marginals` (`trace_posterior`, `sites=None`, `validate_args=None`)

Bases: `object`

Holds the marginal distribution over one or more sites from the `TracePosterior`’s model. This is a convenience container class, which can be extended by `TracePosterior` subclasses. e.g. for implementing diagnostics.

Parameters

- **trace_posterior** (`TracePosterior`) – a `TracePosterior` instance representing a Monte Carlo posterior.
- **sites** (`list`) – optional list of sites for which we need to generate the marginal distribution.

empirical

A dictionary of sites’ names and their corresponding `EmpiricalMarginal` distribution.

Type `OrderedDict`

support (`flatten=False`)

Gets support of this marginal distribution.

Parameters **flatten** (`bool`) – A flag to decide if we want to flatten `batch_shape` when the marginal distribution is collected from the posterior with `num_chains > 1`. Defaults to `False`.

Returns a dict with keys are sites’ names and values are sites’ supports.

Return type `OrderedDict`

class `TracePosterior` (`num_chains=1`)

Bases: `object`

Abstract `TracePosterior` object from which posterior inference algorithms inherit. When run, collects a bag of execution traces from the approximate posterior. This is designed to be used by other utility classes like `EmpiricalMarginal`, that need access to the collected execution traces.

information_criterion (`pointwise=False`)

Computes information criterion of the model. Currently, returns only “Widely Applicable/Watanabe-Akaike Information Criterion” (WAIC) and the corresponding effective number of parameters.

Reference:

[1] *Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC*, Aki Vehtari, Andrew Gelman, and Jonah Gabry

Parameters **pointwise** (`bool`) – a flag to decide if we want to get a vectorized WAIC or not. When `pointwise=False`, returns the sum.

Returns a dictionary containing values of WAIC and its effective number of parameters.

Return type `OrderedDict`

marginal (`sites=None`)

Generates the marginal distribution of this posterior.

Parameters `sites` (*list*) – optional list of sites for which we need to generate the marginal distribution.

Returns A *Marginals* class instance.

Return type *Marginals*

run (**args*, ***kwargs*)

Calls `self._traces` to populate execution traces from a stochastic Pyro model.

Parameters

- **args** – optional args taken by `self._traces`.
- **kwargs** – optional keywords args taken by `self._traces`.

class `TracePredictive` (*model*, *posterior*, *num_samples*, *keep_sites=None*)

Bases: `pyro.infer.abstract_infer.TracePosterior`

Generates and holds traces from the posterior predictive distribution, given model execution traces from the approximate posterior. This is achieved by constraining latent sites to randomly sampled parameter values from the model execution traces and running the model forward to generate traces with new response (“_RETURN”) sites. :param *model*: arbitrary Python callable containing Pyro primitives. :param *TracePosterior* *posterior*: trace posterior instance holding samples from the model’s approximate posterior. :param *int* *num_samples*: number of samples to generate. :param *keep_sites*: The sites which should be sampled from posterior distribution (default: all)

marginal (*sites=None*)

Gets marginal distribution for this predictive posterior distribution.

4.6 MCMC

4.6.1 MCMC

class `MCMC` (*kernel*, *num_samples*, *warmup_steps=None*, *initial_params=None*, *num_chains=1*, *hook_fn=None*, *mp_context=None*, *disable_progbar=False*, *disable_validation=True*, *transforms=None*)

Bases: `object`

Wrapper class for Markov Chain Monte Carlo algorithms. Specific MCMC algorithms are `TraceKernel` instances and need to be supplied as a `kernel` argument to the constructor.

Note: The case of `num_chains > 1` uses python multiprocessing to run parallel chains in multiple processes. This goes with the usual caveats around multiprocessing in python, e.g. the model used to initialize the `kernel` must be serializable via `pickle`, and the performance / constraints will be platform dependent (e.g. only the “spawn” context is available in Windows). This has also not been extensively tested on the Windows platform.

Parameters

- **kernel** – An instance of the `TraceKernel` class, which when given an execution trace returns another sample trace from the target (posterior) distribution.
- **num_samples** (*int*) – The number of samples that need to be generated, excluding the samples discarded during the warmup phase.
- **warmup_steps** (*int*) – Number of warmup iterations. The samples generated during the warmup phase are discarded. If not provided, default is half of *num_samples*.

- **num_chains** (*int*) – Number of MCMC chains to run in parallel. Depending on whether *num_chains* is 1 or more than 1, this class internally dispatches to either *_UnarySampler* or *_MultiSampler*.
- **initial_params** (*dict*) – dict containing initial tensors in unconstrained space to initiate the markov chain. The leading dimension's size must match that of *num_chains*. If not specified, parameter values will be sampled from the prior.
- **hook_fn** – Python callable that takes in (*kernel*, *samples*, *stage*, *i*) as arguments. *stage* is either *sample* or *warmup* and *i* refers to the *i*'th sample for the given stage. This can be used to implement additional logging, or more generally, run arbitrary code per generated sample.
- **mp_context** (*str*) – Multiprocessing context to use when *num_chains* > 1. Only applicable for Python 3.5 and above. Use *mp_context*="spawn" for CUDA.
- **disable_progbar** (*bool*) – Disable progress bar and diagnostics update.
- **disable_validation** (*bool*) – Disables distribution validation check. This is disabled by default, since divergent transitions will lead to exceptions. Switch to *True* for debugging purposes.
- **transforms** (*dict*) – dictionary that specifies a transform for a sample site with constrained support to unconstrained space.

diagnostics ()

Gets some diagnostics statistics such as effective sample size, split Gelman-Rubin, or divergent transitions from the sampler.

get_samples (*num_samples=None*, *group_by_chain=False*)

Get samples from the MCMC run, potentially resampling with replacement.

Parameters

- **num_samples** (*int*) – Number of samples to return. If *None*, all the samples from an MCMC chain are returned in their original ordering.
- **group_by_chain** (*bool*) – Whether to preserve the chain dimension. If *True*, all samples will have *num_chains* as the size of their leading dimension.

Returns dictionary of samples keyed by site name.

run (*args, **kwargs)

summary (*prob=0.9*)

Prints a summary table displaying diagnostics of samples obtained from posterior. The diagnostics displayed are mean, standard deviation, median, the 90% Credibility Interval, *effective_sample_size()*, *split_gelman_rubin()*.

Parameters **prob** (*float*) – the probability mass of samples within the credibility interval.

4.6.2 HMC

```
class HMC(model=None, potential_fn=None, step_size=1, trajectory_length=None, num_steps=None,
          adapt_step_size=True, adapt_mass_matrix=True, full_mass=False, transforms=None,
          max_plate_nesting=None, jit_compile=False, jit_options=None, ignore_jit_warnings=False,
          target_accept_prob=0.8)
```

Bases: `pyro.infer.mcmc.mcmc_kernel.MCMCKernel`

Simple Hamiltonian Monte Carlo kernel, where *step_size* and *num_steps* need to be explicitly specified by the user.

References

[1] *MCMC Using Hamiltonian Dynamics*, Radford M. Neal

Parameters

- **model** – Python callable containing Pyro primitives.
- **potential_fn** – Python callable calculating potential energy with input is a dict of real support parameters.
- **step_size** (*float*) – Determines the size of a single step taken by the verlet integrator while computing the trajectory using Hamiltonian dynamics. If not specified, it will be set to 1.
- **trajectory_length** (*float*) – Length of a MCMC trajectory. If not specified, it will be set to `step_size × num_steps`. In case `num_steps` is not specified, it will be set to 2π .
- **num_steps** (*int*) – The number of discrete steps over which to simulate Hamiltonian dynamics. The state at the end of the trajectory is returned as the proposal. This value is always equal to `int(trajectory_length / step_size)`.
- **adapt_step_size** (*bool*) – A flag to decide if we want to adapt `step_size` during warm-up phase using Dual Averaging scheme.
- **adapt_mass_matrix** (*bool*) – A flag to decide if we want to adapt mass matrix during warm-up phase using Welford scheme.
- **full_mass** (*bool*) – A flag to decide if mass matrix is dense or diagonal.
- **transforms** (*dict*) – Optional dictionary that specifies a transform for a sample site with constrained support to unconstrained space. The transform should be invertible, and implement `log_abs_det_jacobian`. If not specified and the model has sites with constrained support, automatic transformations will be applied, as specified in `torch.distributions.constraint_registry`.
- **max_plate_nesting** (*int*) – Optional bound on max number of nested `pyro.plate()` contexts. This is required if model contains discrete sample sites that can be enumerated over in parallel.
- **jit_compile** (*bool*) – Optional parameter denoting whether to use the PyTorch JIT to trace the log density computation, and use this optimized executable trace in the integrator.
- **jit_options** (*dict*) – A dictionary contains optional arguments for `torch.jit.trace()` function.
- **ignore_jit_warnings** (*bool*) – Flag to ignore warnings from the JIT tracer when `jit_compile=True`. Default is False.
- **target_accept_prob** (*float*) – Increasing this value will lead to a smaller step size, hence the sampling will be slower and more robust. Default to 0.8.

Note: Internally, the mass matrix will be ordered according to the order of the names of latent variables, not the order of their appearance in the model.

Example:

```
>>> true_coefs = torch.tensor([1., 2., 3.])
>>> data = torch.randn(2000, 3)
>>> dim = 3
```

(continues on next page)

(continued from previous page)

```

>>> labels = dist.Bernoulli(logits=(true_coefs * data).sum(-1)).sample()
>>>
>>> def model(data):
...     coefs_mean = torch.zeros(dim)
...     coefs = pyro.sample('beta', dist.Normal(coefs_mean, torch.ones(3)))
...     y = pyro.sample('y', dist.Bernoulli(logits=(coefs * data).sum(-1)),
... ↪obs=labels)
...     return y
>>>
>>> hmc_kernel = HMC(model, step_size=0.0855, num_steps=4)
>>> mcmc = MCMC(hmc_kernel, num_samples=500, warmup_steps=100)
>>> mcmc.run(data)
>>> mcmc.get_samples()['beta'].mean(0) # doctest: +SKIP
tensor([ 0.9819,  1.9258,  2.9737])

```

`cleanup()``clear_cache()``diagnostics()``initial_params``inverse_mass_matrix``logging()``num_steps``sample(params)``setup(warmup_steps, *args, **kwargs)``step_size`

4.6.3 NUTS

```

class NUTS(model=None, potential_fn=None, step_size=1, adapt_step_size=True,
            adapt_mass_matrix=True, full_mass=False, use_multinomial_sampling=True, trans-
            forms=None, max_plate_nesting=None, jit_compile=False, jit_options=None, ig-
            nore_jit_warnings=False, target_accept_prob=0.8, max_tree_depth=10)
Bases: pyro.infer.mcmc.hmc.HMC

```

No-U-Turn Sampler kernel, which provides an efficient and convenient way to run Hamiltonian Monte Carlo. The number of steps taken by the integrator is dynamically adjusted on each call to `sample` to ensure an optimal length for the Hamiltonian trajectory [1]. As such, the samples generated will typically have lower autocorrelation than those generated by the `HMC` kernel. Optionally, the NUTS kernel also provides the ability to adapt step size during the warmup phase.

Refer to the [baseball example](#) to see how to do Bayesian inference in Pyro using NUTS.

References

- [1] *The No-U-turn sampler: adaptively setting path lengths in Hamiltonian Monte Carlo*, Matthew D. Hoffman, and Andrew Gelman.
- [2] *A Conceptual Introduction to Hamiltonian Monte Carlo*, Michael Betancourt
- [3] *Slice Sampling*, Radford M. Neal

Parameters

- **model** – Python callable containing Pyro primitives.
- **potential_fn** – Python callable calculating potential energy with input is a dict of real support parameters.
- **step_size** (*float*) – Determines the size of a single step taken by the verlet integrator while computing the trajectory using Hamiltonian dynamics. If not specified, it will be set to 1.
- **adapt_step_size** (*bool*) – A flag to decide if we want to adapt step_size during warm-up phase using Dual Averaging scheme.
- **adapt_mass_matrix** (*bool*) – A flag to decide if we want to adapt mass matrix during warm-up phase using Welford scheme.
- **full_mass** (*bool*) – A flag to decide if mass matrix is dense or diagonal.
- **use_multinomial_sampling** (*bool*) – A flag to decide if we want to sample candidates along its trajectory using “multinomial sampling” or using “slice sampling”. Slice sampling is used in the original NUTS paper [1], while multinomial sampling is suggested in [2]. By default, this flag is set to True. If it is set to *False*, NUTS uses slice sampling.
- **transforms** (*dict*) – Optional dictionary that specifies a transform for a sample site with constrained support to unconstrained space. The transform should be invertible, and implement *log_abs_det_jacobian*. If not specified and the model has sites with constrained support, automatic transformations will be applied, as specified in `torch.distributions.constraint_registry`.
- **max_plate_nesting** (*int*) – Optional bound on max number of nested `pyro.plate()` contexts. This is required if model contains discrete sample sites that can be enumerated over in parallel.
- **jit_compile** (*bool*) – Optional parameter denoting whether to use the PyTorch JIT to trace the log density computation, and use this optimized executable trace in the integrator.
- **jit_options** (*dict*) – A dictionary contains optional arguments for `torch.jit.trace()` function.
- **ignore_jit_warnings** (*bool*) – Flag to ignore warnings from the JIT tracer when `jit_compile=True`. Default is False.
- **target_accept_prob** (*float*) – Target acceptance probability of step size adaptation scheme. Increasing this value will lead to a smaller step size, so the sampling will be slower but more robust. Default to 0.8.
- **max_tree_depth** (*int*) – Max depth of the binary tree created during the doubling scheme of NUTS sampler. Default to 10.

Example:

```
>>> true_coefs = torch.tensor([1., 2., 3.])
>>> data = torch.randn(2000, 3)
>>> dim = 3
>>> labels = dist.Bernoulli(logits=(true_coefs * data).sum(-1)).sample()
>>>
>>> def model(data):
...     coefs_mean = torch.zeros(dim)
...     coefs = pyro.sample('beta', dist.Normal(coefs_mean, torch.ones(3)))
...     y = pyro.sample('y', dist.Bernoulli(logits=(coefs * data).sum(-1)),
... ↪obs=labels)
...     return y
```

(continues on next page)

(continued from previous page)

```

>>>
>>> nuts_kernel = NUTS(model, adapt_step_size=True)
>>> mcmc = MCMC(nuts_kernel, num_samples=500, warmup_steps=300)
>>> mcmc.run(data)
>>> mcmc.get_samples()['beta'].mean(0) # doctest: +SKIP
tensor([ 0.9221,  1.9464,  2.9228])

```

sample (*params*)

4.6.4 Utilities

initialize_model (*model*, *model_args*=(), *model_kwargs*={}, *transforms*=None, *max_plate_nesting*=None, *jit_compile*=False, *jit_options*=None, *skip_jit_warnings*=False, *num_chains*=1)

Given a Python callable with Pyro primitives, generates the following model-specific properties needed for inference using HMC/NUTS kernels:

- initial parameters to be sampled using a HMC kernel,
- a potential function whose input is a dict of parameters in unconstrained space,
- transforms to transform latent sites of *model* to unconstrained space,
- a prototype trace to be used in MCMC to consume traces from sampled parameters.

Parameters

- **model** – a Pyro model which contains Pyro primitives.
- **model_args** (*tuple*) – optional args taken by *model*.
- **model_kwargs** (*dict*) – optional kwargs taken by *model*.
- **transforms** (*dict*) – Optional dictionary that specifies a transform for a sample site with constrained support to unconstrained space. The transform should be invertible, and implement *log_abs_det_jacobian*. If not specified and the model has sites with constrained support, automatic transformations will be applied, as specified in `torch.distributions.constraint_registry`.
- **max_plate_nesting** (*int*) – Optional bound on max number of nested *pyro.plate()* contexts. This is required if model contains discrete sample sites that can be enumerated over in parallel.
- **jit_compile** (*bool*) – Optional parameter denoting whether to use the PyTorch JIT to trace the log density computation, and use this optimized executable trace in the integrator.
- **jit_options** (*dict*) – A dictionary contains optional arguments for `torch.jit.trace()` function.
- **ignore_jit_warnings** (*bool*) – Flag to ignore warnings from the JIT tracer when `jit_compile=True`. Default is False.
- **num_chains** (*int*) – Number of parallel chains. If *num_chains* > 1, the returned *initial_params* will be a list with *num_chains* elements.

Returns a tuple of (*initial_params*, *potential_fn*, *transforms*, *prototype_trace*)

diagnostics (*samples*, *num_chains=1*)

Gets diagnostics statistics such as effective sample size and split Gelman-Rubin using the samples drawn from the posterior distribution.

Parameters

- **samples** (*dict*) – dictionary of samples keyed by site name.
- **num_chains** (*int*) – number of chains. For more than a single chain, the leading dimension of samples in *samples* must match the number of chains.

Returns dictionary of diagnostic stats for each sample site.

predictive (*model*, *posterior_samples*, **args*, ***kwargs*)

Run model by sampling latent parameters from *posterior_samples*, and return values at sample sites from the forward run. By default, only sites not contained in *posterior_samples* are returned. This can be modified by changing the *return_sites* keyword argument.

Warning: The interface for the *predictive* class is experimental, and might change in the future. e.g. a unified interface for predictive with SVI.

Parameters

- **model** – Python callable containing Pyro primitives.
- **posterior_samples** (*dict*) – dictionary of samples from the posterior.
- **args** – model arguments.
- **kwargs** – model kwargs; and other keyword arguments (see below).

Keyword Arguments

- **num_samples** (*int*) - number of samples to draw from the predictive distribution. This argument has no effect if *posterior_samples* is non-empty, in which case, the leading dimension size of samples in *posterior_samples* is used.
- **return_sites** (*list*) - sites to return; by default only sample sites not present in *posterior_samples* are returned.
- **return_trace** (*bool*) - whether to return the full trace. Note that this is vectorized over *num_samples*.

Returns dict of samples from the predictive distribution, or a single vectorized *trace* (if *return_trace=True*).

5.1 PyTorch Distributions

Most distributions in Pyro are thin wrappers around PyTorch distributions. For details on the PyTorch distribution interface, see `torch.distributions.distribution.Distribution`. For differences between the Pyro and PyTorch interfaces, see *TorchDistributionMixin*.

5.1.1 Bernoulli

```
class Bernoulli (probs=None, logits=None, validate_args=None)  
    Wraps torch.distributions.bernoulli.Bernoulli with TorchDistributionMixin.
```

5.1.2 Beta

```
class Beta (concentration1, concentration0, validate_args=None)  
    Wraps torch.distributions.beta.Beta with TorchDistributionMixin.
```

5.1.3 Binomial

```
class Binomial (total_count=1, probs=None, logits=None, validate_args=None)  
    Wraps torch.distributions.binomial.Binomial with TorchDistributionMixin.
```

5.1.4 Categorical

```
class Categorical (probs=None, logits=None, validate_args=None)  
    Wraps torch.distributions.categorical.Categorical with TorchDistributionMixin.
```

5.1.5 Cauchy

```
class Cauchy (loc, scale, validate_args=None)  
    Wraps torch.distributions.cauchy.Cauchy with TorchDistributionMixin.
```

5.1.6 Chi2

```
class Chi2 (df, validate_args=None)  
    Wraps torch.distributions.chi2.Chi2 with TorchDistributionMixin.
```

5.1.7 Dirichlet

```
class Dirichlet (concentration, validate_args=None)  
    Wraps torch.distributions.dirichlet.Dirichlet with TorchDistributionMixin.
```

5.1.8 Exponential

```
class Exponential (rate, validate_args=None)  
    Wraps torch.distributions.exponential.Exponential with TorchDistributionMixin.
```

5.1.9 ExponentialFamily

```
class ExponentialFamily (batch_shape=torch.Size([]), event_shape=torch.Size([]),  
                        date_args=None)  
    Wraps torch.distributions.exp_family.ExponentialFamily with  
    TorchDistributionMixin.
```

5.1.10 FisherSnedecor

```
class FisherSnedecor (df1, df2, validate_args=None)  
    Wraps torch.distributions.fishersnedecor.FisherSnedecor with  
    TorchDistributionMixin.
```

5.1.11 Gamma

```
class Gamma (concentration, rate, validate_args=None)  
    Wraps torch.distributions.gamma.Gamma with TorchDistributionMixin.
```

5.1.12 Geometric

```
class Geometric (probs=None, logits=None, validate_args=None)  
    Wraps torch.distributions.geometric.Geometric with TorchDistributionMixin.
```

5.1.13 Gumbel

```
class Gumbel (loc, scale, validate_args=None)  
    Wraps torch.distributions.gumbel.Gumbel with TorchDistributionMixin.
```

5.1.14 HalfCauchy

```
class HalfCauchy (scale, validate_args=None)
    Wraps torch.distributions.half_cauchy.HalfCauchy with TorchDistributionMixin.
```

5.1.15 HalfNormal

```
class HalfNormal (scale, validate_args=None)
    Wraps torch.distributions.half_normal.HalfNormal with TorchDistributionMixin.
```

5.1.16 Independent

```
class Independent (base_distribution, reinterpreted_batch_ndims, validate_args=None)
    Wraps torch.distributions.independent.Independent with TorchDistributionMixin.
```

5.1.17 Laplace

```
class Laplace (loc, scale, validate_args=None)
    Wraps torch.distributions.laplace.Laplace with TorchDistributionMixin.
```

5.1.18 LogNormal

```
class LogNormal (loc, scale, validate_args=None)
    Wraps torch.distributions.log_normal.LogNormal with TorchDistributionMixin.
```

5.1.19 LogisticNormal

```
class LogisticNormal (loc, scale, validate_args=None)
    Wraps torch.distributions.logistic_normal.LogisticNormal with
    TorchDistributionMixin.
```

5.1.20 LowRankMultivariateNormal

```
class LowRankMultivariateNormal (loc, cov_factor, cov_diag, validate_args=None)
    Wraps torch.distributions.lowrank_multivariate_normal.
    LowRankMultivariateNormal with TorchDistributionMixin.
```

5.1.21 Multinomial

```
class Multinomial (total_count=1, probs=None, logits=None, validate_args=None)
    Wraps torch.distributions.multinomial.Multinomial with TorchDistributionMixin.
```

5.1.22 MultivariateNormal

```
class MultivariateNormal (loc, covariance_matrix=None, precision_matrix=None, scale_tril=None,
    validate_args=None)
    Wraps torch.distributions.multivariate_normal.MultivariateNormal with
    TorchDistributionMixin.
```

5.1.23 NegativeBinomial

```
class NegativeBinomial (total_count, probs=None, logits=None, validate_args=None)
    Wraps      torch.distributions.negative_binomial.NegativeBinomial      with
    TorchDistributionMixin.
```

5.1.24 Normal

```
class Normal (loc, scale, validate_args=None)
    Wraps torch.distributions.normal.Normal with TorchDistributionMixin.
```

5.1.25 OneHotCategorical

```
class OneHotCategorical (probs=None, logits=None, validate_args=None)
    Wraps      torch.distributions.one_hot_categorical.OneHotCategorical      with
    TorchDistributionMixin.
```

5.1.26 Pareto

```
class Pareto (scale, alpha, validate_args=None)
    Wraps torch.distributions.pareto.Pareto with TorchDistributionMixin.
```

5.1.27 Poisson

```
class Poisson (rate, validate_args=None)
    Wraps torch.distributions.poisson.Poisson with TorchDistributionMixin.
```

5.1.28 RelaxedBernoulli

```
class RelaxedBernoulli (temperature, probs=None, logits=None, validate_args=None)
    Wraps      torch.distributions.relaxed_bernoulli.RelaxedBernoulli      with
    TorchDistributionMixin.
```

5.1.29 RelaxedOneHotCategorical

```
class RelaxedOneHotCategorical (temperature, probs=None, logits=None, validate_args=None)
    Wraps torch.distributions.relaxed_categorical.RelaxedOneHotCategorical with
    TorchDistributionMixin.
```

5.1.30 StudentT

```
class StudentT (df, loc=0.0, scale=1.0, validate_args=None)
    Wraps torch.distributions.studentT.StudentT with TorchDistributionMixin.
```

5.1.31 TransformedDistribution

```
class TransformedDistribution (base_distribution, transforms, validate_args=None)
    Wraps torch.distributions.transformed_distribution.TransformedDistribution
    with TorchDistributionMixin.
```

5.1.32 Uniform

```
class Uniform (low, high, validate_args=None)
    Wraps torch.distributions.uniform.Uniform with TorchDistributionMixin.
```

5.1.33 Weibull

```
class Weibull (scale, concentration, validate_args=None)
    Wraps torch.distributions.weibull.Weibull with TorchDistributionMixin.
```

5.2 Pyro Distributions

5.2.1 Abstract Distribution

```
class Distribution
```

Bases: `object`

Base class for parameterized probability distributions.

Distributions in Pyro are stochastic function objects with `sample()` and `log_prob()` methods. Distribution are stochastic functions with fixed parameters:

```
d = dist.Bernoulli(param)
x = d()                                # Draws a random sample.
p = d.log_prob(x)                      # Evaluates log probability of x.
```

Implementing New Distributions:

Derived classes must implement the methods: `sample()`, `log_prob()`.

Examples:

Take a look at the [examples](#) to see how they interact with inference algorithms.

```
__call__ (*args, **kwargs)
```

Samples a random value (just an alias for `.sample(*args, **kwargs)`).

For tensor distributions, the returned tensor should have the same `.shape` as the parameters.

Returns A random value.

Return type `torch.Tensor`

```
enumerate_support (expand=True)
```

Returns a representation of the parametrized distribution's support, along the first dimension. This is implemented only by discrete distributions.

Note that this returns support values of all the batched RVs in lock-step, rather than the full cartesian product.

Parameters `expand (bool)` – whether to expand the result to a tensor of shape $(n,) + \text{batch_shape} + \text{event_shape}$. If false, the return value has unexpanded shape $(n,) + (1,) * \text{len}(\text{batch_shape}) + \text{event_shape}$ which can be broadcasted to the full shape.

Returns An iterator over the distribution’s discrete support.

Return type iterator

`has_enumerate_support = False`

`has_rsample = False`

`log_prob (x, *args, **kwargs)`

Evaluates log probability densities for each of a batch of samples.

Parameters `x (torch.Tensor)` – A single value or a batch of values batched along axis 0.

Returns log probability densities as a one-dimensional `Tensor` with same batch size as value and params. The shape of the result should be `self.batch_size`.

Return type `torch.Tensor`

`sample (*args, **kwargs)`

Samples a random value.

For tensor distributions, the returned tensor should have the same `.shape` as the parameters, unless otherwise noted.

Parameters `sample_shape (torch.Size)` – the size of the iid batch to be drawn from the distribution.

Returns A random value or batch of random values (if parameters are batched). The shape of the result should be `self.shape()`.

Return type `torch.Tensor`

`score_parts (x, *args, **kwargs)`

Computes ingredients for stochastic gradient estimators of ELBO.

The default implementation is correct both for non-reparameterized and for fully reparameterized distributions. Partially reparameterized distributions should override this method to compute correct `.score_function` and `.entropy_term` parts.

Parameters `x (torch.Tensor)` – A single value or batch of values.

Returns A `ScoreParts` object containing parts of the ELBO estimator.

Return type `ScoreParts`

5.2.2 TorchDistributionMixin

class `TorchDistributionMixin`

Bases: `pyro.distributions.distribution.Distribution`

Mixin to provide Pyro compatibility for PyTorch distributions.

You should instead use *TorchDistribution* for new distribution classes.

This is mainly useful for wrapping existing PyTorch distributions for use in Pyro. Derived classes must first inherit from `torch.distributions.distribution.Distribution` and then inherit from *TorchDistributionMixin*.

__call__ (*sample_shape=torch.Size([])*)

Samples a random value.

This is reparameterized whenever possible, calling `rsample()` for reparameterized distributions and `sample()` for non-reparameterized distributions.

Parameters **sample_shape** (*torch.Size*) – the size of the iid batch to be drawn from the distribution.

Returns A random value or batch of random values (if parameters are batched). The shape of the result should be *self.shape()*.

Return type `torch.Tensor`

event_dim

Returns Number of dimensions of individual events.

Return type `int`

shape (*sample_shape=torch.Size([])*)

The tensor shape of samples from this distribution.

Samples are of shape:

```
d.shape(sample_shape) == sample_shape + d.batch_shape + d.event_shape
```

Parameters **sample_shape** (*torch.Size*) – the size of the iid batch to be drawn from the distribution.

Returns Tensor shape of samples.

Return type `torch.Size`

expand_by (*sample_shape*)

Expands a distribution by adding *sample_shape* to the left side of its *batch_shape*.

To expand internal dims of *self.batch_shape* from 1 to something larger, use `expand()` instead.

Parameters **sample_shape** (*torch.Size*) – The size of the iid batch to be drawn from the distribution.

Returns An expanded version of this distribution.

Return type `ReshapedDistribution`

reshape (*sample_shape=None, extra_event_dims=None*)

to_event (*reinterpreted_batch_ndims=None*)

Reinterprets the *n* rightmost dimensions of this distributions *batch_shape* as event dims, adding them to the left side of *event_shape*.

Example:

```
>>> [d1.batch_shape, d1.event_shape]
[torch.Size([2, 3]), torch.Size([4, 5])]
>>> d2 = d1.to_event(1)
>>> [d2.batch_shape, d2.event_shape]
[torch.Size([2]), torch.Size([3, 4, 5])]
>>> d3 = d1.to_event(2)
>>> [d3.batch_shape, d3.event_shape]
[torch.Size([]), torch.Size([2, 3, 4, 5])]
```

Parameters `reinterpreted_batch_ndims` (*int*) – The number of batch dimensions to reinterpret as event dimensions.

Returns A reshaped version of this distribution.

Return type `pyro.distributions.torch.Independent`

independent (*reinterpreted_batch_ndims=None*)

mask (*mask*)

Masks a distribution by a zero-one tensor that is broadcastable to the distributions `batch_shape`.

Parameters `mask` (*torch.Tensor*) – A zero-one valued float tensor.

Returns A masked copy of this distribution.

Return type `MaskedDistribution`

5.2.3 TorchDistribution

class `TorchDistribution` (*batch_shape=torch.Size([])*, *event_shape=torch.Size([])*, *validate_args=None*)

Bases: `torch.distributions.distribution.Distribution`, `pyro.distributions.torch_distribution.TorchDistributionMixin`

Base class for PyTorch-compatible distributions with Pyro support.

This should be the base class for almost all new Pyro distributions.

Note: Parameters and data should be of type `Tensor` and all methods return type `Tensor` unless otherwise noted.

Tensor Shapes:

TorchDistributions provide a method `.shape()` for the tensor shape of samples:

```
x = d.sample(sample_shape)
assert x.shape == d.shape(sample_shape)
```

Pyro follows the same distribution shape semantics as PyTorch. It distinguishes between three different roles for tensor shapes of samples:

- *sample shape* corresponds to the shape of the iid samples drawn from the distribution. This is taken as an argument by the distribution's *sample* method.
- *batch shape* corresponds to non-identical (independent) parameterizations of the distribution, inferred from the distribution's parameter shapes. This is fixed for a distribution instance.
- *event shape* corresponds to the event dimensions of the distribution, which is fixed for a distribution class. These are collapsed when we try to score a sample from the distribution via *d.log_prob(x)*.

These shapes are related by the equation:

```
assert d.shape(sample_shape) == sample_shape + d.batch_shape + d.event_shape
```

Distributions provide a vectorized `log_prob()` method that evaluates the log probability density of each event in a batch independently, returning a tensor of shape `sample_shape + d.batch_shape`:


```
x = d.sample(sample_shape)
assert x.shape == d.shape(sample_shape)
log_p = d.log_prob(x)
assert log_p.shape == sample_shape + d.batch_shape
```

Implementing New Distributions:

Derived classes must implement the methods `sample()` (or `rsample()` if `.has_rsample == True`) and `log_prob()`, and must implement the properties `batch_shape`, and `event_shape`. Discrete classes may also implement the `enumerate_support()` method to improve gradient estimates and set `.has_enumerate_support = True`.

5.2.4 AVFMultivariateNormal

class AVFMultivariateNormal (*loc, scale_tril, control_var*)

Bases: `pyro.distributions.torch.MultivariateNormal`

Multivariate normal (Gaussian) distribution with transport equation inspired control variates (adaptive velocity fields).

A distribution over vectors in which all the elements have a joint Gaussian density.

Parameters

- **loc** (*torch.Tensor*) – D-dimensional mean vector.
- **scale_tril** (*torch.Tensor*) – Cholesky of Covariance matrix; D x D matrix.
- **control_var** (*torch.Tensor*) – 2 x L x D tensor that parameterizes the control variate; L is an arbitrary positive integer. This parameter needs to be learned (i.e. adapted) to achieve lower variance gradients. In a typical use case this parameter will be adapted concurrently with the *loc* and *scale_tril* that define the distribution.

Example usage:

```
control_var = torch.tensor(0.1 * torch.ones(2, 1, D), requires_grad=True)
opt_cv = torch.optim.Adam([control_var], lr=0.1, betas=(0.5, 0.999))

for _ in range(1000):
    d = AVFMultivariateNormal(loc, scale_tril, control_var)
    z = d.rsample()
    cost = torch.pow(z, 2.0).sum()
    cost.backward()
    opt_cv.step()
    opt_cv.zero_grad()
```

```
arg_constraints = {'control_var': Real(), 'loc': Real(), 'scale_tril': LowerTriangular}
rsample (sample_shape=torch.Size([]))
```

5.2.5 BetaBinomial

class BetaBinomial (*concentration1, concentration0, total_count=1, validate_args=None*)

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Compound distribution comprising of a beta-binomial pair. The probability of success (`probs` for the *Binomial* distribution) is unknown and randomly drawn from a *Beta* distribution prior to a certain number of Bernoulli trials given by `total_count`.

Parameters

- or `torch.Tensor concentration1 (float)` – 1st concentration parameter (alpha) for the Beta distribution.
- or `torch.Tensor concentration0 (float)` – 2nd concentration parameter (beta) for the Beta distribution.
- or `torch.Tensor total_count (int)` – number of Bernoulli trials.

```
arg_constraints = {'concentration0': GreaterThan(lower_bound=0.0), 'concentration1':  
concentration0  
concentration1  
enumerate_support (expand=True)  
expand (batch_shape, _instance=None)  
has_enumerate_support = True  
log_prob (value)  
mean  
sample (sample_shape=())  
support  
variance
```

5.2.6 Delta

```
class Delta (v, log_density=0.0, event_dim=0, validate_args=None)
```

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Degenerate discrete distribution (a single point).

Discrete distribution that assigns probability one to the single element in its support. Delta distribution parameterized by a random choice should not be used with MCMC based inference, as doing so produces incorrect results.

Parameters

- `v (torch.Tensor)` – The single support element.
- `log_density (torch.Tensor)` – An optional density for this Delta. This is useful to keep the class of *Delta* distributions closed under differentiable transformation.
- `event_dim (int)` – Optional event dimension, defaults to zero.

```
arg_constraints = {'log_density': Real(), 'v': Real()}  
expand (batch_shape, _instance=None)  
has_rsample = True  
log_prob (x)  
mean  
rsample (sample_shape=torch.Size([]))  
support = Real()  
variance
```

5.2.7 DirichletMultinomial

```
class DirichletMultinomial (concentration, total_count=1, is_sparse=False, validate_args=None)
```

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Compound distribution comprising of a dirichlet-multinomial pair. The probability of classes (`probs` for the *Multinomial* distribution) is unknown and randomly drawn from a *Dirichlet* distribution prior to a certain number of Categorical trials given by `total_count`.

Parameters

- **or** `torch.Tensor concentration` (*float*) – concentration parameter (alpha) for the Dirichlet distribution.
- **or** `torch.Tensor total_count` (*int*) – number of Categorical trials.
- **is_sparse** (*bool*) – Whether to assume value is mostly zero when computing `log_prob()`, which can speed up computation when data is sparse.

```
arg_constraints = {'concentration': GreaterThan(lower_bound=0.0), 'total_count': Int
```

```
concentration
```

```
expand (batch_shape, _instance=None)
```

```
log_prob (value)
```

```
mean
```

```
sample (sample_shape=())
```

```
support
```

```
variance
```

5.2.8 EmpiricalDistribution

```
class Empirical (samples, log_weights, validate_args=None)
```

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Empirical distribution associated with the sampled data. Note that the shape requirement for `log_weights` is that its shape must match the leftmost shape of `samples`. Samples are aggregated along the `aggregation_dim`, which is the rightmost dim of `log_weights`.

Example:

```
>>> emp_dist = Empirical(torch.randn(2, 3, 10), torch.ones(2, 3))
>>> emp_dist.batch_shape
torch.Size([2])
>>> emp_dist.event_shape
torch.Size([10])
```

```
>>> single_sample = emp_dist.sample()
>>> single_sample.shape
torch.Size([2, 10])
>>> batch_sample = emp_dist.sample((100,))
>>> batch_sample.shape
torch.Size([100, 2, 10])
```

```
>>> emp_dist.log_prob(single_sample).shape
torch.Size([2])
>>> # Vectorized samples cannot be scored by log_prob.
>>> with pyro.validation_enabled():
...     emp_dist.log_prob(batch_sample).shape
Traceback (most recent call last):
...
ValueError: ``value.shape`` must be torch.Size([2, 10])
```

Parameters

- **samples** (*torch.Tensor*) – samples from the empirical distribution.
- **log_weights** (*torch.Tensor*) – log weights (optional) corresponding to the samples.

arg_constraints = {}

enumerate_support (*expand=True*)

See `pyro.distributions.torch_distribution.TorchDistribution.enumerate_support()`

event_shape

See `pyro.distributions.torch_distribution.TorchDistribution.event_shape()`

has_enumerate_support = True

log_prob (*value*)

Returns the log of the probability mass function evaluated at *value*. Note that this currently only supports scoring values with empty *sample_shape*.

Parameters *value* (*torch.Tensor*) – scalar or tensor value to be scored.

log_weights

mean

See `pyro.distributions.torch_distribution.TorchDistribution.mean()`

sample (*sample_shape=torch.Size([])*)

See `pyro.distributions.torch_distribution.TorchDistribution.sample()`

sample_size

Number of samples that constitute the empirical distribution.

Return int number of samples collected.

support = Real()

variance

See `pyro.distributions.torch_distribution.TorchDistribution.variance()`

5.2.9 GammaPoisson

class GammaPoisson (*concentration, rate, validate_args=None*)

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Compound distribution comprising of a gamma-poisson pair, also referred to as a gamma-poisson mixture. The *rate* parameter for the *Poisson* distribution is unknown and randomly drawn from a *Gamma* distribution.

Note: This can be treated as an alternate parametrization of the [NegativeBinomial](#) (`total_count`, `probs`) distribution, with $concentration = total_count$ and $rate = (1 - probs) / probs$.

Parameters

- or `torch.Tensor concentration (float)` – shape parameter (alpha) of the Gamma distribution.
- or `torch.Tensor rate (float)` – rate parameter (beta) for the Gamma distribution.

```
arg_constraints = {'concentration': GreaterThan(lower_bound=0.0), 'rate': GreaterThan(
concentration
expand (batch_shape, _instance=None)
log_prob (value)
mean
rate
sample (sample_shape=())
support = IntegerGreaterThan(lower_bound=0)
variance
```

5.2.10 GaussianScaleMixture

```
class GaussianScaleMixture (coord_scale, component_logits, component_scale)
Bases: pyro.distributions.torch_distribution.TorchDistribution
```

Mixture of Normal distributions with zero mean and diagonal covariance matrices.

That is, this distribution is a mixture with K components, where each component distribution is a D -dimensional Normal distribution with zero mean and a D -dimensional diagonal covariance matrix. The K different covariance matrices are controlled by the parameters `coord_scale` and `component_scale`. That is, the covariance matrix of the k 'th component is given by

$$\text{Sigma}_{ii} = (\text{component_scale}_k * \text{coord_scale}_i) ** 2 \quad (i = 1, \dots, D)$$

where `component_scale_k` is a positive scale factor and `coord_scale_i` are positive scale parameters shared between all K components. The mixture weights are controlled by a K -dimensional vector of softmax logits, `component_logits`. This distribution implements pathwise derivatives for samples from the distribution. This distribution does not currently support batched parameters.

See reference [1] for details on the implementations of the pathwise derivative. Please consider citing this reference if you use the pathwise derivative in your research.

[1] Pathwise Derivatives for Multivariate Distributions, Martin Jankowiak & Theofanis Karaletsos. arXiv:1806.01856

Note that this distribution supports both even and odd dimensions, but the former should be more a bit higher precision, since it doesn't use any erfs in the backward call. Also note that this distribution does not support $D = 1$.

Parameters

- `coord_scale (torch.tensor)` – D -dimensional vector of scales

- **component_logits** (*torch.tensor*) – K-dimensional vector of logits
- **component_scale** (*torch.tensor*) – K-dimensional vector of scale multipliers

```
arg_constraints = {'component_logits': Real(), 'component_scale': GreaterThan(lower_bound=0.0)}
has_rsample = True
log_prob(value)
rsample(sample_shape=torch.Size([]))
```

5.2.11 InverseGamma

```
class InverseGamma(concentration, rate, validate_args=None)
    Bases: pyro.distributions.torch.TransformedDistribution
    Creates an inverse-gamma distribution parameterized by concentration and rate.
```

$X \sim \text{Gamma}(\text{concentration}, \text{rate})$ $Y = 1/X \sim \text{InverseGamma}(\text{concentration}, \text{rate})$

Parameters

- **concentration** (*torch.Tensor*) – the concentration parameter (i.e. alpha).
- **rate** (*torch.Tensor*) – the rate parameter (i.e. beta).

```
arg_constraints = {'concentration': GreaterThan(lower_bound=0.0), 'rate': GreaterThan(lower_bound=0.0)}
concentration
expand(batch_shape, _instance=None)
has_rsample = True
rate
support = GreaterThan(lower_bound=0.0)
```

5.2.12 LKJCorrCholesky

```
class LKJCorrCholesky(d, eta, validate_args=None)
    Bases: pyro.distributions.torch_distribution.TorchDistribution
```

Generates cholesky factors of correlation matrices using an LKJ prior.

The expected use is to combine it with a vector of variances and pass it to the `scale_tril` parameter of a multivariate distribution such as `MultivariateNormal`.

E.g., if `theta` is a (positive) vector of covariances with the same dimensionality as this distribution, and `Omega` is sampled from this distribution, `scale_tril=torch.mm(torch.diag(sqrt(theta)), Omega)`

Note that the *event_shape* of this distribution is $[d, d]$

Note: When using this distribution with HMC/NUTS, it is important to use a *step_size* such as $1e-4$. If not, you are likely to experience LAPACK errors regarding positive-definiteness.

For example usage, refer to [pyro/examples/lkj.py](#).

Parameters

- **d** (*int*) – Dimensionality of the matrix
- **eta** (*torch.Tensor*) – A single positive number parameterizing the distribution.

```

arg_constraints = {'eta': GreaterThan(lower_bound=0.0)}
expand(batch_shape, _instance=None)
has_rsample = False
lkj_constant(eta, K)
log_prob(x)
sample(sample_shape=torch.Size([]))
support = CorrCholesky()

```

5.2.13 MaskedMixture

class MaskedMixture (*mask, component0, component1, validate_args=None*)
 Bases: `pyro.distributions.torch_distribution.TorchDistribution`

A masked deterministic mixture of two distributions.

This is useful when the mask is sampled from another distribution, possibly correlated across the batch. Often the mask can be marginalized out via enumeration.

Example:

```

change_point = pyro.sample("change_point",
                           dist.Categorical(torch.ones(len(data) + 1)),
                           infer={'enumerate': 'parallel'})
mask = torch.arange(len(data), dtype=torch.long) >= change_point
with pyro.plate("data", len(data)):
    pyro.sample("obs", MaskedMixture(mask, dist1, dist2), obs=data)

```

Parameters

- **mask** (*torch.Tensor*) – A byte tensor toggling between `component0` and `component1`.
- **component0** (*pyro.distributions.TorchDistribution*) – a distribution for batch elements `mask == 0`.
- **component1** (*pyro.distributions.TorchDistribution*) – a distribution for batch elements `mask == 1`.

```

arg_constraints = {}
expand(batch_shape)
has_rsample
log_prob(value)
mean
rsample(sample_shape=torch.Size([]))
sample(sample_shape=torch.Size([]))
support

```

variance

5.2.14 MixtureOfDiagNormals

class MixtureOfDiagNormals (*locs*, *coord_scale*, *component_logits*)

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Mixture of Normal distributions with arbitrary means and arbitrary diagonal covariance matrices.

That is, this distribution is a mixture with K components, where each component distribution is a D -dimensional Normal distribution with a D -dimensional mean parameter and a D -dimensional diagonal covariance matrix. The K different component means are gathered into the $K \times D$ dimensional parameter *locs* and the K different scale parameters are gathered into the $K \times D$ dimensional parameter *coord_scale*. The mixture weights are controlled by a K -dimensional vector of softmax logits, *component_logits*. This distribution implements pathwise derivatives for samples from the distribution.

See reference [1] for details on the implementations of the pathwise derivative. Please consider citing this reference if you use the pathwise derivative in your research. Note that this distribution does not support dimension $D = 1$.

[1] Pathwise Derivatives for Multivariate Distributions, Martin Jankowiak & Theofanis Karaletsos. arXiv:1806.01856

Parameters

- **locs** (*torch.Tensor*) – $K \times D$ mean matrix
- **coord_scale** (*torch.Tensor*) – $K \times D$ scale matrix
- **component_logits** (*torch.Tensor*) – K -dimensional vector of softmax logits

arg_constraints = {'component_logits': `Real()`, 'coord_scale': `GreaterThan(lower_bound=0)`}

expand (*batch_shape*, *_instance=None*)

has_rsample = `True`

log_prob (*value*)

rsample (*sample_shape=torch.Size([])*)

5.2.15 MixtureOfDiagNormalsSharedCovariance

class MixtureOfDiagNormalsSharedCovariance (*locs*, *coord_scale*, *component_logits*)

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Mixture of Normal distributions with diagonal covariance matrices.

That is, this distribution is a mixture with K components, where each component distribution is a D -dimensional Normal distribution with a D -dimensional mean parameter *loc* and a D -dimensional diagonal covariance matrix specified by a scale parameter *coord_scale*. The K different component means are gathered into the parameter *locs* and the scale parameter is shared between all K components. The mixture weights are controlled by a K -dimensional vector of softmax logits, *component_logits*. This distribution implements pathwise derivatives for samples from the distribution.

See reference [1] for details on the implementations of the pathwise derivative. Please consider citing this reference if you use the pathwise derivative in your research. Note that this distribution does not support dimension $D = 1$.

[1] Pathwise Derivatives for Multivariate Distributions, Martin Jankowiak & Theofanis Karaletsos. arXiv:1806.01856

Parameters

- **locs** (*torch.Tensor*) – K x D mean matrix
- **coord_scale** (*torch.Tensor*) – shared D-dimensional scale vector
- **component_logits** (*torch.Tensor*) – K-dimensional vector of softmax logits

```
arg_constraints = {'component_logits': Real(), 'coord_scale': GreaterThan(lower_bound)}
expand(batch_shape, _instance=None)
has_rsample = True
log_prob(value)
rsample(sample_shape=torch.Size([]))
```

5.2.16 OMTMultivariateNormal

```
class OMTMultivariateNormal(loc, scale_tril)
Bases: pyro.distributions.torch.MultivariateNormal
```

Multivariate normal (Gaussian) distribution with OMT gradients w.r.t. both parameters. Note the gradient computation w.r.t. the Cholesky factor has cost $O(D^3)$, although the resulting gradient variance is generally expected to be lower.

A distribution over vectors in which all the elements have a joint Gaussian density.

Parameters

- **loc** (*torch.Tensor*) – Mean.
- **scale_tril** (*torch.Tensor*) – Cholesky of Covariance matrix.

```
arg_constraints = {'loc': Real(), 'scale_tril': LowerTriangular()}
rsample(sample_shape=torch.Size([]))
```

5.2.17 RelaxedBernoulliStraightThrough

```
class RelaxedBernoulliStraightThrough(temperature, probs=None, logits=None, validate_args=None)
Bases: pyro.distributions.torch.RelaxedBernoulli
```

An implementation of `RelaxedBernoulli` with a straight-through gradient estimator.

This distribution has the following properties:

- The samples returned by the `rsample()` method are discrete/quantized.
- The `log_prob()` method returns the log probability of the relaxed/unquantized sample using the GumbelSoftmax distribution.
- In the backward pass the gradient of the sample with respect to the parameters of the distribution uses the relaxed/unquantized sample.

References:

- [1] **The Concrete Distribution: A Continuous Relaxation of Discrete Random Variables**, Chris J. Maddison, Andriy Mnih, Yee Whye Teh
- [2] **Categorical Reparameterization with Gumbel-Softmax**, Eric Jang, Shixiang Gu, Ben Poole

`log_prob(value)`

See `pyro.distributions.torch.RelaxedBernoulli.log_prob()`

`rsample(sample_shape=torch.Size([]))`

See `pyro.distributions.torch.RelaxedBernoulli.rsample()`

5.2.18 RelaxedOneHotCategoricalStraightThrough

class RelaxedOneHotCategoricalStraightThrough(*temperature, probs=None, logits=None, validate_args=None*)

Bases: `pyro.distributions.torch.RelaxedOneHotCategorical`

An implementation of `RelaxedOneHotCategorical` with a straight-through gradient estimator.

This distribution has the following properties:

- The samples returned by the `rsample()` method are discrete/quantized.
- The `log_prob()` method returns the log probability of the relaxed/unquantized sample using the GumbelSoftmax distribution.
- In the backward pass the gradient of the sample with respect to the parameters of the distribution uses the relaxed/unquantized sample.

References:

[1] **The Concrete Distribution: A Continuous Relaxation of Discrete Random Variables**, Chris J. Maddison, Andriy Mnih, Yee Whye Teh

[2] **Categorical Reparameterization with Gumbel-Softmax**, Eric Jang, Shixiang Gu, Ben Poole

`log_prob(value)`

See `pyro.distributions.torch.RelaxedOneHotCategorical.log_prob()`

`rsample(sample_shape=torch.Size([]))`

See `pyro.distributions.torch.RelaxedOneHotCategorical.rsample()`

5.2.19 Rejector

class Rejector(*propose, log_prob_accept, log_scale*)

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Rejection sampled distribution given an acceptance rate function.

Parameters

- **propose** (`Distribution`) – A proposal distribution that samples batched proposals via `propose()`. `rsample()` supports a `sample_shape` arg only if `propose()` supports a `sample_shape` arg.
- **log_prob_accept** (`callable`) – A callable that inputs a batch of proposals and returns a batch of log acceptance probabilities.
- **log_scale** – Total log probability of acceptance.

`has_rsample = True`

`log_prob(x)`

`rsample(sample_shape=torch.Size([]))`

`score_parts(x)`

5.2.20 SpanningTree

class SpanningTree (*edge_logits*, *sampler_options=None*, *validate_args=None*)
 Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Distribution over spanning trees on a fixed number V of vertices.

A tree is represented as `torch.LongTensor edges` of shape $(V-1, 2)$ satisfying the following properties:

1. The edges constitute a tree, i.e. are connected and cycle free.
2. Each edge $(v1, v2) = \text{edges}[e]$ is sorted, i.e. $v1 < v2$.
3. The entire tensor is sorted in colexicographic order.

Use `validate_edges()` to verify *edges* are correctly formed.

The `edge_logits` tensor has one entry for each of the $V*(V-1)//2$ edges in the complete graph on V vertices, where edges are each sorted and the edge order is colexicographic:

```
(0,1), (0,2), (1,2), (0,3), (1,3), (2,3), (0,4), (1,4), (2,4), ...
```

This ordering corresponds to the size-independent pairing function:

```
k = v1 + v2 * (v2 - 1) // 2
```

where k is the rank of the edge $(v1, v2)$ in the complete graph. To convert a matrix of edge logits to the linear representation used here:

```
assert my_matrix.shape == (V, V)
i, j = make_complete_graph(V)
edge_logits = my_matrix[i, j]
```

Parameters

- **edge_logits** (*torch.Tensor*) – A tensor of length $V*(V-1)//2$ containing logits (aka negative energies) of all edges in the complete graph on V vertices. See above comment for edge ordering.
- **sampler_options** (*dict*) – An optional dict of sampler options including: `mcmc_steps` defaulting to a single MCMC step (which is pretty good); `initial_edges` defaulting to a cheap approximate sample; `backend` one of “python” or “cpp”, defaulting to “python”.

arg_constraints = {'edge_logits': `Real()`}

enumerate_support (*expand=True*)

This is implemented for trees with up to 6 vertices (and 5 edges).

has_enumerate_support = `True`

log_partition_function

log_prob (*edges*)

sample (*sample_shape=torch.Size([])*)

This sampler is implemented using MCMC run for a small number of steps after being initialized by a cheap approximate sampler. This sampler is approximate and cubic time. This is faster than the classic Aldous-Broder sampler [1,2], especially for graphs with large mixing time. Recent research [3,4] proposes samplers that run in sub-matrix-multiply time but are more complex to implement.

References

- [1] *Generating random spanning trees* Andrei Broder (1989)
- [2] *The Random Walk Construction of Uniform Spanning Trees and Uniform Labelled Trees*, David J. Aldous (1990)
- [3] *Sampling Random Spanning Trees Faster than Matrix Multiplication*, David Durfee, Rasmus Kyng, John Peebles, Anup B. Rao, Sushant Sachdeva (2017) <https://arxiv.org/abs/1611.07451>
- [4] *An almost-linear time algorithm for uniform random spanning tree generation*, Aaron Schild (2017) <https://arxiv.org/abs/1711.06455>

support = IntegerGreaterThan(lower_bound=0)

validate_edges (*edges*)

Validates a batch of edges tensors, as returned by `sample()` or `enumerate_support()` or as input to `log_prob()`.

Parameters *edges* (`torch.LongTensor`) – A batch of edges.

Raises ValueError

Returns None

5.2.21 VonMises

class VonMises (*loc, concentration, validate_args=None*)

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

A circular von Mises distribution.

This implementation uses polar coordinates. The `loc` and `value` args can be any real number (to facilitate unconstrained optimization), but are interpreted as angles modulo 2π .

See `VonMises3D` for a 3D cartesian coordinate cousin of this distribution.

Parameters

- **loc** (`torch.Tensor`) – an angle in radians.
- **concentration** (`torch.Tensor`) – concentration parameter

arg_constraints = {'concentration': GreaterThan(lower_bound=0.0), 'loc': Real() }

expand (*batch_shape*)

has_rsample = False

log_prob (*value*)

mean

The provided mean is the circular one.

sample (***kwargs*)

The sampling algorithm for the von Mises distribution is based on the following paper: Best, D. J., and Nicholas I. Fisher. “Efficient simulation of the von Mises distribution.” *Applied Statistics* (1979): 152-157.

support = Real()

variance

The provided variance is the circular one.

5.2.22 VonMises3D

class VonMises3D (*concentration, validate_args=None*)

Bases: `pyro.distributions.torch_distribution.TorchDistribution`

Spherical von Mises distribution.

This implementation combines the direction parameter and concentration parameter into a single combined parameter that contains both direction and magnitude. The `value` arg is represented in cartesian coordinates: it must be a normalized 3-vector that lies on the 2-sphere.

See [VonMises](#) for a 2D polar coordinate cousin of this distribution.

Currently only `log_prob()` is implemented.

Parameters `concentration` (`torch.Tensor`) – A combined location-and-concentration vector. The direction of this vector is the location, and its magnitude is the concentration.

arg_constraints = {'concentration': `Real()`}

expand (`batch_shape`)

log_prob (`value`)

support = `Real()`

5.3 Transformed Distributions

5.3.1 AffineCoupling

class AffineCoupling (*split_dim, hypernet, log_scale_min_clip=-5.0, log_scale_max_clip=3.0*)

Bases: `pyro.distributions.torch_transform.TransformModule`

An implementation of the affine coupling layer of RealNVP (Dinh et al., 2017) that uses the transformation,

$$\mathbf{y}_{1:d} = \mathbf{x}_{1:d} \quad \mathbf{y}_{(d+1):D} = \mu + \sigma \odot \mathbf{x}_{(d+1):D}$$

where \mathbf{x} are the inputs, \mathbf{y} are the outputs, e.g. $\mathbf{x}_{1:d}$ represents the first d elements of the inputs, and μ, σ are shift and translation parameters calculated as the output of a function inputting only $\mathbf{x}_{1:d}$.

That is, the first d components remain unchanged, and the subsequent $D - d$ are shifted and translated by a function of the previous components.

Together with *TransformedDistribution* this provides a way to create richer variational approximations.

Example usage:

```
>>> from pyro.nn import DenseNN
>>> input_dim = 10
>>> split_dim = 6
>>> base_dist = dist.Normal(torch.zeros(input_dim), torch.ones(input_dim))
>>> hypernet = DenseNN(split_dim, [10*input_dim], [input_dim-split_dim, input_dim-
→split_dim])
>>> flow = AffineCoupling(split_dim, hypernet)
>>> pyro.module("my_flow", flow) # doctest: +SKIP
>>> flow_dist = dist.TransformedDistribution(base_dist, [flow])
>>> flow_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
         0.1389, -0.4629,  0.0986])
```

The inverse of the Bijector is required when, e.g., scoring the log density of a sample with *TransformedDistribution*. This implementation caches the inverse of the Bijector when its forward operation is called, e.g., when sampling from *TransformedDistribution*. However, if the cached value isn't available, either because it was overwritten during sampling a new value or an arbitrary value is being scored, it will calculate it manually.

This is an operation that scales as $O(1)$, i.e. constant in the input dimension. So in general, it is cheap to sample and score (an arbitrary value) from *AffineCoupling*.

Parameters

- **split_dim** (*int*) – Zero-indexed dimension d upon which to perform input/output split for transformation.
- **hypernet** (*callable*) – an autoregressive neural network whose forward call returns a real-valued mean and logit-scale as a tuple. The input should have final dimension `split_dim` and the output final dimension `input_dim-split_dim` for each member of the tuple.
- **log_scale_min_clip** (*float*) – The minimum value for clipping the $\log(\text{scale})$ from the autoregressive NN
- **log_scale_max_clip** (*float*) – The maximum value for clipping the $\log(\text{scale})$ from the autoregressive NN

References:

Laurent Dinh, Jascha Sohl-Dickstein, and Samy Bengio. Density estimation using Real NVP. ICLR 2017.

bijective = True

codomain = `Real()`

domain = `Real()`

event_dim = 1

log_abs_det_jacobian (x, y)

Calculates the elementwise determinant of the log jacobian

5.3.2 BatchNormTransform

class BatchNormTransform (*input_dim, momentum=0.1, epsilon=1e-05*)

Bases: `pyro.distributions.torch_transform.TransformModule`

A type of batch normalization that can be used to stabilize training in normalizing flows. The inverse operation is defined as

$$x = (y - \hat{\mu}) \oslash \sqrt{\hat{\sigma}^2} \otimes \gamma + \beta$$

that is, the standard batch norm equation, where x is the input, y is the output, γ, β are learnable parameters, and $\hat{\mu}/\hat{\sigma}^2$ are smoothed running averages of the sample mean and variance, respectively. The constraint $\gamma > 0$ is enforced to ease calculation of the log-det-Jacobian term.

This is an element-wise transform, and when applied to a vector, learns two parameters (γ, β) for each dimension of the input.

When the module is set to training mode, the moving averages of the sample mean and variance are updated every time the inverse operator is called, e.g., when a normalizing flow scores a minibatch with the *log_prob* method.

Also, when the module is set to training mode, the sample mean and variance on the current minibatch are used in place of the smoothed averages, $\hat{\mu}$ and $\hat{\sigma}^2$, for the inverse operator. For this reason it is not the case that $x = g(g^{-1}(x))$ during training, i.e., that the inverse operation is the inverse of the forward one.

Example usage:

```
>>> from pyro.nn import AutoRegressiveNN
>>> from pyro.distributions import InverseAutoregressiveFlow
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> iaafs = [InverseAutoregressiveFlow(AutoRegressiveNN(10, [40])) for _ in
↳ range(2)]
>>> bn = BatchNormTransform(10)
>>> flow_dist = dist.TransformedDistribution(base_dist, [iaafs[0], bn, iaafs[1]])
>>> flow_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])
```

Parameters

- **input_dim** (*int*) – the dimension of the input
- **momentum** (*float*) – momentum parameter for updating moving averages
- **epsilon** (*float*) – small number to add to variances to ensure numerical stability

References:

- [1] Sergey Ioffe and Christian Szegedy. Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift. In International Conference on Machine Learning, 2015. <https://arxiv.org/abs/1502.03167>
- [2] Laurent Dinh, Jascha Sohl-Dickstein, and Samy Bengio. Density Estimation using Real NVP. In International Conference on Learning Representations, 2017. <https://arxiv.org/abs/1605.08803>
- [3] George Papamakarios, Theo Pavlakou, and Iain Murray. Masked Autoregressive Flow for Density Estimation. In Neural Information Processing Systems, 2017. <https://arxiv.org/abs/1705.07057>

bijective = True

codomain = Real()

constrained_gamma

domain = Real()

event_dim = 0

log_abs_det_jacobian (*x*, *y*)

Calculates the elementwise determinant of the log jacobian, dx/dy

5.3.3 BlockAutoregressive

class BlockAutoregressive (*input_dim*, *hidden_factors*=[8, 8], *activation*='tanh', *residual*=None)

Bases: `pyro.distributions.torch_transform.TransformModule`

An implementation of Block Neural Autoregressive Flow (block-NAF) (De Cao et al., 2019) transformation. Block-NAF uses a similar transformation to deep dense NAF, building the autoregressive NN into the structure of the flow, in a sense.

Together with *TransformedDistribution* this provides a way to create richer variational approximations.

Example usage:

```

>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> naf = BlockAutoregressive(input_dim=10)
>>> pyro.module("my_naf", naf) # doctest: +SKIP
>>> naf_dist = dist.TransformedDistribution(base_dist, [naf])
>>> naf_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])

```

The inverse operation is not implemented. This would require numerical inversion, e.g., using a root finding method - a possibility for a future implementation.

Parameters

- **input_dim** (*int*) – The dimensionality of the input and output variables.
- **hidden_factors** (*list*) – Hidden layer *i* has `hidden_factors[i]` hidden units per input dimension. This corresponds to both *a* and *b* in De Cao et al. (2019). The elements of `hidden_factors` must be integers.
- **activation** (*string*) – Activation function to use. One of ‘ELU’, ‘LeakyReLU’, ‘sigmoid’, or ‘tanh’.
- **residual** (*string*) – Type of residual connections to use. Choices are “None”, “normal” for $y + f(y)$, and “gated” for $\alpha y + (1 - \alpha)y$ for learnable parameter α .

References:

Block Neural Autoregressive Flow [arXiv:1904.04676] Nicola De Cao, Ivan Titov, Wilker Aziz

autoregressive = True

bijective = True

codomain = Real()

domain = Real()

event_dim = 1

log_abs_det_jacobian (*x, y*)

Calculates the elementwise determinant of the log jacobian

5.3.4 DeepELUFlow

class DeepELUFlow (*autoregressive_nn, hidden_units=16*)

Bases: `pyro.distributions.transforms.naf.ELUMixin`, `pyro.distributions.transforms.naf.DeepNAFFlow`

An implementation of deep ELU flow (DSF) Neural Autoregressive Flow (NAF), of the “IAF flavour” that can be used for sampling and scoring samples drawn from it (but not arbitrary ones). This flow is suggested in Huang et al., 2018, section 3.3, but left for future experiments.

Example usage:

```

>>> from pyro.nn import AutoRegressiveNN
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> arn = AutoRegressiveNN(10, [40], param_dims=[16]*3)
>>> naf = DeepELUFlow(arn, hidden_units=16)
>>> pyro.module("my_naf", naf) # doctest: +SKIP
>>> naf_dist = dist.TransformedDistribution(base_dist, [naf])

```

(continues on next page)

(continued from previous page)

```
>>> naf_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])
```

The inverse operation is not implemented. This would require numerical inversion, e.g., using a root finding method - a possibility for a future implementation.

Parameters

- **autoregressive_nn** (*nn.Module*) – an autoregressive neural network whose forward call returns a tuple of three real-valued tensors, whose last dimension is the input dimension, and whose penultimate dimension is equal to `hidden_units`.
- **hidden_units** (*int*) – the number of hidden units to use in the NAF transformation (see Eq (8) in reference)

Reference:

Neural Autoregressive Flows [arXiv:1804.00779] Chin-Wei Huang, David Krueger, Alexandre Lacoste, Aaron Courville

5.3.5 DeepLeakyReLUFlow

class DeepLeakyReLUFlow (*autoregressive_nn*, *hidden_units=16*)

Bases: `pyro.distributions.transforms.naf.LeakyReLUMixin`, `pyro.distributions.transforms.naf.DeepNAFFlow`

An implementation of deep leaky ReLU flow (DSF) Neural Autoregressive Flow (NAF), of the “IAF flavour” that can be used for sampling and scoring samples drawn from it (but not arbitrary ones). This flow is suggested in Huang et al., 2018, section 3.3, but left for future experiments.

Example usage:

```
>>> from pyro.nn import AutoRegressiveNN
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> arn = AutoRegressiveNN(10, [40], param_dims=[16]*3)
>>> naf = DeepLeakyReLUFlow(arn, hidden_units=16)
>>> pyro.module("my_naf", naf) # doctest: +SKIP
>>> naf_dist = dist.TransformedDistribution(base_dist, [naf])
>>> naf_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])
```

The inverse operation is not implemented. This would require numerical inversion, e.g., using a root finding method - a possibility for a future implementation.

Parameters

- **autoregressive_nn** (*nn.Module*) – an autoregressive neural network whose forward call returns a tuple of three real-valued tensors, whose last dimension is the input dimension, and whose penultimate dimension is equal to `hidden_units`.
- **hidden_units** (*int*) – the number of hidden units to use in the NAF transformation (see Eq (8) in reference)

Reference:

Neural Autoregressive Flows [arXiv:1804.00779] Chin-Wei Huang, David Krueger, Alexandre Lacoste, Aaron Courville

5.3.6 DeepSigmoidalFlow

class DeepSigmoidalFlow(*autoregressive_nn*, *hidden_units=16*)

Bases: `pyro.distributions.transforms.naf.SigmoidalMixin`, `pyro.distributions.transforms.naf.DeepNAFFlow`

An implementation of deep sigmoidal flow (DSF) Neural Autoregressive Flow (NAF), of the “IAF flavour” that can be used for sampling and scoring samples drawn from it (but not arbitrary ones).

Example usage:

```
>>> from pyro.nn import AutoRegressiveNN
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> arn = AutoRegressiveNN(10, [40], param_dims=[16]*3)
>>> naf = DeepSigmoidalFlow(arn, hidden_units=16)
>>> pyro.module("my_naf", naf) # doctest: +SKIP
>>> naf_dist = dist.TransformedDistribution(base_dist, [naf])
>>> naf_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
         0.1389, -0.4629,  0.0986])
```

The inverse operation is not implemented. This would require numerical inversion, e.g., using a root finding method - a possibility for a future implementation.

Parameters

- **autoregressive_nn** (*nn.Module*) – an autoregressive neural network whose forward call returns a tuple of three real-valued tensors, whose last dimension is the input dimension, and whose penultimate dimension is equal to *hidden_units*.
- **hidden_units** (*int*) – the number of hidden units to use in the NAF transformation (see Eq (8) in reference)

Reference:

Neural Autoregressive Flows [arXiv:1804.00779] Chin-Wei Huang, David Krueger, Alexandre Lacoste, Aaron Courville

5.3.7 HouseholderFlow

class HouseholderFlow(*input_dim*, *count_transforms=1*)

Bases: `pyro.distributions.torch_transform.TransformModule`

A flow formed from multiple applications of the Householder transformation. A single Householder transformation takes the form,

$$\mathbf{y} = (I - 2 * \frac{\mathbf{u}\mathbf{u}^T}{\|\mathbf{u}\|^2})\mathbf{x}$$

where \mathbf{x} are the inputs, \mathbf{y} are the outputs, and the learnable parameters are $\mathbf{u} \in \mathbb{R}^D$ for input dimension D .

The transformation represents the reflection of \mathbf{x} through the plane passing through the origin with normal \mathbf{u} .

D applications of this transformation are able to transform standard i.i.d. standard Gaussian noise into a Gaussian variable with an arbitrary covariance matrix. With $K < D$ transformations, one is able to approximate a full-rank Gaussian distribution using a linear transformation of rank K .

Together with *TransformedDistribution* this provides a way to create richer variational approximations.

Example usage:

```

>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> flow = HouseholderFlow(10, count_transforms=5)
>>> pyro.module("my_flow", p) # doctest: +SKIP
>>> flow_dist = dist.TransformedDistribution(base_dist, flow)
>>> flow_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])

```

Parameters

- **input_dim** (*int*) – the dimension of the input (and output) variable.
- **count_transforms** (*int*) – number of applications of Householder transformation to apply.

References:

Improving Variational Auto-Encoders using Householder Flow, [arXiv:1611.09630] Tomczak, J. M., & Welling, M.

bijjective = True

codomain = Real()

domain = Real()

event_dim = 1

log_abs_det_jacobian (*x*, *y*)

Calculates the elementwise determinant of the log jacobian. Householder flow is measure preserving, so $\log(|\det J|) = 0$

reset_parameters ()

u ()

volume_preserving = True

5.3.8 InverseAutoRegressiveFlow

class InverseAutoRegressiveFlow (*autoregressive_nn*, *log_scale_min_clip=-5.0*, *log_scale_max_clip=3.0*)

Bases: `pyro.distributions.torch_transform.TransformModule`

An implementation of Inverse Autoregressive Flow, using Eq (10) from Kingma Et Al., 2016,

$$\mathbf{y} = \mu_t + \sigma_t \odot \mathbf{x}$$

where \mathbf{x} are the inputs, \mathbf{y} are the outputs, μ_t, σ_t are calculated from an autoregressive network on \mathbf{x} , and $\sigma_t > 0$.

Together with *TransformedDistribution* this provides a way to create richer variational approximations.

Example usage:

```

>>> from pyro.nn import AutoRegressiveNN
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> iaf = InverseAutoRegressiveFlow(AutoRegressiveNN(10, [40]))
>>> pyro.module("my_iaf", iaf) # doctest: +SKIP
>>> iaf_dist = dist.TransformedDistribution(base_dist, [iaf])
>>> iaf_dist.sample() # doctest: +SKIP

```

(continues on next page)

(continued from previous page)

```
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])
```

The inverse of the Bijector is required when, e.g., scoring the log density of a sample with *TransformedDistribution*. This implementation caches the inverse of the Bijector when its forward operation is called, e.g., when sampling from *TransformedDistribution*. However, if the cached value isn't available, either because it was overwritten during sampling a new value or an arbitrary value is being scored, it will calculate it manually. Note that this is an operation that scales as $O(D)$ where D is the input dimension, and so should be avoided for large dimensional uses. So in general, it is cheap to sample from IAF and score a value that was sampled by IAF, but expensive to score an arbitrary value.

Parameters

- **autoregressive_nn** (*nn.Module*) – an autoregressive neural network whose forward call returns a real-valued mean and logit-scale as a tuple
- **log_scale_min_clip** (*float*) – The minimum value for clipping the log(scale) from the autoregressive NN
- **log_scale_max_clip** (*float*) – The maximum value for clipping the log(scale) from the autoregressive NN

References:

1. Improving Variational Inference with Inverse Autoregressive Flow [arXiv:1606.04934] Diederik P. Kingma, Tim Salimans, Rafal Jozefowicz, Xi Chen, Ilya Sutskever, Max Welling
2. Variational Inference with Normalizing Flows [arXiv:1505.05770] Danilo Jimenez Rezende, Shakir Mohamed
3. MADE: Masked Autoencoder for Distribution Estimation [arXiv:1502.03509] Mathieu Germain, Karol Gregor, Iain Murray, Hugo Larochelle

autoregressive = True

bijection = True

codomain = Real()

domain = Real()

event_dim = 1

log_abs_det_jacobian (*x, y*)

Calculates the elementwise determinant of the log jacobian

5.3.9 InverseAutoRegressiveFlowStable

class InverseAutoRegressiveFlowStable (*autoregressive_nn, sigmoid_bias=2.0*)

Bases: `pyro.distributions.torch_transform.TransformModule`

An implementation of an Inverse Autoregressive Flow, using Eqs (13)/(14) from Kingma Et Al., 2016,

$$\mathbf{y} = \sigma_t \odot \mathbf{x} + (1 - \sigma_t) \odot \mu_t$$

where \mathbf{x} are the inputs, \mathbf{y} are the outputs, μ_t, σ_t are calculated from an autoregressive network on \mathbf{x} , and σ_t is restricted to $(0, 1)$.

This variant of IAF is claimed by the authors to be more numerically stable than one using Eq (10), although in practice it leads to a restriction on the distributions that can be represented, presumably since the input is restricted to rescaling by a number on $(0, 1)$.

Example usage:

```
>>> from pyro.nn import AutoRegressiveNN
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> iaf = InverseAutoregressiveFlowStable(AutoRegressiveNN(10, [40]))
>>> iaf_module = pyro.module("my_iaf", iaf)
>>> iaf_dist = dist.TransformedDistribution(base_dist, [iaf])
>>> iaf_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])
```

See *InverseAutoregressiveFlow* docs for a discussion of the running cost.

Parameters

- **autoregressive_nn** (*nn.Module*) – an autoregressive neural network whose forward call returns a real-valued mean and logit-scale as a tuple
- **sigmoid_bias** (*float*) – bias on the hidden units fed into the sigmoid; default='2.0'

References:

1. Improving Variational Inference with Inverse Autoregressive Flow [arXiv:1606.04934] Diederik P. Kingma, Tim Salimans, Rafal Jozefowicz, Xi Chen, Ilya Sutskever, Max Welling
2. Variational Inference with Normalizing Flows [arXiv:1505.05770] Danilo Jimenez Rezende, Shakir Mohamed
3. MADE: Masked Autoencoder for Distribution Estimation [arXiv:1502.03509] Mathieu Germain, Karol Gregor, Iain Murray, Hugo Larochelle

bijective = True

codomain = `Real()`

domain = `Real()`

event_dim = 1

log_abs_det_jacobian (*x, y*)

Calculates the elementwise determinant of the log jacobian

5.3.10 PermuteTransform

class PermuteTransform (*permutation*)

Bases: `torch.distributions.transforms.Transform`

A bijection that reorders the input dimensions, that is, multiplies the input by a permutation matrix. This is useful in between *InverseAutoregressiveFlow* transforms to increase the flexibility of the resulting distribution and stabilize learning. Whilst not being an autoregressive transform, the log absolute determinant of the Jacobian is easily calculable as 0. Note that reordering the input dimension between two layers of *InverseAutoregressiveFlow* is not equivalent to reordering the dimension inside the MADE networks that those IAFs use; using a *PermuteTransform* results in a distribution with more flexibility.

Example usage:

```
>>> from pyro.nn import AutoRegressiveNN
>>> from pyro.distributions import InverseAutoregressiveFlow, PermuteTransform
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> iaf1 = InverseAutoregressiveFlow(AutoRegressiveNN(10, [40]))
>>> ff = PermuteTransform(torch.randperm(10, dtype=torch.long))
```

(continues on next page)

(continued from previous page)

```
>>> iaf2 = InverseAutoregressiveFlow(AutoRegressiveNN(10, [40]))
>>> iaf_dist = dist.TransformedDistribution(base_dist, [iaf1, ff, iaf2])
>>> iaf_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])
```

Parameters `permutation` (`torch.LongTensor`) – a permutation ordering that is applied to the inputs.

bijection = `True`

codomain = `Real()`

event_dim = `1`

inv_permutation

log_abs_det_jacobian (`x`, `y`)

Calculates the elementwise determinant of the log Jacobian, i.e. $\log(\text{abs}([\text{dy}_0/\text{dx}_0, \dots, \text{dy}_{\{N-1\}}/\text{dx}_{\{N-1\}}]))$. Note that this type of transform is not autoregressive, so the log Jacobian is not the sum of the previous expression. However, it turns out it's always 0 (since the determinant is -1 or +1), and so returning a vector of zeros works.

volume_preserving = `True`

5.3.11 PlanarFlow

class `PlanarFlow` (`input_dim`)

Bases: `pyro.distributions.torch_transform.TransformModule`

A ‘planar’ normalizing flow that uses the transformation

$$\mathbf{y} = \mathbf{x} + \mathbf{u} \tanh(\mathbf{w}^T \mathbf{z} + b)$$

where \mathbf{x} are the inputs, \mathbf{y} are the outputs, and the learnable parameters are $b \in \mathbb{R}$, $\mathbf{u} \in \mathbb{R}^D$, $\mathbf{w} \in \mathbb{R}^D$ for input dimension D . For this to be an invertible transformation, the condition $\mathbf{w}^T \mathbf{u} > -1$ is enforced.

Together with `TransformedDistribution` this provides a way to create richer variational approximations.

Example usage:

```
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> plf = PlanarFlow(10)
>>> pyro.module("my_plf", plf) # doctest: +SKIP
>>> plf_dist = dist.TransformedDistribution(base_dist, [plf])
>>> plf_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])
```

The inverse of this transform does not possess an analytical solution and is left unimplemented. However, the inverse is cached when the forward operation is called during sampling, and so samples drawn using planar flow can be scored.

Parameters `input_dim` (`int`) – the dimension of the input (and output) variable.

References:

Variational Inference with Normalizing Flows [arXiv:1505.05770] Danilo Jimenez Rezende, Shakir Mohamed

```

bijective = True
codomain = Real()
domain = Real()
event_dim = 1
log_abs_det_jacobian(x, y)
    Calculates the elementwise determinant of the log jacobian
reset_parameters()
u_hat()

```

5.3.12 PolynomialFlow

class PolynomialFlow (*autoregressive_nn, input_dim, count_degree, count_sum*)
 Bases: `pyro.distributions.torch_transform.TransformModule`

An autoregressive normalizing flow as described in Jaini et al. (2019) using the element-wise transformation

$$y_n = c_n + \int_0^{x_n} \sum_{k=1}^K \left(\sum_{r=0}^R a_{r,k}^{(n)} u^r \right) du$$

where x_n is the n is the n , $a_{r,k}^{(n)} \in \mathbb{R}$ are learnable parameters that are the output of an autoregressive NN inputting $x_{\prec n} = x_1, x_2, \dots, x_{n-1}$.

Together with *TransformedDistribution* this provides a way to create richer variational approximations.

Example usage:

```

>>> from pyro.nn import AutoRegressiveNN
>>> input_dim = 10
>>> count_degree = 4
>>> count_sum = 3
>>> base_dist = dist.Normal(torch.zeros(input_dim), torch.ones(input_dim))
>>> arn = AutoRegressiveNN(input_dim, [input_dim*10], param_dims=[(count_degree + 1)*count_sum])
>>> flow = PolynomialFlow(arn, input_dim=input_dim, count_degree=count_degree, count_sum=count_sum)
>>> pyro.module("my_flow", flow) # doctest: +SKIP
>>> flow_dist = dist.TransformedDistribution(base_dist, [flow])
>>> flow_dist.sample() # doctest: +SKIP
tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
        0.1389, -0.4629,  0.0986])

```

The inverse of this transform does not possess an analytical solution and is left unimplemented. However, the inverse is cached when the forward operation is called during sampling, and so samples drawn using polynomial flow can be scored.

Parameters

- **autoregressive_nn** (*nn.Module*) – an autoregressive neural network whose forward call returns a tensor of real-valued numbers of size (batch_size, (count_degree+1)*count_sum, input_dim)
- **count_degree** (*int*) – The degree of the polynomial to use for each element-wise transformation.
- **count_sum** (*int*) – The number of polynomials to sum in each element-wise transformation.

References:

Sum-of-squares polynomial flow. [arXiv:1905.02325] Priyank Jaini, Kira A. Shelby, Yaoliang Yu

```
autoregressive = True
bijective = True
codomain = Real()
domain = Real()
event_dim = 1
log_abs_det_jacobian(x, y)
    Calculates the elementwise determinant of the log Jacobian
reset_parameters()
```

5.3.13 RadialFlow

```
class RadialFlow(input_dim)
    Bases: pyro.distributions.torch_transform.TransformModule
```

A ‘radial’ normalizing flow that uses the transformation

$$\mathbf{y} = \mathbf{x} + \beta h(\alpha, r)(\mathbf{x} - \mathbf{x}_0)$$

where \mathbf{x} are the inputs, \mathbf{y} are the outputs, and the learnable parameters are $\alpha \in \mathbb{R}^+$, $\beta \in \mathbb{R}$, $\mathbf{x}_0 \in \mathbb{R}^D$, for input dimension D , $r = \|\mathbf{x} - \mathbf{x}_0\|_2$, $h(\alpha, r) = 1/(\alpha + r)$. For this to be an invertible transformation, the condition $\beta > -\alpha$ is enforced.

Together with *TransformedDistribution* this provides a way to create richer variational approximations.

Example usage:

```
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> flow = RadialFlow(10)
>>> pyro.module("my_flow", flow) # doctest: +SKIP
>>> flow_dist = dist.TransformedDistribution(base_dist, [flow])
>>> flow_dist.sample() # doctest: +SKIP
    tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
            0.1389, -0.4629,  0.0986])
```

The inverse of this transform does not possess an analytical solution and is left unimplemented. However, the inverse is cached when the forward operation is called during sampling, and so samples drawn using radial flow can be scored.

Parameters `input_dim` (*int*) – the dimension of the input (and output) variable.

References:

Variational Inference with Normalizing Flows [arXiv:1505.05770] Danilo Jimenez Rezende, Shakir Mohamed

```
bijective = True
codomain = Real()
domain = Real()
event_dim = 1
log_abs_det_jacobian(x, y)
    Calculates the elementwise determinant of the log jacobian
```


`reset_parameters()`

5.3.14 SylvesterFlow

class `SylvesterFlow(input_dim, count_transforms=1)`

Bases: `pyro.distributions.transforms.householder.HouseholderFlow`

An implementation of Sylvester flow of the Householder variety (Van den Berg Et Al., 2018),

$$\mathbf{y} = \mathbf{x} + QR \tanh(SQ^T \mathbf{x} + \mathbf{b})$$

where \mathbf{x} are the inputs, \mathbf{y} are the outputs, $R, S \sim D \times D$ are upper triangular matrices for input dimension D , $Q \sim D \times D$ is an orthogonal matrix, and $\mathbf{b} \sim D$ is learnable bias term.

Sylvester flow is a generalization of Planar flow. In the Householder type of Sylvester flow, the orthogonality of Q is enforced by representing it as the product of Householder transformations

Together with *TransformedDistribution* it provides a way to create richer variational approximations.

Example usage:

```
>>> base_dist = dist.Normal(torch.zeros(10), torch.ones(10))
>>> hsf = SylvesterFlow(10, count_transforms=4)
>>> pyro.module("my_hsf", hsf) # doctest: +SKIP
>>> hsf_dist = dist.TransformedDistribution(base_dist, [hsf])
>>> hsf_dist.sample() # doctest: +SKIP
      tensor([-0.4071, -0.5030,  0.7924, -0.2366, -0.2387, -0.1417,  0.0868,
              0.1389, -0.4629,  0.0986])
```

The inverse of this transform does not possess an analytical solution and is left unimplemented. However, the inverse is cached when the forward operation is called during sampling, and so samples drawn using Sylvester flow can be scored.

References:

Rianne van den Berg, Leonard Hasenclever, Jakub M. Tomczak, Max Welling. Sylvester Normalizing Flows for Variational Inference. In proceedings of The 34th Conference on Uncertainty in Artificial Intelligence (UAI 2018).

$Q(x)$

$R()$

$S()$

`bijective = True`

`codomain = Real()`

`domain = Real()`

`dtanh_dx(x)`

`event_dim = 1`

`log_abs_det_jacobian(x, y)`

Calculates the elementwise determinant of the log jacobian

`reset_parameters2()`

5.3.15 TransformModule

class TransformModule(*args, **kwargs)

Bases: `torch.distributions.transforms.Transform`, `torch.nn.modules.module.Module`

Transforms with learnable parameters such as normalizing flows should inherit from this class rather than *Transform* so they are also a subclass of *nn.Module* and inherit all the useful methods of that class.

Parameters in Pyro are basically thin wrappers around PyTorch Tensors that carry unique names. As such Parameters are the primary stateful objects in Pyro. Users typically interact with parameters via the Pyro primitive *pyro.param*. Parameters play a central role in stochastic variational inference, where they are used to represent point estimates for the parameters in parameterized families of models and guides.

6.1 ParamStore

class ParamStoreDict

Bases: `object`

Global store for parameters in Pyro. This is basically a key-value store. The typical user interacts with the ParamStore primarily through the primitive *pyro.param*.

See [Intro Part II](#) for further discussion and [SVI Part I](#) for some examples.

Some things to bear in mind when using parameters in Pyro:

- parameters must be assigned unique names
- the *init_tensor* argument to *pyro.param* is only used the first time that a given (named) parameter is registered with Pyro.
- for this reason, a user may need to use the *clear()* method if working in a REPL in order to get the desired behavior. this method can also be invoked with *pyro.clear_param_store()*.
- the internal name of a parameter within a PyTorch *nn.Module* that has been registered with Pyro is prepended with the Pyro name of the module. so nothing prevents the user from having two different modules each of which contains a parameter named *weight*. by contrast, a user can only have one top-level parameter named *weight* (outside of any module).
- parameters can be saved and loaded from disk using *save* and *load*.

clear()

Clear the ParamStore

items()

Iterate over (name, constrained_param) pairs.

keys()

Iterate over param names.

values()

Iterate over constrained parameter values.

setdefault (name, init_constrained_value, constraint=Real())

Retrieve a constrained parameter value from the if it exists, otherwise set the initial value. Note that this is a little fancier than `dict.setdefault()`.

If the parameter already exists, `init_constrained_tensor` will be ignored. To avoid expensive creation of `init_constrained_tensor` you can wrap it in a `lambda` that will only be evaluated if the parameter does not already exist:

```
param_store.get("foo", lambda: (0.001 * torch.randn(1000, 1000)).exp(),
                 constraint=constraints.positive)
```

Parameters

- **name** (*str*) – parameter name
- **init_constrained_value** (*torch.Tensor* or callable returning a *torch.Tensor*) – initial constrained value
- **constraint** (*torch.distributions.constraints.Constraint*) – torch constraint object

Returns constrained parameter value

Return type `torch.Tensor`

named_parameters()

Returns an iterator over (name, unconstrained_value) tuples for each parameter in the ParamStore.

get_all_param_names()

replace_param (param_name, new_param, old_param)

get_param (name, init_tensor=None, constraint=Real(), event_dim=None)

Get parameter from its name. If it does not yet exist in the ParamStore, it will be created and stored. The Pyro primitive `pyro.param` dispatches to this method.

Parameters

- **name** (*str*) – parameter name
- **init_tensor** (*torch.Tensor*) – initial tensor
- **constraint** (*torch.distributions.constraints.Constraint*) – torch constraint
- **event_dim** (*int*) – (ignored)

Returns parameter

Return type `torch.Tensor`

match (name)

Get all parameters that match regex. The parameter must exist.

Parameters `name` (*str*) – regular expression

Returns dict with key param name and value torch Tensor

param_name (*p*)

Get parameter name from parameter

Parameters `p` – parameter

Returns parameter name

get_state ()

Get the ParamStore state.

set_state (*state*)

Set the ParamStore state using state from a previous `get_state()` call

save (*filename*)

Save parameters to disk

Parameters `filename` (*str*) – file name to save to

load (*filename*, *map_location=None*)

Loads parameters from disk

Note: If using `pyro.module()` on parameters loaded from disk, be sure to set the `update_module_params` flag:

```
pyro.get_param_store().load('saved_params.save')
pyro.module('module', nn, update_module_params=True)
```

Parameters

- **filename** (*str*) – file name to load from
- **map_location** (*function, torch.device, string or a dict*) – specifies how to remap storage locations

param_with_module_name (*pyro_name*, *param_name*)

module_from_param_with_module_name (*param_name*)

user_param_name (*param_name*)

The module *pyro.nn* provides implementations of neural network modules that are useful in the context of deep probabilistic programming. None of these modules is really part of the core language.

7.1 AutoRegressiveNN

```
class AutoRegressiveNN(input_dim, hidden_dims, param_dims=[1, 1], permutation=None,
                        skip_connections=False, nonlinearity=ReLU())
Bases: torch.nn.modules.module.Module
```

An implementation of a MADE-like auto-regressive neural network.

Example usage:

```
>>> x = torch.randn(100, 10)
>>> arn = AutoRegressiveNN(10, [50], param_dims=[1])
>>> p = arn(x) # 1 parameters of size (100, 10)
>>> arn = AutoRegressiveNN(10, [50], param_dims=[1, 1])
>>> m, s = arn(x) # 2 parameters of size (100, 10)
>>> arn = AutoRegressiveNN(10, [50], param_dims=[1, 5, 3])
>>> a, b, c = arn(x) # 3 parameters of sizes, (100, 1, 10), (100, 5, 10), (100, 3,
↪ 10)
```

Parameters

- **input_dim** (*int*) – the dimensionality of the input
- **hidden_dims** (*list[int]*) – the dimensionality of the hidden units per layer
- **param_dims** (*list[int]*) – shape the output into parameters of dimension (p_n, input_dim) for p_n in param_dims when p_n > 1 and dimension (input_dim) when p_n == 1. The default is [1, 1], i.e. output two parameters of dimension (input_dim), which is useful for inverse autoregressive flow.

- **permutation** (*torch.LongTensor*) – an optional permutation that is applied to the inputs and controls the order of the autoregressive factorization. in particular for the identity permutation the autoregressive structure is such that the Jacobian is upper triangular. By default this is chosen at random.
- **skip_connections** (*bool*) – Whether to add skip connections from the input to the output.
- **nonlinearity** (*torch.nn.module*) – The nonlinearity to use in the feedforward network such as `torch.nn.ReLU()`. Note that no nonlinearity is applied to the final network output, so the output is an unbounded real number.

Reference:

MADE: Masked Autoencoder for Distribution Estimation [arXiv:1502.03509] Mathieu Germain, Karol Gregor, Iain Murray, Hugo Larochelle

forward (*x*)

The forward method

get_permutation ()

Get the permutation applied to the inputs (by default this is chosen at random)

class MaskedLinear (*in_features, out_features, mask, bias=True*)

Bases: `torch.nn.modules.linear.Linear`

A linear mapping with a given mask on the weights (arbitrary bias)

Parameters

- **in_features** (*int*) – the number of input features
- **out_features** (*int*) – the number of output features
- **mask** (*torch.Tensor*) – the mask to apply to the `in_features` x `out_features` weight matrix
- **bias** (*bool*) – whether or not *MaskedLinear* should include a bias term. defaults to *True*

forward (*_input*)

the forward method that does the masked linear computation and returns the result

create_mask (*input_dim, observed_dim, hidden_dims, permutation, output_dim_multiplier*)

Creates MADE masks for a conditional distribution

Parameters

- **input_dim** (*int*) – the dimensionality of the input variable
- **observed_dim** (*int*) – the dimensionality of the variable that is conditioned on (for conditional densities)
- **hidden_dims** (*list[int]*) – the dimensionality of the hidden layers(s)
- **permutation** (*torch.LongTensor*) – the order of the input variables
- **output_dim_multiplier** (*int*) – tiles the output (e.g. for when a separate mean and scale parameter are desired)

sample_mask_indices (*input_dim, hidden_dim, simple=True*)

Samples the indices assigned to hidden units during the construction of MADE masks

Parameters

- **input_dim** (*int*) – the dimensionality of the input variable

- **hidden_dim** (*int*) – the dimensionality of the hidden layer
- **simple** (*bool*) – True to space fractional indices by rounding to nearest int, false round randomly

The module `pyro.optim` provides support for optimization in Pyro. In particular it provides *PyroOptim*, which is used to wrap PyTorch optimizers and manage optimizers for dynamically generated parameters (see the tutorial [SVI Part I](#) for a discussion). Any custom optimization algorithms are also to be found here.

8.1 Pyro Optimizers

class `PyroOptim`(*optim_constructor*, *optim_args*)

Bases: `object`

A wrapper for `torch.optim.Optimizer` objects that helps with managing dynamically generated parameters.

Parameters

- **`optim_constructor`** – a `torch.optim.Optimizer`
- **`optim_args`** – a dictionary of learning arguments for the optimizer or a callable that returns such dictionaries

`__call__`(*params*, **args*, ***kwargs*)

Parameters *params* (*an iterable of strings*) – a list of parameters

Do an optimization step for each param in *params*. If a given param has never been seen before, initialize an optimizer for it.

`get_state`()

Get state associated with all the optimizers in the form of a dictionary with key-value pairs (parameter name, optim state dicts)

`set_state`(*state_dict*)

Set the state associated with all the optimizers using the state obtained from a previous call to `get_state()`

`save`(*filename*)

Parameters *filename* – file name to save to

Save optimizer state to disk

`load(filename)`

Parameters `filename` – file name to load from

Load optimizer state from disk

AdagradRMSProp(*optim_args*)

Wraps `pyro.optim.adagrad_rmsprop.AdagradRMSProp` with `PyroOptim`.

ClippedAdam(*optim_args*)

Wraps `pyro.optim.clipped_adam.ClippedAdam` with `PyroOptim`.

class PyroLRScheduler(*scheduler_constructor*, *optim_args*)

Bases: `pyro.optim.optim.PyroOptim`

A wrapper for `lr_scheduler` objects that adjusts learning rates for dynamically generated parameters.

Parameters

- **scheduler_constructor** – a `lr_scheduler`
- **optim_args** – a dictionary of learning arguments for the optimizer or a callable that returns such dictionaries. must contain the key ‘optimizer’ with pytorch optimizer value

Example:

```
optimizer = torch.optim.SGD
scheduler = pyro.optim.ExponentialLR({'optimizer': optimizer, 'optim_args': {'lr
→': 0.01}, 'gamma': 0.1})
svi = SVI(model, guide, pyro_scheduler, loss=TraceGraph_ELBO())
for i in range(epochs):
    for minibatch in DataLoader(dataset, batch_size):
        svi.step(minibatch)
    scheduler.step(epoch=i)
```

`__call__`(*params*, **args*, ***kwargs*)

step(**args*, ***kwargs*)

Takes the same arguments as the PyTorch scheduler (optional `epoch` kwarg or `loss` in for `ReduceLROnPlateau`)

class AdagradRMSProp(*params*, *eta=1.0*, *delta=1e-16*, *t=0.1*)

Bases: `torch.optim.optimizer.Optimizer`

Implements a mash-up of the Adagrad algorithm and RMSProp. For the precise update equation see equations 10 and 11 in reference [1].

References: [1] ‘Automatic Differentiation Variational Inference’, Alp Kucukelbir, Dustin Tran, Rajesh Ranganath, Andrew Gelman, David M. Blei URL: <https://arxiv.org/abs/1603.00788> [2] ‘Lecture 6.5 RmsProp: Divide the gradient by a running average of its recent magnitude’, Tieleman, T. and Hinton, G., COURSERA: Neural Networks for Machine Learning. [3] ‘Adaptive subgradient methods for online learning and stochastic optimization’, Duchi, John, Hazan, E and Singer, Y.

Arguments:

Parameters

- **params** – iterable of parameters to optimize or dicts defining parameter groups
- **eta** (*float*) – sets the step size scale (optional; default: 1.0)
- **t** (*float*) – t, optional): momentum parameter (optional; default: 0.1)

- **delta** (*float*) – modulates the exponent that controls how the step size scales (optional: default: 1e-16)

share_memory ()

step (*closure=None*)

Performs a single optimization step.

Parameters **closure** – A (optional) closure that reevaluates the model and returns the loss.

class **ClippedAdam** (*params*, *lr=0.001*, *betas=(0.9, 0.999)*, *eps=1e-08*, *weight_decay=0*, *clip_norm=10.0*, *lrd=1.0*)

Bases: `torch.optim.optimizer.Optimizer`

Parameters

- **params** – iterable of parameters to optimize or dicts defining parameter groups
- **lr** – learning rate (default: 1e-3)
- **betas** (*Tuple*) – coefficients used for computing running averages of gradient and its square (default: (0.9, 0.999))
- **eps** – term added to the denominator to improve numerical stability (default: 1e-8)
- **weight_decay** – weight decay (L2 penalty) (default: 0)
- **clip_norm** – magnitude of norm to which gradients are clipped (default: 10.0)
- **lrd** – rate at which learning rate decays (default: 1.0)

Small modification to the Adam algorithm implemented in `torch.optim.Adam` to include gradient clipping and learning rate decay.

Reference

A Method for Stochastic Optimization, Diederik P. Kingma, Jimmy Ba <https://arxiv.org/abs/1412.6980>

step (*closure=None*)

Parameters **closure** – An optional closure that reevaluates the model and returns the loss.

Performs a single optimization step.

8.2 PyTorch Optimizers

Adamax (*optim_args*)

Wraps `torch.optim.Adamax` with *PyroOptim*.

Adagrad (*optim_args*)

Wraps `torch.optim.Adagrad` with *PyroOptim*.

SGD (*optim_args*)

Wraps `torch.optim.SGD` with *PyroOptim*.

Adam (*optim_args*)

Wraps `torch.optim.Adam` with *PyroOptim*.

Rprop (*optim_args*)

Wraps `torch.optim.Rprop` with *PyroOptim*.

ASGD (*optim_args*)

Wraps `torch.optim.ASGD` with *PyroOptim*.

RMSprop (*optim_args*)Wraps `torch.optim.RMSprop` with *PyroOptim*.**SparseAdam** (*optim_args*)Wraps `torch.optim.SparseAdam` with *PyroOptim*.**Adadelta** (*optim_args*)Wraps `torch.optim.Adadelta` with *PyroOptim*.**MultiStepLR** (*optim_args*)Wraps `torch.optim.MultiStepLR` with *PyroLRScheduler*.**ReduceLROnPlateau** (*optim_args*)Wraps `torch.optim.ReduceLROnPlateau` with *PyroLRScheduler*.**StepLR** (*optim_args*)Wraps `torch.optim.StepLR` with *PyroLRScheduler*.**CosineAnnealingWarmRestarts** (*optim_args*)Wraps `torch.optim.CosineAnnealingWarmRestarts` with *PyroLRScheduler*.**CosineAnnealingLR** (*optim_args*)Wraps `torch.optim.CosineAnnealingLR` with *PyroLRScheduler*.**CyclicLR** (*optim_args*)Wraps `torch.optim.CyclicLR` with *PyroLRScheduler*.**LambdaLR** (*optim_args*)Wraps `torch.optim.LambdaLR` with *PyroLRScheduler*.**ExponentialLR** (*optim_args*)Wraps `torch.optim.ExponentialLR` with *PyroLRScheduler*.

8.3 Higher-Order Optimizers

class MultiOptimizerBases: `object`

Base class of optimizers that make use of higher-order derivatives.

Higher-order optimizers generally use `torch.autograd.grad()` rather than `torch.Tensor.backward()`, and therefore require a different interface from usual Pyro and PyTorch optimizers. In this interface, the `step()` method inputs a loss tensor to be differentiated, and backpropagation is triggered one or more times inside the optimizer.

Derived classes must implement `step()` to compute derivatives and update parameters in-place.

Example:

```
tr = poutine.trace(model).get_trace(*args, **kwargs)
loss = -tr.log_prob_sum()
params = {name: site['value'].unconstrained()
          for name, site in tr.nodes.items()
          if site['type'] == 'param'}
optim.step(loss, params)
```

step (*loss*, *params*)

Performs an in-place optimization step on parameters given a differentiable loss tensor.

Note that this detaches the updated tensors.

Parameters

- **loss** (*torch.Tensor*) – A differentiable tensor to be minimized. Some optimizers require this to be differentiable multiple times.
- **params** (*dict*) – A dictionary mapping param name to unconstrained value as stored in the param store.

get_step (*loss, params*)

Computes an optimization step of parameters given a differentiable `loss` tensor, returning the updated values.

Note that this preserves derivatives on the updated tensors.

Parameters

- **loss** (*torch.Tensor*) – A differentiable tensor to be minimized. Some optimizers require this to be differentiable multiple times.
- **params** (*dict*) – A dictionary mapping param name to unconstrained value as stored in the param store.

Returns A dictionary mapping param name to updated unconstrained value.

Return type *dict*

class PyroMultiOptimizer (*optim*)

Bases: *pyro.optim.multi.MultiOptimizer*

Facade to wrap *PyroOptim* objects in a *MultiOptimizer* interface.

step (*loss, params*)

class TorchMultiOptimizer (*optim_constructor, optim_args*)

Bases: *pyro.optim.multi.PyroMultiOptimizer*

Facade to wrap *Optimizer* objects in a *MultiOptimizer* interface.

class MixedMultiOptimizer (*parts*)

Bases: *pyro.optim.multi.MultiOptimizer*

Container class to combine different *MultiOptimizer* instances for different parameters.

Parameters **parts** (*list*) – A list of (*names*, *optim*) pairs, where each *names* is a list of parameter names, and each *optim* is a *MultiOptimizer* or *PyroOptim* object to be used for the named parameters. Together the *names* should partition up all desired parameters to optimize.

Raises **ValueError** – if any name is optimized by multiple optimizers.

step (*loss, params*)

get_step (*loss, params*)

class Newton (*trust_radii={}*)

Bases: *pyro.optim.multi.MultiOptimizer*

Implementation of *MultiOptimizer* that performs a Newton update on batched low-dimensional variables, optionally regularizing via a per-parameter `trust_radius`. See *newton_step()* for details.

The result of *get_step()* will be differentiable, however the updated values from *step()* will be detached.

Parameters **trust_radii** (*dict*) – a dict mapping parameter name to radius of trust region. Missing names will use unregularized Newton update, equivalent to infinite trust radius.

get_step (*loss, params*)

Poutine (Effect handlers)

Beneath the built-in inference algorithms, Pyro has a library of composable effect handlers for creating new inference algorithms and working with probabilistic programs. Pyro's inference algorithms are all built by applying these handlers to stochastic functions.

9.1 Handlers

Poutine is a library of composable effect handlers for recording and modifying the behavior of Pyro programs. These lower-level ingredients simplify the implementation of new inference algorithms and behavior.

Handlers can be used as higher-order functions, decorators, or context managers to modify the behavior of functions or blocks of code:

For example, consider the following Pyro program:

```
>>> def model(x):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(x, s))
...     return z ** 2
```

We can mark sample sites as observed using `condition`, which returns a callable with the same input and output signatures as `model`:

```
>>> conditioned_model = poutine.condition(model, data={"z": 1.0})
```

We can also use handlers as decorators:

```
>>> @pyro.condition(data={"z": 1.0})
... def model(x):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(x, s))
...     return z ** 2
```

Or as context managers:

```
>>> with pyro.condition(data={"z": 1.0}):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(0., s))
...     y = z ** 2
```

Handlers compose freely:

```
>>> conditioned_model = poutine.condition(model, data={"z": 1.0})
>>> traced_model = poutine.trace(conditioned_model)
```

Many inference algorithms or algorithmic components can be implemented in just a few lines of code:

```
guide_tr = poutine.trace(guide).get_trace(...)
model_tr = poutine.trace(poutine.replay(conditioned_model, trace=guide_tr)).get_
↳ trace(...)
monte_carlo_elbo = model_tr.log_prob_sum() - guide_tr.log_prob_sum()
```

block (*fn=None, hide_fn=None, expose_fn=None, hide=None, expose=None, hide_types=None, expose_types=None*)

This handler selectively hides Pyro primitive sites from the outside world. Default behavior: block everything.

A site is hidden if at least one of the following holds:

0. `hide_fn(msg)` is `True` or `(not expose_fn(msg))` is `True`
1. `msg["name"]` in `hide`
2. `msg["type"]` in `hide_types`
3. `msg["name"]` not in `expose` and `msg["type"]` not in `expose_types`
4. `hide`, `hide_types`, and `expose_types` are all `None`

For example, suppose the stochastic function `fn` has two sample sites “a” and “b”. Then any effect outside of `BlockMessenger(fn, hide=["a"])` will not be applied to site “a” and will only see site “b”:

```
>>> def fn():
...     a = pyro.sample("a", dist.Normal(0., 1.))
...     return pyro.sample("b", dist.Normal(a, 1.))
>>> fn_inner = trace(fn)
>>> fn_outer = trace(block(fn_inner, hide=["a"]))
>>> trace_inner = fn_inner.get_trace()
>>> trace_outer = fn_outer.get_trace()
>>> "a" in trace_inner
True
>>> "a" in trace_outer
False
>>> "b" in trace_inner
True
>>> "b" in trace_outer
True
```

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **hide_fn** – function that takes a site and returns `True` to hide the site or `False/None` to expose it. If specified, all other parameters are ignored. Only specify one of `hide_fn` or `expose_fn`, not both.

- **expose_fn** – function that takes a site and returns True to expose the site or False/None to hide it. If specified, all other parameters are ignored. Only specify one of `hide_fn` or `expose_fn`, not both.
- **hide** – list of site names to hide
- **expose** – list of site names to be exposed while all others hidden
- **hide_types** – list of site types to be hidden
- **expose_types** – list of site types to be exposed while all others hidden

Returns stochastic function decorated with a *BlockMessenger*

broadcast (*fn=None*)

Automatically broadcasts the batch shape of the stochastic function at a sample site when inside a single or nested plate context. The existing *batch_shape* must be broadcastable with the size of the *plate* contexts installed in the *cond_indep_stack*.

Notice how *model_automatic_broadcast* below automates expanding of distribution batch shapes. This makes it easy to modularize a Pyro model as the sub-components are agnostic of the wrapping *plate* contexts.

```
>>> def model_broadcast_by_hand():
...     with IndepMessenger("batch", 100, dim=-2):
...         with IndepMessenger("components", 3, dim=-1):
...             sample = pyro.sample("sample", dist.Bernoulli(torch.ones(3) * 0.5)
...                                     .expand_by(100))
...             assert sample.shape == torch.Size((100, 3))
...     return sample
```

```
>>> @poutine.broadcast
... def model_automatic_broadcast():
...     with IndepMessenger("batch", 100, dim=-2):
...         with IndepMessenger("components", 3, dim=-1):
...             sample = pyro.sample("sample", dist.Bernoulli(torch.tensor(0.5)))
...             assert sample.shape == torch.Size((100, 3))
...     return sample
```

condition (*fn=None, data=None*)

Given a stochastic function with some sample statements and a dictionary of observations at names, change the sample statements at those names into observes with those values.

Consider the following Pyro program:

```
>>> def model(x):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(x, s))
...     return z ** 2
```

To observe a value for site *z*, we can write

```
>>> conditioned_model = condition(model, data={"z": torch.tensor(1.)})
```

This is equivalent to adding *obs=value* as a keyword argument to *pyro.sample("z", ...)* in *model*.

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **data** – a dict or a *Trace*

Returns stochastic function decorated with a *ConditionMessenger*

do (*fn=None, data=None*)

Given a stochastic function with some sample statements and a dictionary of values at names, set the return values of those sites equal to the values and hide them from the rest of the stack as if they were hard-coded to those values by using `block`.

Consider the following Pyro program:

```
>>> def model(x):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(x, s))
...     return z ** 2
```

To intervene with a value for site `z`, we can write

```
>>> intervened_model = do(model, data={"z": torch.tensor(1.)})
```

This is equivalent to replacing `z = pyro.sample("z", ...)` with `z = value`.

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **data** – a dict or a *Trace*

Returns stochastic function decorated with a *BlockMessenger* and *pyro.poutine.condition_messenger.ConditionMessenger*

enum (*fn=None, first_available_dim=None*)

Enumerates in parallel over discrete sample sites marked `infer={"enumerate": "parallel"}`.

Parameters **first_available_dim** (*int*) – The first tensor dimension (counting from the right) that is available for parallel enumeration. This dimension and all dimensions left may be used internally by Pyro. This should be a negative integer.

escape (*fn=None, escape_fn=None*)

Given a callable that contains Pyro primitive calls, evaluate `escape_fn` on each site, and if the result is `True`, raise a *NonlocalExit* exception that stops execution and returns the offending site.

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **escape_fn** – function that takes a partial trace and a site, and returns a boolean value to decide whether to exit at that site

Returns stochastic function decorated with *EscapeMessenger*

infer_config (*fn=None, config_fn=None*)

Given a callable that contains Pyro primitive calls and a callable taking a trace site and returning a dictionary, updates the value of the `infer` kwarg at a sample site to `config_fn(site)`.

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **config_fn** – a callable taking a site and returning an `infer` dict

Returns stochastic function decorated with *InferConfigMessenger*

lift (*fn=None, prior=None*)

Given a stochastic function with `param` calls and a prior distribution, create a stochastic function where all `param` calls are replaced by sampling from `prior`. `Prior` should be a callable or a dict of names to callables.

Consider the following Pyro program:

```
>>> def model(x):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(x, s))
...     return z ** 2
>>> lifted_model = lift(model, prior={"s": dist.Exponential(0.3)})
```

`lift` makes param statements behave like sample statements using the distributions in `prior`. In this example, site `s` will now behave as if it was replaced with `s = pyro.sample("s", dist.Exponential(0.3))`:

```
>>> tr = trace(lifted_model).get_trace(0.0)
>>> tr.nodes["s"]["type"] == "sample"
True
>>> tr2 = trace(lifted_model).get_trace(0.0)
>>> bool((tr2.nodes["s"]["value"] == tr.nodes["s"]["value"]).all())
False
```

Parameters

- **fn** – function whose parameters will be lifted to random values
- **prior** – prior function in the form of a Distribution or a dict of stochastic fns

Returns `fn` decorated with a `LiftMessenger`

markov (*fn=None, history=1, keep=False*)

Markov dependency declaration.

This can be used in a variety of ways: - as a context manager - as a decorator for recursive functions - as an iterator for markov chains

Parameters

- **history** (*int*) – The number of previous contexts visible from the current context. Defaults to 1. If zero, this is similar to `pyro.plate`.
- **keep** (*bool*) – If true, frames are replayable. This is important when branching: if `keep=True`, neighboring branches at the same level can depend on each other; if `keep=False`, neighboring branches are independent (conditioned on their share)

mask (*fn=None, mask=None*)

Given a stochastic function with some batched sample statements and masking tensor, mask out some of the sample statements elementwise.

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **mask** (*torch.ByteTensor*) – a `{0, 1}`-valued masking tensor (1 includes a site, 0 excludes a site)

Returns stochastic function decorated with a `MaskMessenger`

queue (*fn=None, queue=None, max_tries=None, extend_fn=None, escape_fn=None, num_samples=None*)

Used in sequential enumeration over discrete variables.

Given a stochastic function and a queue, return a return value from a complete trace in the queue.

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)

- **queue** – a queue data structure like `multiprocessing.Queue` to hold partial traces
- **max_tries** – maximum number of attempts to compute a single complete trace
- **extend_fn** – function (possibly stochastic) that takes a partial trace and a site, and returns a list of extended traces
- **escape_fn** – function (possibly stochastic) that takes a partial trace and a site, and returns a boolean value to decide whether to exit
- **num_samples** – optional number of extended traces for `extend_fn` to return

Returns stochastic function decorated with poutine logic

replay (*fn=None, trace=None, params=None*)

Given a callable that contains Pyro primitive calls, return a callable that runs the original, reusing the values at sites in `trace` at those sites in the new trace

Consider the following Pyro program:

```
>>> def model(x):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(x, s))
...     return z ** 2
```

`replay` makes sample statements behave as if they had sampled the values at the corresponding sites in the trace:

```
>>> old_trace = trace(model).get_trace(1.0)
>>> replayed_model = replay(model, trace=old_trace)
>>> bool(replayed_model(0.0) == old_trace.nodes["_RETURN"]["value"])
True
```

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **trace** – a `Trace` data structure to replay against
- **params** – dict of names of param sites and constrained values in `fn` to replay against

Returns a stochastic function decorated with a `ReplayMessenger`

scale (*fn=None, scale=None*)

Given a stochastic function with some sample statements and a positive scale factor, scale the score of all sample and observe sites in the function.

Consider the following Pyro program:

```
>>> def model(x):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(x, s), obs=1.0)
...     return z ** 2
```

`scale` multiplicatively scales the log-probabilities of sample sites:

```
>>> scaled_model = scale(model, scale=0.5)
>>> scaled_tr = trace(scaled_model).get_trace(0.0)
>>> unscaled_tr = trace(model).get_trace(0.0)
>>> bool((scaled_tr.log_prob_sum() == 0.5 * unscaled_tr.log_prob_sum()).all())
True
```

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **scale** – a positive scaling factor

Returns stochastic function decorated with a *ScaleMessenger*

trace (*fn=None, graph_type=None, param_only=None*)

Return a handler that records the inputs and outputs of primitive calls and their dependencies.

Consider the following Pyro program:

```
>>> def model(x):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(x, s))
...     return z ** 2
```

We can record its execution using `trace` and use the resulting data structure to compute the log-joint probability of all of the sample sites in the execution or extract all parameters.

```
>>> trace = trace(model).get_trace(0.0)
>>> logp = trace.log_prob_sum()
>>> params = [trace.nodes[name]["value"].unconstrained() for name in trace.param_
↪nodes]
```

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **graph_type** – string that specifies the kind of graph to construct
- **param_only** – if true, only records params and not samples

Returns stochastic function decorated with a *TraceMessenger*

config_enumerate (*guide=None, default='parallel', expand=False, num_samples=None*)

Configures enumeration for all relevant sites in a guide. This is mainly used in conjunction with *TraceEnum_ELBO*.

When configuring for exhaustive enumeration of discrete variables, this configures all sample sites whose distribution satisfies `.has_enumerate_support == True`. When configuring for local parallel Monte Carlo sampling via `default="parallel"`, `num_samples=n`, this configures all sample sites. This does not overwrite existing annotations `infer={"enumerate": ...}`.

This can be used as either a function:

```
guide = config_enumerate(guide)
```

or as a decorator:

```
@config_enumerate
def guide1(*args, **kwargs):
    ...

@config_enumerate(default="sequential", expand=True)
def guide2(*args, **kwargs):
    ...
```

Parameters

- **guide** (*callable*) – a pyro model that will be used as a guide in *SVI*.
- **default** (*str*) – Which enumerate strategy to use, one of “sequential”, “parallel”, or None. Defaults to “parallel”.
- **expand** (*bool*) – Whether to expand enumerated sample values. See *enumerate_support()* for details. This only applies to exhaustive enumeration, where *num_samples*=None. If *num_samples* is not None, then this samples will always be expanded.
- **num_samples** (*int* or *None*) – if not None, use local Monte Carlo sampling rather than exhaustive enumeration. This makes sense for both continuous and discrete distributions.

Returns an annotated guide

Return type callable

9.2 Trace

class **Trace** (*graph_type*='flat')

Bases: *object*

Graph data structure denoting the relationships amongst different pyro primitives in the execution trace.

An execution trace of a Pyro program is a record of every call to `pyro.sample()` and `pyro.param()` in a single execution of that program. Traces are directed graphs whose nodes represent primitive calls or input/output, and whose edges represent conditional dependence relationships between those primitive calls. They are created and populated by `poutine.trace`.

Each node (or site) in a trace contains the name, input and output value of the site, as well as additional metadata added by inference algorithms or user annotation. In the case of `pyro.sample`, the trace also includes the stochastic function at the site, and any observed data added by users.

Consider the following Pyro program:

```
>>> def model(x):
...     s = pyro.param("s", torch.tensor(0.5))
...     z = pyro.sample("z", dist.Normal(x, s))
...     return z ** 2
```

We can record its execution using `pyro.poutine.trace` and use the resulting data structure to compute the log-joint probability of all of the sample sites in the execution or extract all parameters.

```
>>> trace = pyro.poutine.trace(model).get_trace(0.0)
>>> logp = trace.log_prob_sum()
>>> params = [trace.nodes[name]["value"].unconstrained() for name in trace.param_
↳ nodes]
```

We can also inspect or manipulate individual nodes in the trace. `trace.nodes` contains a `collections.OrderedDict` of site names and metadata corresponding to `x`, `s`, `z`, and the return value:

```
>>> list(name for name in trace.nodes.keys()) # doctest: +SKIP
['_INPUT', 's', 'z', '_RETURN']
```

Values of `trace.nodes` are dictionaries of node metadata:


```
>>> trace.nodes["z"] # doctest: +SKIP
{'type': 'sample', 'name': 'z', 'is_observed': False,
 'fn': Normal(), 'value': tensor(0.6480), 'args': (), 'kwargs': {},
 'infer': {}, 'scale': 1.0, 'cond_indep_stack': (),
 'done': True, 'stop': False, 'continuation': None}
```

'infer' is a dictionary of user- or algorithm-specified metadata. 'args' and 'kwargs' are the arguments passed via `pyro.sample` to `fn.__call__` or `fn.log_prob`. 'scale' is used to scale the log-probability of the site when computing the log-joint. 'cond_indep_stack' contains data structures corresponding to `pyro.plate` contexts appearing in the execution. 'done', 'stop', and 'continuation' are only used by Pyro's internals.

Parameters `graph_type` (*string*) – string specifying the kind of trace graph to construct

add_edge (*site1*, *site2*)

add_node (*site_name*, ***kwargs*)

Parameters `site_name` (*string*) – the name of the site to be added

Adds a site to the trace.

Raises an error when attempting to add a duplicate node instead of silently overwriting.

compute_log_prob (*site_filter*=<function <lambda>>)

Compute the site-wise log probabilities of the trace. Each `log_prob` has shape equal to the corresponding `batch_shape`. Each `log_prob_sum` is a scalar. Both computations are memoized.

compute_score_parts ()

Compute the batched local score parts at each site of the trace. Each `log_prob` has shape equal to the corresponding `batch_shape`. Each `log_prob_sum` is a scalar. All computations are memoized.

copy ()

Makes a shallow copy of self with nodes and edges preserved.

edges

format_shapes (*title*='Trace Shapes:', *last_site*=None)

Returns a string showing a table of the shapes of all sites in the trace.

iter_stochastic_nodes ()

Returns an iterator over stochastic nodes in the trace.

log_prob_sum (*site_filter*=<function <lambda>>)

Compute the site-wise log probabilities of the trace. Each `log_prob` has shape equal to the corresponding `batch_shape`. Each `log_prob_sum` is a scalar. The computation of `log_prob_sum` is memoized.

Returns total log probability.

Return type `torch.Tensor`

nonreparam_stochastic_nodes

Returns a list of names of sample sites whose stochastic functions are not reparameterizable primitive distributions

observation_nodes

Returns a list of names of observe sites

pack_tensors (*plate_to_symbol*=None)

Computes packed representations of tensors in the trace. This should be called after `compute_log_prob()` or `compute_score_parts()`.

param_nodes

Returns a list of names of param sites

predecessors (*site_name*)**remove_node** (*site_name*)**reparameterized_nodes**

Returns a list of names of sample sites whose stochastic functions are reparameterizable primitive distributions

stochastic_nodes

Returns a list of names of sample sites

successors (*site_name*)**symbolize_dims** (*plate_to_symbol=None*)

Assign unique symbols to all tensor dimensions.

topological_sort (*reverse=False*)

Return a list of nodes (site names) in topologically sorted order.

Parameters **reverse** (*bool*) – Return the list in reverse order.

Returns list of topologically sorted nodes (site names).

9.3 Messengers

Messenger objects contain the implementations of the effects exposed by handlers. Advanced users may modify the implementations of messengers behind existing handlers or write new messengers that implement new effects and compose correctly with the rest of the library.

9.3.1 Messenger

class Messenger

Bases: `object`

Context manager class that modifies behavior and adds side effects to stochastic functions i.e. callables containing Pyro primitive statements.

This is the base Messenger class. It implements the default behavior for all Pyro primitives, so that the joint distribution induced by a stochastic function `fn` is identical to the joint distribution induced by `Messenger() (fn)`.

Class of transformers for messages passed during inference. Most inference operations are implemented in subclasses of this.

classmethod register (*fn=None, type=None, post=None*)**Parameters**

- **fn** – function implementing operation
- **type** (*str*) – name of the operation (also passed to `effective()`)
- **post** (*bool*) – if `True`, use this operation as postprocess

Dynamically add operations to an effect. Useful for generating wrappers for libraries.

Example:

```
@SomeMessengerClass.register
def some_function(msg):
    ...do_something...
    return msg
```

classmethod unregister (*fn=None, type=None*)

Parameters

- **fn** – function implementing operation
- **type** (*str*) – name of the operation (also passed to *effective()*)

Dynamically remove operations from an effect. Useful for removing wrappers from libraries.

Example:

```
SomeMessengerClass.unregister(some_function, "name")
```

9.3.2 BlockMessenger

class BlockMessenger (*hide_fn=None, expose_fn=None, hide_all=True, expose_all=False, hide=None, expose=None, hide_types=None, expose_types=None*)

Bases: *pyro.poutine.messenger.Messenger*

This Messenger selectively hides Pyro primitive sites from the outside world. Default behavior: block everything. BlockMessenger has a flexible interface that allows users to specify in several different ways which sites should be hidden or exposed.

A site is hidden if at least one of the following holds:

0. `hide_fn(msg)` is `True` or `(not expose_fn(msg))` is `True`
1. `msg["name"]` in `hide`
2. `msg["type"]` in `hide_types`
3. `msg["name"]` not in `expose` and `msg["type"]` not in `expose_types`
4. `hide`, `hide_types`, and `expose_types` are all `None`

For example, suppose the stochastic function `fn` has two sample sites “a” and “b”. Then any poutine outside of `BlockMessenger(fn, hide=["a"])` will not be applied to site “a” and will only see site “b”:

```
>>> def fn():
...     a = pyro.sample("a", dist.Normal(0., 1.))
...     return pyro.sample("b", dist.Normal(a, 1.))
```

```
>>> fn_inner = TraceMessenger()(fn)
>>> fn_outer = TraceMessenger()(BlockMessenger(hide=["a"])(TraceMessenger()(fn)))
>>> trace_inner = fn_inner.get_trace()
>>> trace_outer = fn_outer.get_trace()
>>> "a" in trace_inner
True
>>> "a" in trace_outer
False
>>> "b" in trace_inner
```

(continues on next page)

(continued from previous page)

```
True
>>> "b" in trace_outer
True
```

See the constructor for details.

Parameters

- **hide_fn** – function that takes a site and returns True to hide the site or False/None to expose it. If specified, all other parameters are ignored. Only specify one of `hide_fn` or `expose_fn`, not both.
- **expose_fn** – function that takes a site and returns True to expose the site or False/None to hide it. If specified, all other parameters are ignored. Only specify one of `hide_fn` or `expose_fn`, not both.
- **hide_all** (*bool*) – hide all sites
- **expose_all** (*bool*) – expose all sites normally
- **hide** (*list*) – list of site names to hide, rest will be exposed normally
- **expose** (*list*) – list of site names to expose, rest will be hidden
- **hide_types** (*list*) – list of site types to hide, rest will be exposed normally
- **expose_types** (*list*) – list of site types to expose normally, rest will be hidden

9.3.3 BroadcastMessenger

class BroadcastMessenger

Bases: `pyro.poutine.messenger.Messenger`

BroadcastMessenger automatically broadcasts the batch shape of the stochastic function at a sample site when inside a single or nested plate context. The existing *batch_shape* must be broadcastable with the size of the *plate* contexts installed in the *cond_indep_stack*.

9.3.4 ConditionMessenger

class ConditionMessenger (*data*)

Bases: `pyro.poutine.messenger.Messenger`

Adds values at observe sites to condition on data and override sampling

9.3.5 EscapeMessenger

class EscapeMessenger (*escape_fn*)

Bases: `pyro.poutine.messenger.Messenger`

Messenger that does a nonlocal exit by raising a `util.NonlocalExit` exception

9.3.6 IndepMessenger

class CondIndepStackFrame

Bases: `pyro.poutine.indep_messenger.CondIndepStackFrame`

vectorized

class IndepMessenger (*name=None, size=None, dim=None, device=None*)

Bases: `pyro.poutine.messenger.Messenger`

This messenger keeps track of stack of independence information declared by nested plate contexts. This information is stored in a `cond_indep_stack` at each sample/observe site for consumption by `TraceMessenger`.

Example:

```
x_axis = IndepMessenger('outer', 320, dim=-1)
y_axis = IndepMessenger('inner', 200, dim=-2)
with x_axis:
    x_noise = sample("x_noise", dist.Normal(loc, scale).expand_by([320]))
with y_axis:
    y_noise = sample("y_noise", dist.Normal(loc, scale).expand_by([200, 1]))
with x_axis, y_axis:
    xy_noise = sample("xy_noise", dist.Normal(loc, scale).expand_by([200, 320]))
```

indices

next_context ()

Increments the counter.

9.3.7 LiftMessenger

class LiftMessenger (*prior*)

Bases: `pyro.poutine.messenger.Messenger`

Messenger which “lifts” parameters to random samples. Given a stochastic function with param calls and a prior, creates a stochastic function where all param calls are replaced by sampling from prior.

Prior should be a callable or a dict of names to callables.

9.3.8 ReplayMessenger

class ReplayMessenger (*trace=None, params=None*)

Bases: `pyro.poutine.messenger.Messenger`

Messenger for replaying from an existing execution trace.

9.3.9 ScaleMessenger

class ScaleMessenger (*scale*)

Bases: `pyro.poutine.messenger.Messenger`

This messenger rescales the log probability score.

This is typically used for data subsampling or for stratified sampling of data (e.g. in fraud detection where negatives vastly outnumber positives).

Parameters `scale` (*float* or *torch.Tensor*) – a positive scaling factor

9.3.10 TraceMessenger

class `TraceHandler` (*msngr, fn*)

Bases: `object`

Execution trace poutine.

A `TraceHandler` records the input and output to every Pyro primitive and stores them as a site in a `Trace()`. This should, in theory, be sufficient information for every inference algorithm (along with the implicit computational graph in the `Variables?`)

We can also use this for visualization.

get_trace (**args, **kwargs*)

Returns data structure

Return type `pyro.poutine.Trace`

Helper method for a very common use case. Calls this poutine and returns its trace instead of the function's return value.

trace

class `TraceMessenger` (*graph_type=None, param_only=None*)

Bases: `pyro.poutine.messenger.Messenger`

Execution trace messenger.

A `TraceMessenger` records the input and output to every Pyro primitive and stores them as a site in a `Trace()`. This should, in theory, be sufficient information for every inference algorithm (along with the implicit computational graph in the `Variables?`)

We can also use this for visualization.

get_trace ()

Returns data structure

Return type `pyro.poutine.Trace`

Helper method for a very common use case. Returns a shallow copy of `self.trace`.

identify_dense_edges (*trace*)

Modifies a trace in-place by adding all edges based on the `cond_indep_stack` information stored at each site.

9.4 Runtime

exception `NonlocalExit` (*site, *args, **kwargs*)

Bases: `exceptions.Exception`

Exception for exiting nonlocally from poutine execution.

Used by `poutine.EscapeMessenger` to return site information.

reset_stack ()

Reset the state of the frames remaining in the stack. Necessary for multiple re-executions in `poutine.queue`.

am_i_wrapped ()

Checks whether the current computation is wrapped in a poutine. :returns: bool

apply_stack (*initial_msg*)

Execute the effect stack at a single site according to the following scheme:

1. For each `Messenger` in the stack from bottom to top, execute `Messenger._process_message` with the message; if the message field “stop” is True, stop; otherwise, continue
2. Apply default behavior (`default_process_message`) to finish remaining site execution
3. For each `Messenger` in the stack from top to bottom, execute `_postprocess_message` to update the message and internal messenger state with the site results
4. If the message field “continuation” is not `None`, call it with the message

Parameters `initial_msg` (*dict*) – the starting version of the trace site

Returns `None`

default_process_message (*msg*)

Default method for processing messages in inference.

Parameters `msg` – a message to be processed

Returns `None`

effectful (*fn=None, type=None*)

Parameters

- **fn** – function or callable that performs an effectful computation
- **type** (*str*) – the type label of the operation, e.g. “sample”

Wrapper for calling `apply_stack()` to apply any active effects.

9.5 Utilities

all_escape (*trace, msg*)

Parameters

- **trace** – a partial trace
- **msg** – the message at a Pyro primitive site

Returns boolean decision value

Utility function that checks if a site is not already in a trace.

Used by `EscapeMessenger` to decide whether to do a nonlocal exit at a site. Subroutine for approximately integrating out variables for variance reduction.

discrete_escape (*trace, msg*)

Parameters

- **trace** – a partial trace
- **msg** – the message at a Pyro primitive site

Returns boolean decision value

Utility function that checks if a sample site is discrete and not already in a trace.

Used by `EscapeMessenger` to decide whether to do a nonlocal exit at a site. Subroutine for integrating out discrete variables for variance reduction.

enable_validation (*is_validate*)

enum_extend (*trace, msg, num_samples=None*)

Parameters

- **trace** – a partial trace
- **msg** – the message at a Pyro primitive site
- **num_samples** – maximum number of extended traces to return.

Returns a list of traces, copies of input trace with one extra site

Utility function to copy and extend a trace with sites based on the input site whose values are enumerated from the support of the input site's distribution.

Used for exact inference and integrating out discrete variables.

is_validation_enabled ()

mc_extend (*trace, msg, num_samples=None*)

Parameters

- **trace** – a partial trace
- **msg** – the message at a Pyro primitive site
- **num_samples** – maximum number of extended traces to return.

Returns a list of traces, copies of input trace with one extra site

Utility function to copy and extend a trace with sites based on the input site whose values are sampled from the input site's function.

Used for Monte Carlo marginalization of individual sample sites.

prune_subsample_sites (*trace*)

Copies and removes all subsample sites from a trace.

site_is_subsample (*site*)

Determines whether a trace site originated from a subsample statement inside an *plate*.

The `pyro.ops` module implements tensor utilities that are mostly independent of the rest of Pyro.

10.1 Utilities for HMC

class `DualAveraging` (`prox_center=0`, `t0=10`, `kappa=0.75`, `gamma=0.05`)

Bases: `object`

Dual Averaging is a scheme to solve convex optimization problems. It belongs to a class of subgradient methods which uses subgradients to update parameters (in primal space) of a model. Under some conditions, the averages of generated parameters during the scheme are guaranteed to converge to an optimal value. However, a counter-intuitive aspect of traditional subgradient methods is “new subgradients enter the model with decreasing weights” (see [1]). Dual Averaging scheme solves that phenomenon by updating parameters using weights equally for subgradients (which lie in a dual space), hence we have the name “dual averaging”.

This class implements a dual averaging scheme which is adapted for Markov chain Monte Carlo (MCMC) algorithms. To be more precise, we will replace subgradients by some statistics calculated during an MCMC trajectory. In addition, introducing some free parameters such as `t0` and `kappa` is helpful and still guarantees the convergence of the scheme.

References

[1] *Primal-dual subgradient methods for convex problems*, Yurii Nesterov

[2] *The No-U-turn sampler: adaptively setting path lengths in Hamiltonian Monte Carlo*, Matthew D. Hoffman, Andrew Gelman

Parameters

- **`prox_center`** (`float`) – A “prox-center” parameter introduced in [1] which pulls the primal sequence towards it.
- **`t0`** (`float`) – A free parameter introduced in [2] that stabilizes the initial steps of the scheme.

- **kappa** (*float*) – A free parameter introduced in [2] that controls the weights of steps of the scheme. For a small **kappa**, the scheme will quickly forget states from early steps. This should be a number in (0.5, 1].
- **gamma** (*float*) – A free parameter which controls the speed of the convergence of the scheme.

reset ()

step (*g*)

Updates states of the scheme given a new statistic/subgradient *g*.

Parameters *g* (*float*) – A statistic calculated during an MCMC trajectory or subgradient.

get_state ()

Returns the latest x_t and average of $\{x_i\}_{i=1}^t$ in primal space.

velocity_verlet (*z*, *r*, *potential_fn*, *inverse_mass_matrix*, *step_size*, *num_steps=1*, *z_grads=None*)

Second order symplectic integrator that uses the velocity verlet algorithm.

Parameters

- **z** (*dict*) – dictionary of sample site names and their current values (type *Tensor*).
- **r** (*dict*) – dictionary of sample site names and corresponding momenta (type *Tensor*).
- **potential_fn** (*callable*) – function that returns potential energy given *z* for each sample site. The negative gradient of the function with respect to *z* determines the rate of change of the corresponding sites' momenta *r*.
- **inverse_mass_matrix** (*torch.Tensor*) – a tensor M^{-1} which is used to calculate kinetic energy: $E_{kinetic} = \frac{1}{2}z^T M^{-1}z$. Here *M* can be a 1D tensor (diagonal matrix) or a 2D tensor (dense matrix).
- **step_size** (*float*) – step size for each time step iteration.
- **num_steps** (*int*) – number of discrete time steps over which to integrate.
- **z_grads** (*torch.Tensor*) – optional gradients of potential energy at current *z*.

Return tuple (*z_next*, *r_next*, *z_grads*, *potential_energy*) next position and momenta, together with the potential energy and its gradient w.r.t. *z_next*.

potential_grad (*potential_fn*, *z*)

Gradient of *potential_fn* w.r.t. parameters *z*.

Parameters

- **potential_fn** – python callable that takes in a dictionary of parameters and returns the potential energy.
- **z** (*dict*) – dictionary of parameter values keyed by site name.

Returns tuple of (*z_grads*, *potential_energy*), where *z_grads* is a dictionary with the same keys as *z* containing gradients and *potential_energy* is a torch scalar.

class WelfordCovariance (*diagonal=True*)

Bases: *object*

Implements Welford's online scheme for estimating (co)variance (see [1]). Useful for adapting diagonal and dense mass structures for HMC.

References

[1] *The Art of Computer Programming*, Donald E. Knuth

```

reset ()
update (sample)
get_covariance (regularize=True)

```

10.2 Newton Optimizers

newton_step (*loss*, *x*, *trust_radius=None*)

Performs a Newton update step to minimize loss on a batch of variables, optionally constraining to a trust region [1].

This is especially useful because the final solution of newton iteration is differentiable wrt the inputs, even when all but the final *x* is detached, due to this method's quadratic convergence [2]. *loss* must be twice-differentiable as a function of *x*. If *loss* is 2+d-times differentiable, then the return value of this function is d-times differentiable.

When *loss* is interpreted as a negative log probability density, then the return values *mode*, *cov* of this function can be used to construct a Laplace approximation `MultivariateNormal(mode, cov)`.

Warning: Take care to detach the result of this function when used in an optimization loop. If you forget to detach the result of this function during optimization, then backprop will propagate through the entire iteration process, and worse will compute two extra derivatives for each step.

Example use inside a loop:

```

x = torch.zeros(1000, 2) # arbitrary initial value
for step in range(100):
    x = x.detach() # block gradients through previous steps
    x.requires_grad = True # ensure loss is differentiable wrt x
    loss = my_loss_function(x)
    x = newton_step(loss, x, trust_radius=1.0)
# the final x is still differentiable

```

[1] Yuan, Ya-xiang. ICIAM. Vol. 99. 2000. "A review of trust region algorithms for optimization." <ftp://ftp.cc.ac.cn/pub/yyx/papers/p995.pdf>

[2] Christianson, Bruce. Optimization Methods and Software 3.4 (1994) "Reverse accumulation and attractive fixed points." <http://uhra.herts.ac.uk/bitstream/handle/2299/4338/903839.pdf>

Parameters

- **loss** (*torch.Tensor*) – A scalar function of *x* to be minimized.
- **x** (*torch.Tensor*) – A dependent variable of shape (N, D) where N is the batch size and D is a small number.
- **trust_radius** (*float*) – An optional trust region *trust_radius*. The updated value *mode* of this function will be within *trust_radius* of the input *x*.

Returns A pair (*mode*, *cov*) where *mode* is an updated tensor of the same shape as the original value *x*, and *cov* is an estimate of the covariance DxD matrix with *cov.shape* == *x.shape[: -1] + (D, D)*.

Return type tuple

newton_step_1d (*loss*, *x*, *trust_radius=None*)

Performs a Newton update step to minimize loss on a batch of 1-dimensional variables, optionally regularizing to constrain to a trust region.

See [newton_step\(\)](#) for details.

Parameters

- **loss** (*torch.Tensor*) – A scalar function of *x* to be minimized.
- **x** (*torch.Tensor*) – A dependent variable with rightmost size of 1.
- **trust_radius** (*float*) – An optional trust region *trust_radius*. The updated value mode of this function will be within *trust_radius* of the input *x*.

Returns A pair (*mode*, *cov*) where *mode* is an updated tensor of the same shape as the original value *x*, and *cov* is an estimate of the covariance 1x1 matrix with *cov.shape == x.shape[:-1] + (1,1)*.

Return type *tuple*

newton_step_2d (*loss*, *x*, *trust_radius=None*)

Performs a Newton update step to minimize loss on a batch of 2-dimensional variables, optionally regularizing to constrain to a trust region.

See [newton_step\(\)](#) for details.

Parameters

- **loss** (*torch.Tensor*) – A scalar function of *x* to be minimized.
- **x** (*torch.Tensor*) – A dependent variable with rightmost size of 2.
- **trust_radius** (*float*) – An optional trust region *trust_radius*. The updated value mode of this function will be within *trust_radius* of the input *x*.

Returns A pair (*mode*, *cov*) where *mode* is an updated tensor of the same shape as the original value *x*, and *cov* is an estimate of the covariance 2x2 matrix with *cov.shape == x.shape[:-1] + (2,2)*.

Return type *tuple*

newton_step_3d (*loss*, *x*, *trust_radius=None*)

Performs a Newton update step to minimize loss on a batch of 3-dimensional variables, optionally regularizing to constrain to a trust region.

See [newton_step\(\)](#) for details.

Parameters

- **loss** (*torch.Tensor*) – A scalar function of *x* to be minimized.
- **x** (*torch.Tensor*) – A dependent variable with rightmost size of 2.
- **trust_radius** (*float*) – An optional trust region *trust_radius*. The updated value mode of this function will be within *trust_radius* of the input *x*.

Returns A pair (*mode*, *cov*) where *mode* is an updated tensor of the same shape as the original value *x*, and *cov* is an estimate of the covariance 3x3 matrix with *cov.shape == x.shape[:-1] + (3,3)*.

Return type *tuple*

10.3 Tensor Indexing

vindex(*tensor*, *args*)

Vectorized advanced indexing with broadcasting semantics.

See also the convenience wrapper *Vindex*.

This is useful for writing indexing code that is compatible with batching and enumeration, especially for selecting mixture components with discrete random variables.

For example suppose *x* is a parameter with *x.dim()* == 3 and we wish to generalize the expression *x*[*i*, :, *j*] from integer *i*, *j* to tensors *i*, *j* with batch dims and enum dims (but no event dims). Then we can write the generalize version using *Vindex*

```
xij = Vindex(x)[i, :, j]

batch_shape = broadcast_shape(i.shape, j.shape)
event_shape = (x.size(1),)
assert xij.shape == batch_shape + event_shape
```

To handle the case when *x* may also contain batch dimensions (e.g. if *x* was sampled in a plated context as when using vectorized particles), *vindex()* uses the special convention that *Ellipsis* denotes batch dimensions (hence ... can appear only on the left, never in the middle or in the right). Suppose *x* has event dim 3. Then we can write:

```
old_batch_shape = x.shape[:-3]
old_event_shape = x.shape[-3:]

xij = Vindex(x)[..., i, :, j] # The ... denotes unknown batch shape.

new_batch_shape = broadcast_shape(old_batch_shape, i.shape, j.shape)
new_event_shape = (x.size(1),)
assert xij.shape == new_batch_shape + new_event_shape
```

Note that this special handling of *Ellipsis* differs from the NEP [1].

Formally, this function assumes:

1. Each *arg* is either *Ellipsis*, *slice*(*None*), an integer, or a batched *torch.LongTensor* (i.e. with empty event shape). This function does not support Nontrivial slices or *torch.ByteTensor* masks. *Ellipsis* can only appear on the left as *args*[0].
2. If *args*[0] is not *Ellipsis* then *tensor* is not batched, and its event dim is equal to *len(args)*.
3. If *args*[0] is *Ellipsis* then *tensor* is batched and its event dim is equal to *len(args[1:])*. Dims of *tensor* to the left of the event dims are considered batch dims and will be broadcasted with dims of *tensor* *args*.

Note that if none of the *args* is a tensor with *.dim()* > 0, then this function behaves like standard indexing:

```
if not any(isinstance(a, torch.Tensor) and a.dim() for a in args):
    assert Vindex(x)[args] == x[args]
```

References

- [1] <https://www.numpy.org/neps/nep-0021-advanced-indexing.html> introduces *vindex* as a helper for vectorized indexing. The Pyro implementation is similar to the proposed notation *x.vindex[]* except for slightly different handling of *Ellipsis*.

Parameters

- **tensor** (*torch.Tensor*) – A tensor to be indexed.
- **args** (*tuple*) – An index, as args to `__getitem__`.

Returns A nonstandard interpretation of `tensor[args]`.

Return type `torch.Tensor`

class `Vindex` (*tensor*)

Bases: `object`

Convenience wrapper around `vindex()`.

The following are equivalent:

```
Vindex(x)[..., i, j, :]  
vindex(x, (Ellipsis, i, j, slice(None)))
```

Parameters **tensor** (*torch.Tensor*) – A tensor to be indexed.

Returns An object with a special `__getitem__()` method.

10.4 Tensor Contraction

contract_expression (*equation, *shapes, **kwargs*)

Wrapper around `opt_einsum.contract_expression()` that optionally uses Pyro’s cheap optimizer and optionally caches contraction paths.

Parameters **cache_path** (*bool*) – whether to cache the contraction path. Defaults to True.

contract (*equation, *operands, **kwargs*)

Wrapper around `opt_einsum.contract()` that optionally uses Pyro’s cheap optimizer and optionally caches contraction paths.

Parameters **cache_path** (*bool*) – whether to cache the contraction path. Defaults to True.

einsum (*equation, *operands, **kwargs*)

Generalized plated sum-product algorithm via tensor variable elimination.

This generalizes `contract()` in two ways:

1. Multiple outputs are allowed, and intermediate results can be shared.
2. Inputs and outputs can be plated along symbols given in `plates`; reductions along `plates` are product reductions.

The best way to understand this function is to try the examples below, which show how `einsum()` calls can be implemented as multiple calls to `contract()` (which is generally more expensive).

To illustrate multiple outputs, note that the following are equivalent:

```
z1, z2, z3 = einsum('ab,bc->a,b,c', x, y) # multiple outputs  
  
z1 = contract('ab,bc->a', x, y)  
z2 = contract('ab,bc->b', x, y)  
z3 = contract('ab,bc->c', x, y)
```

To illustrate plated inputs, note that the following are equivalent:

```
assert len(x) == 3 and len(y) == 3
z = einsum('ab,ai,bi->b', w, x, y, plates='i')

z = contract('ab,a,a,a,b,b,b->b', w, *x, *y)
```

When a sum dimension a always appears with a plate dimension i , then a corresponds to a distinct symbol for each slice of a . Thus the following are equivalent:

```
assert len(x) == 3 and len(y) == 3
z = einsum('ai,ai->', x, y, plates='i')

z = contract('a,b,c,a,b,c->', *x, *y)
```

When such a sum dimension appears in the output, it must be accompanied by all of its plate dimensions, e.g. the following are equivalent:

```
assert len(x) == 3 and len(y) == 3
z = einsum('abi,abi->bi', x, y, plates='i')

z0 = contract('ab,ac,ad,ab,ac,ad->b', *x, *y)
z1 = contract('ab,ac,ad,ab,ac,ad->c', *x, *y)
z2 = contract('ab,ac,ad,ab,ac,ad->d', *x, *y)
z = torch.stack([z0, z1, z2])
```

Note that each plate slice through the output is multilinear in all plate slices through all inputs, thus e.g. batch matrix multiply would be implemented *without* plates, so the following are all equivalent:

```
xy = einsum('abc,acd->abd', x, y, plates='')
xy = torch.stack([xa.mm(ya) for xa, ya in zip(x, y)])
xy = torch.bmm(x, y)
```

Among all valid equations, some computations are polynomial in the sizes of the input tensors and other computations are exponential in the sizes of the input tensors. This function raises `NotImplementedError` whenever the computation is exponential.

Parameters

- **equation** (*str*) – An einsum equation, optionally with multiple outputs.
- **operands** (*torch.Tensor*) – A collection of tensors.
- **plates** (*str*) – An optional string of plate symbols.
- **backend** (*str*) – An optional einsum backend, defaults to ‘torch’.
- **cache** (*dict*) – An optional `shared_intermediates()` cache.
- **modulo_total** (*bool*) – Optionally allow einsum to arbitrarily scale each result plate, which can significantly reduce computation. This is safe to set whenever each result plate denotes a nonnormalized probability distribution whose total is not of interest.

Returns a tuple of tensors of requested shape, one entry per output.

Return type *tuple*

Raises

- **ValueError** – if tensor sizes mismatch or an output requests a plated dim without that dim’s plates.
- **NotImplementedError** – if contraction would have cost exponential in the size of any input tensor.

ubersum (*equation*, **operands*, ***kwargs*)
Deprecated, use `einsum()` instead.

10.5 Statistical Utilities

gelman_rubin (*input*, *chain_dim*=0, *sample_dim*=1)
Computes R-hat over chains of samples. It is required that `input.size(sample_dim) >= 2` and `input.size(chain_dim) >= 2`.

Parameters

- **input** (*torch.Tensor*) – the input tensor.
- **chain_dim** (*int*) – the chain dimension.
- **sample_dim** (*int*) – the sample dimension.

Returns **torch.Tensor** R-hat of *input*.

split_gelman_rubin (*input*, *chain_dim*=0, *sample_dim*=1)
Computes R-hat over chains of samples. It is required that `input.size(sample_dim) >= 4`.

Parameters

- **input** (*torch.Tensor*) – the input tensor.
- **chain_dim** (*int*) – the chain dimension.
- **sample_dim** (*int*) – the sample dimension.

Returns **torch.Tensor** split R-hat of *input*.

autocorrelation (*input*, *dim*=0)
Computes the autocorrelation of samples at dimension *dim*.
Reference: https://en.wikipedia.org/wiki/Autocorrelation#Efficient_computation

Parameters

- **input** (*torch.Tensor*) – the input tensor.
- **dim** (*int*) – the dimension to calculate autocorrelation.

Returns **torch.Tensor** autocorrelation of *input*.

autocovariance (*input*, *dim*=0)
Computes the autocovariance of samples at dimension *dim*.

Parameters

- **input** (*torch.Tensor*) – the input tensor.
- **dim** (*int*) – the dimension to calculate autocorrelation.

Returns **torch.Tensor** autocorrelation of *input*.

effective_sample_size (*input*, *chain_dim*=0, *sample_dim*=1)
Computes effective sample size of *input*.

Reference:

[1] *Introduction to Markov Chain Monte Carlo*, Charles J. Geyer

[2] *Stan Reference Manual version 2.18*, Stan Development Team

Parameters

- **input** (*torch.Tensor*) – the input tensor.
- **chain_dim** (*int*) – the chain dimension.
- **sample_dim** (*int*) – the sample dimension.

Returns torch.Tensor effective sample size of input.

resample (*input, num_samples, dim=0, replacement=False*)
 Draws *num_samples* samples from *input* at dimension *dim*.

Parameters

- **input** (*torch.Tensor*) – the input tensor.
- **num_samples** (*int*) – the number of samples to draw from *input*.
- **dim** (*int*) – dimension to draw from *input*.

Returns torch.Tensor samples drawn randomly from *input*.

quantile (*input, probs, dim=0*)
 Computes quantiles of *input* at *probs*. If *probs* is a scalar, the output will be squeezed at *dim*.

Parameters

- **input** (*torch.Tensor*) – the input tensor.
- **probs** (*list*) – quantile positions.
- **dim** (*int*) – dimension to take quantiles from *input*.

Returns torch.Tensor quantiles of *input* at *probs*.

pi (*input, prob, dim=0*)
 Computes percentile interval which assigns equal probability mass to each tail of the interval.

Parameters

- **input** (*torch.Tensor*) – the input tensor.
- **prob** (*float*) – the probability mass of samples within the interval.
- **dim** (*int*) – dimension to calculate percentile interval from *input*.

Returns torch.Tensor quantiles of *input* at *probs*.

hpdi (*input, prob, dim=0*)
 Computes “highest posterior density interval” which is the narrowest interval with probability mass *prob*.

Parameters

- **input** (*torch.Tensor*) – the input tensor.
- **prob** (*float*) – the probability mass of samples within the interval.
- **dim** (*int*) – dimension to calculate percentile interval from *input*.

Returns torch.Tensor quantiles of *input* at *probs*.

waic (*input, log_weights=None, pointwise=False, dim=0*)
 Computes “Widely Applicable/Watanabe-Akaike Information Criterion” (WAIC) and its corresponding effective number of parameters.

Reference:

[1] *WAIC and cross-validation in Stan*, Aki Vehtari, Andrew Gelman

Parameters

- **input** (*torch.Tensor*) – the input tensor, which is log likelihood of a model.
- **log_weights** (*torch.Tensor*) – weights of samples along `dim`.
- **dim** (*int*) – the sample dimension of `input`.

Returns tuple tuple of WAIC and effective number of parameters.

fit_generalized_pareto (*X*)

Given a dataset *X* assumed to be drawn from the Generalized Pareto Distribution, estimate the distributional parameters *k*, *sigma* using a variant of the technique described in reference [1], as described in reference [2].

References [1] ‘A new and efficient estimation method for the generalized Pareto distribution.’ Zhang, J. and Stephens, M.A. (2009). [2] ‘Pareto Smoothed Importance Sampling.’ Aki Vehtari, Andrew Gelman, Jonah Gabry

Parameters *torch.Tensor* – the input data *X*

Returns tuple tuple of floats (*k*, *sigma*) corresponding to the fit parameters

CHAPTER 11

Generic Interface

The `pyro.generic` module provides an interface to dynamically dispatch Pyro code to custom backends.

class GenericModule (*name, default_backend*)

Bases: `object`

Wrapper for a module that can be dynamically routed to a custom backend.

current_backend = {'distributions': 'pyro.distributions', 'infer': 'pyro.infer', 'op

pyro_backend (**args, **kws*)

Context manager to set a custom backend for Pyro models.

Backends can be specified either by name (for standard backends) or by providing a dict mapping module name to backend module name. Standard backends include: `pyro`, `minipyro`, `funsor`, and `numpy`.

Automatic Guide Generation

The `pyro.contrib.autoguide` module provides algorithms to automatically generate guides from simple models, for use in *SVI*. For example to generate a mean field Gaussian guide:

```
def model():
    ...

guide = AutoDiagonalNormal(model)  # a mean field guide
svi = SVI(model, guide, Adam({'lr': 1e-3}), Trace_ELBO())
```

Automatic guides can also be combined using `pyro.poutine.block()` and `AutoGuideList`.

12.1 AutoGuide

class `AutoGuide` (*model*, *prefix*='auto')

Bases: `object`

Base class for automatic guides.

Derived classes must implement the `__call__()` method.

Auto guides can be used individually or combined in an `AutoGuideList` object.

Parameters

- **model** (*callable*) – a pyro model
- **prefix** (*str*) – a prefix that will be prefixed to all param internal sites

__call__ (**args*, ***kwargs*)

A guide with the same **args*, ***kwargs* as the base model.

Returns A dict mapping sample site name to sampled value.

Return type `dict`

median (*args, **kwargs)

Returns the posterior median value of each latent variable.

Returns A dict mapping sample site name to median tensor.

Return type dict

sample_latent (**kwargs)

Samples an encoded latent given the same *args, **kwargs as the base model.

12.2 AutoGuideList

class AutoGuideList (model, prefix='auto')

Bases: `pyro.contrib.autoguide.AutoGuide`

Container class to combine multiple automatic guides.

Example usage:

```
guide = AutoGuideList(my_model)
guide.add(AutoDiagonalNormal(poutine.block(model, hide=["assignment"])))
guide.add(AutoDiscreteParallel(poutine.block(model, expose=["assignment"])))
svi = SVI(model, guide, optim, Trace_ELBO())
```

Parameters

- **model** (*callable*) – a Pyro model
- **prefix** (*str*) – a prefix that will be prefixed to all param internal sites

__call__ (*args, **kwargs)

A composite guide with the same *args, **kwargs as the base model.

Returns A dict mapping sample site name to sampled value.

Return type dict

add (part)

Add an automatic guide for part of the model. The guide should have been created by blocking the model to restrict to a subset of sample sites. No two parts should operate on any one sample site.

Parameters **part** (`AutoGuide` or *callable*) – a partial guide to add

median (*args, **kwargs)

Returns the posterior median value of each latent variable.

Returns A dict mapping sample site name to median tensor.

Return type dict

12.3 AutoCallable

class AutoCallable (model, guide, median=<function <lambda>>)

Bases: `pyro.contrib.autoguide.AutoGuide`

`AutoGuide` wrapper for simple callable guides.

This is used internally for composing autoguides with custom user-defined guides that are simple callables, e.g.:

```
def my_local_guide(*args, **kwargs):
    ...

guide = AutoGuideList(model)
guide.add(AutoDelta(poutine.block(model, expose=['my_global_param'])))
guide.add(my_local_guide) # automatically wrapped in an AutoCallable
```

To specify a median callable, you can instead:

```
def my_local_median(*args, **kwargs)
    ...

guide.add(AutoCallable(model, my_local_guide, my_local_median))
```

For more complex guides that need e.g. access to plates, users should instead subclass `AutoGuide`.

Parameters

- **model** (*callable*) – a Pyro model
- **guide** (*callable*) – a Pyro guide (typically over only part of the model)
- **median** (*callable*) – an optional callable returning a dict mapping sample site name to computed median tensor.

`__call__(*args, **kwargs)`

12.4 AutoDelta

class `AutoDelta(model, prefix='auto', init_loc_fn=<function init_to_median>)`

Bases: `pyro.contrib.autoguide.AutoGuide`

This implementation of `AutoGuide` uses Delta distributions to construct a MAP guide over the entire latent space. The guide does not depend on the model's `*args, **kwargs`.

..note:: This class does MAP inference in constrained space.

Usage:

```
guide = AutoDelta(model)
svi = SVI(model, guide, ...)
```

By default latent variables are initialized using `init_loc_fn()`. To change this default behavior the user should call `pyro.param()` before beginning inference, with "auto_" prefixed to the targeted sample site names e.g. for sample sites named "level" and "concentration", initialize via:

```
pyro.param("auto_level", torch.tensor([-1., 0., 1.]))
pyro.param("auto_concentration", torch.ones(k),
           constraint=constraints.positive)
```

Parameters

- **model** (*callable*) – A Pyro model.
- **init_loc_fn** (*callable*) – A per-site initialization function. See [Initialization](#) section for available functions.

`__call__ (*args, **kwargs)`

An automatic guide with the same `*args, **kwargs` as the base model.

Returns A dict mapping sample site name to sampled value.

Return type `dict`

`median (*args, **kwargs)`

Returns the posterior median value of each latent variable.

Returns A dict mapping sample site name to median tensor.

Return type `dict`

12.5 AutoContinuous

class `AutoContinuous (model, prefix='auto', init_loc_fn=<function init_to_median>)`

Bases: `pyro.contrib.autoguide.AutoGuide`

Base class for implementations of continuous-valued Automatic Differentiation Variational Inference [1].

Each derived class implements its own `get_posterior()` method.

Assumes model structure and latent dimension are fixed, and all latent variables are continuous.

Parameters `model (callable)` – a Pyro model

Reference:

[1] *Automatic Differentiation Variational Inference*, Alp Kucukelbir, Dustin Tran, Rajesh Ranganath, Andrew Gelman, David M. Blei

Parameters

- `model (callable)` – A Pyro model.
- `init_loc_fn (callable)` – A per-site initialization function. See [Initialization](#) section for available functions.

`__call__ (*args, **kwargs)`

An automatic guide with the same `*args, **kwargs` as the base model.

Returns A dict mapping sample site name to sampled value.

Return type `dict`

`get_posterior (*args, **kwargs)`

Returns the posterior distribution.

`median (*args, **kwargs)`

Returns the posterior median value of each latent variable.

Returns A dict mapping sample site name to median tensor.

Return type `dict`

`quantiles (quantiles, *args, **kwargs)`

Returns posterior quantiles each latent variable. Example:

```
print(guide.quantiles([0.05, 0.5, 0.95]))
```


Parameters `quantiles` (*torch.Tensor* or *list*) – A list of requested quantiles between 0 and 1.

Returns A dict mapping sample site name to a list of quantile values.

Return type `dict`

sample_latent (**args, **kwargs*)

Samples an encoded latent given the same **args, **kwargs* as the base model.

12.6 AutoMultivariateNormal

class `AutoMultivariateNormal` (*model, prefix='auto', init_loc_fn=<function init_to_median>*)

Bases: `pyro.contrib.autoguide.AutoContinuous`

This implementation of `AutoContinuous` uses a Cholesky factorization of a Multivariate Normal distribution to construct a guide over the entire latent space. The guide does not depend on the model's **args, **kwargs*.

Usage:

```
guide = AutoMultivariateNormal(model)
svi = SVI(model, guide, ...)
```

By default the mean vector is initialized to zero and the Cholesky factor is initialized to the identity. To change this default behavior the user should call `pyro.param()` before beginning inference, e.g.:

```
latent_dim = 10
pyro.param("auto_loc", torch.randn(latent_dim))
pyro.param("auto_scale_tril", torch.tril(torch.rand(latent_dim)),
          constraint=constraints.lower_cholesky)
```

get_posterior (**args, **kwargs*)

Returns a MultivariateNormal posterior distribution.

12.7 AutoDiagonalNormal

class `AutoDiagonalNormal` (*model, prefix='auto', init_loc_fn=<function init_to_median>*)

Bases: `pyro.contrib.autoguide.AutoContinuous`

This implementation of `AutoContinuous` uses a Normal distribution with a diagonal covariance matrix to construct a guide over the entire latent space. The guide does not depend on the model's **args, **kwargs*.

Usage:

```
guide = AutoDiagonalNormal(model)
svi = SVI(model, guide, ...)
```

By default the mean vector is initialized to zero and the scale is initialized to the identity. To change this default behavior the user should call `pyro.param()` before beginning inference, e.g.:

```
latent_dim = 10
pyro.param("auto_loc", torch.randn(latent_dim))
pyro.param("auto_scale", torch.ones(latent_dim),
          constraint=constraints.positive)
```

get_posterior (*args, **kwargs)
Returns a diagonal Normal posterior distribution.

12.8 AutoLowRankMultivariateNormal

class AutoLowRankMultivariateNormal (model, prefix='auto', init_loc_fn=<function
init_to_median>, rank=1)
Bases: [pyro.contrib.autoguide.AutoContinuous](#)

This implementation of [AutoContinuous](#) uses a low rank plus diagonal Multivariate Normal distribution to construct a guide over the entire latent space. The guide does not depend on the model's *args, **kwargs.

Usage:

```
guide = AutoLowRankMultivariateNormal(model, rank=10)
svi = SVI(model, guide, ...)
```

By default the `cov_diag` is initialized to 1/2 and the `cov_factor` is initialized randomly such that `cov_factor.matmul(cov_factor.t())` is half the identity matrix. To change this default behavior the user should call [pyro.param\(\)](#) before beginning inference, e.g.:

```
latent_dim = 10
pyro.param("auto_loc", torch.randn(latent_dim))
pyro.param("auto_cov_factor", torch.randn(latent_dim, rank))
pyro.param("auto_cov_diag", torch.randn(latent_dim).exp()),
            constraint=constraints.positive)
```

Parameters

- **model** (*callable*) – a generative model
- **rank** (*int*) – the rank of the low-rank part of the covariance matrix
- **init_loc_fn** (*callable*) – A per-site initialization function. See [Initialization](#) section for available functions.
- **prefix** (*str*) – a prefix that will be prefixed to all param internal sites

get_posterior (*args, **kwargs)
Returns a LowRankMultivariateNormal posterior distribution.

12.9 AutoIAFNormal

class AutoIAFNormal (model, hidden_dim=None, prefix='auto', init_loc_fn=<function
init_to_median>)
Bases: [pyro.contrib.autoguide.AutoContinuous](#)

This implementation of [AutoContinuous](#) uses a Diagonal Normal distribution transformed via a InverseAutoregressiveFlow to construct a guide over the entire latent space. The guide does not depend on the model's *args, **kwargs.

Usage:

```
guide = AutoIAFNormal(model, hidden_dim=latent_dim)
svi = SVI(model, guide, ...)
```

Parameters

- **model** (*callable*) – a generative model
- **hidden_dim** (*int*) – number of hidden dimensions in the IAF
- **init_loc_fn** (*callable*) – A per-site initialization function. See [Initialization](#) section for available functions.
- **prefix** (*str*) – a prefix that will be prefixed to all param internal sites

get_posterior (*args, **kwargs)

Returns a diagonal Normal posterior distribution transformed by InverseAutoregressiveFlow.

12.10 AutoLaplaceApproximation

class AutoLaplaceApproximation (*model*, *prefix*='auto', *init_loc_fn*=<function init_to_median>)

Bases: `pyro.contrib.autoguide.AutoContinuous`

Laplace approximation (quadratic approximation) approximates the posterior $\log p(z|x)$ by a multivariate normal distribution in the unconstrained space. Under the hood, it uses Delta distributions to construct a MAP guide over the entire (unconstrained) latent space. Its covariance is given by the inverse of the hessian of $-\log p(x, z)$ at the MAP point of z .

Usage:

```
delta_guide = AutoLaplaceApproximation(model)
svi = SVI(model, delta_guide, ...)
# ...then train the delta_guide...
guide = delta_guide.laplace_approximation()
```

By default the mean vector is initialized to zero. To change this default behavior the user should call `pyro.param()` before beginning inference, e.g.:

```
latent_dim = 10
pyro.param("auto_loc", torch.randn(latent_dim))
```

get_posterior (*args, **kwargs)

Returns a Delta posterior distribution for MAP inference.

laplace_approximation (*args, **kwargs)

Returns a `AutoMultivariateNormal` instance whose posterior's *loc* and *scale_tril* are given by Laplace approximation.

12.11 AutoDiscreteParallel

class AutoDiscreteParallel (*model*, *prefix*='auto')

Bases: `pyro.contrib.autoguide.AutoGuide`

A discrete mean-field guide that learns a latent discrete distribution for each discrete site in the model.

__call__ (*args, **kwargs)

An automatic guide with the same **args*, ***kwargs* as the base model.

Returns A dict mapping sample site name to sampled value.

Return type `dict`

12.12 Initialization

The `pyro.contrib.autoguide` module contains initialization functions for automatic guides.

The standard interface for initialization is a function that inputs a Pyro trace `site` dict and returns an appropriately sized `value` to serve as an initial constrained value for a guide estimate.

`init_to_feasible` (*site*)

Initialize to an arbitrary feasible point, ignoring distribution parameters.

`init_to_sample` (*site*)

Initialize to a random sample from the prior.

`init_to_median` (*site*, *num_samples=15*)

Initialize to the prior median; fallback to a feasible point if median is undefined.

`init_to_mean` (*site*)

Initialize to the prior mean; fallback to median if mean is undefined.

`class InitMessenger` (*init_fn*)

Bases: `pyro.poutine.messenger.Messenger`

Initializes a site by replacing `.sample()` calls with values drawn from an initialization strategy. This is mainly for internal use by autoguide classes.

Parameters `init_fn` (*callable*) – An initialization function.

Automatic Name Generation

The `pyro.contrib.autoname` module provides tools for automatically generating unique, semantically meaningful names for sample sites.

scope (*fn=None, prefix=None, inner=None*)

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)
- **prefix** – a string to prepend to sample names (optional if *fn* is provided)
- **inner** – switch to determine where duplicate name counters appear

Returns *fn* decorated with a `ScopeMessenger`

`scope` prepends a prefix followed by a `/` to the name at a Pyro sample site. It works much like TensorFlow's `name_scope` and `variable_scope`, and can be used as a context manager, a decorator, or a higher-order function.

`scope` is very useful for aligning compositional models with guides or data.

Example:

```
>>> @scope(prefix="a")
... def model():
...     return pyro.sample("x", dist.Bernoulli(0.5))
...
>>> assert "a/x" in poutine.trace(model).get_trace()
```

Example:

```
>>> def model():
...     with scope(prefix="a"):
...         return pyro.sample("x", dist.Bernoulli(0.5))
...
>>> assert "a/x" in poutine.trace(model).get_trace()
```

Scopes compose as expected, with outer scopes appearing before inner scopes in names:

```
>>> @scope(prefix="b")
... def model():
...     with scope(prefix="a"):
...         return pyro.sample("x", dist.Bernoulli(0.5))
...
>>> assert "b/a/x" in poutine.trace(model).get_trace()
```

When used as a decorator or higher-order function, `scope` will use the name of the input function as the prefix if no user-specified prefix is provided.

Example:

```
>>> @scope
... def model():
...     return pyro.sample("x", dist.Bernoulli(0.5))
...
>>> assert "model/x" in poutine.trace(model).get_trace()
```

name_count (*fn=None*)

`name_count` is a very simple autonaming scheme that simply appends a suffix “__” plus a counter to any name that appears multiple times in an execution. Only duplicate instances of a name get a suffix; the first instance is not modified.

Example:

```
>>> @name_count
... def model():
...     for i in range(3):
...         pyro.sample("x", dist.Bernoulli(0.5))
...
>>> assert "x" in poutine.trace(model).get_trace()
>>> assert "x__1" in poutine.trace(model).get_trace()
>>> assert "x__2" in poutine.trace(model).get_trace()
```

`name_count` also composes with `scope()` by adding a suffix to duplicate scope entrances:

Example:

```
>>> @name_count
... def model():
...     for i in range(3):
...         with pyro.contrib.autoname.scope(prefix="a"):
...             pyro.sample("x", dist.Bernoulli(0.5))
...
>>> assert "a/x" in poutine.trace(model).get_trace()
>>> assert "a__1/x" in poutine.trace(model).get_trace()
>>> assert "a__2/x" in poutine.trace(model).get_trace()
```

Example:

```
>>> @name_count
... def model():
...     with pyro.contrib.autoname.scope(prefix="a"):
...         for i in range(3):
...             pyro.sample("x", dist.Bernoulli(0.5))
...
>>> assert "a/x" in poutine.trace(model).get_trace()
```

(continues on next page)

(continued from previous page)

```
>>> assert "a/x__1" in poutine.trace(model).get_trace()
>>> assert "a/x__2" in poutine.trace(model).get_trace()
```

13.1 Named Data Structures

The `pyro.contrib.named` module is a thin syntactic layer on top of Pyro. It allows Pyro models to be written to look like programs with operating on Python data structures like `latent.x.sample(...)`, rather than programs with string-labeled statements like `x = pyro.sample("x", ...)`.

This module provides three container data structures `named.Object`, `named.List`, and `named.Dict`. These data structures are intended to be nested in each other. Together they track the address of each piece of data in each data structure, so that this address can be used as a Pyro site. For example:

```
>>> state = named.Object("state")
>>> print(str(state))
state

>>> z = state.x.y.z # z is just a placeholder.
>>> print(str(z))
state.x.y.z

>>> state.xs = named.List() # Create a contained list.
>>> x0 = state.xs.add()
>>> print(str(x0))
state.xs[0]

>>> state.ys = named.Dict()
>>> foo = state.ys['foo']
>>> print(str(foo))
state.ys['foo']
```

These addresses can now be used inside `sample`, `observe` and `param` statements. These named data structures even provide in-place methods that alias Pyro statements. For example:

```
>>> state = named.Object("state")
>>> loc = state.loc.param_(torch.zeros(1, requires_grad=True))
>>> scale = state.scale.param_(torch.ones(1, requires_grad=True))
>>> z = state.z.sample_(dist.Normal(loc, scale))
>>> obs = state.x.sample_(dist.Normal(loc, scale), obs=z)
```

For deeper examples of how these can be used in model code, see the [Tree Data](#) and [Mixture](#) examples.

Authors: Fritz Obermeyer, Alexander Rush

class `Object` (*name*)

Bases: `object`

Object to hold immutable latent state.

This object can serve either as a container for nested latent state or as a placeholder to be replaced by a tensor via a `named.sample`, `named.observe`, or `named.param` statement. When used as a placeholder, `Object` objects take the place of strings in normal `pyro.sample` statements.

Parameters `name` (*str*) – The name of the object.

Example:

```
state = named.Object("state")
state.x = 0
state.ys = named.List()
state.zs = named.Dict()
state.a.b.c.d.e.f.g = 0 # Creates a chain of named.Objects.
```

Warning: This data structure is write-once: data may be added but may not be mutated or removed. Trying to mutate this data structure may result in silent errors.

sample_(*fn*, **args*, ***kwargs*)

Calls the stochastic function *fn* with additional side-effects depending on *name* and the enclosing context (e.g. an inference algorithm). See [Intro I](#) and [Intro II](#) for a discussion.

Parameters

- **name** – name of sample
- **fn** – distribution class or function
- **obs** – observed datum (optional; should only be used in context of inference) optionally specified in *kwargs*
- **infer** (*dict*) – Optional dictionary of inference parameters specified in *kwargs*. See inference documentation for details.

Returns sample

param_(**args*, ***kwargs*)

Saves the variable as a parameter in the param store. To interact with the param store or write to disk, see [Parameters](#).

Parameters

- **name** (*str*) – name of parameter
- **init_tensor** (*torch.Tensor* or *callable*) – initial tensor or lazy callable that returns a tensor. For large tensors, it may be cheaper to write e.g. `lambda: torch.randn(100000)`, which will only be evaluated on the initial statement.
- **constraint** (*torch.distributions.constraints.Constraint*) – torch constraint, defaults to `constraints.real`.
- **event_dim** (*int*) – (optional) number of rightmost dimensions unrelated to batching. Dimension to the left of this will be considered batch dimensions; if the param statement is inside a subsampled plate, then corresponding batch dimensions of the parameter will be correspondingly subsampled. If unspecified, all dimensions will be considered event dims and no subsampling will be performed.

Returns parameter

Return type `torch.Tensor`

class List (*name=None*)

Bases: `list`

List-like object to hold immutable latent state.

This must either be given a name when constructed:


```
latent = named.List("root")
```

or must be immediately stored in a `named.Object`:

```
latent = named.Object("root")
latent.xs = named.List() # Must be bound to a Object before use.
```

Warning: This data structure is write-once: data may be added but may not be mutated or removed. Trying to mutate this data structure may result in silent errors.

add()

Append one new `named.Object`.

Returns a new latent object at the end

Return type *named.Object*

class Dict (*name=None*)

Bases: `dict`

Dict-like object to hold immutable latent state.

This must either be given a name when constructed:

```
latent = named.Dict("root")
```

or must be immediately stored in a `named.Object`:

```
latent = named.Object("root")
latent.xs = named.Dict() # Must be bound to a Object before use.
```

Warning: This data structure is write-once: data may be added but may not be mutated or removed. Trying to mutate this data structure may result in silent errors.

13.2 Scoping

`pyro.contrib.autoname.scoping` contains the implementation of `pyro.contrib.autoname.scope()`, a tool for automatically appending a semantically meaningful prefix to names of sample sites.

class NameCountMessenger

Bases: `pyro.poutine.messenger.Messenger`

`NameCountMessenger` is the implementation of `pyro.contrib.autoname.name_count()`

class ScopeMessenger (*prefix=None, inner=None*)

Bases: `pyro.poutine.messenger.Messenger`

`ScopeMessenger` is the implementation of `pyro.contrib.autoname.scope()`

scope (*fn=None, prefix=None, inner=None*)

Parameters

- **fn** – a stochastic function (callable containing Pyro primitive calls)

(continued from previous page)

```
>>> assert "x" in poutine.trace(model).get_trace()
>>> assert "x__1" in poutine.trace(model).get_trace()
>>> assert "x__2" in poutine.trace(model).get_trace()
```

`name_count` also composes with `scope()` by adding a suffix to duplicate scope entrances:

Example:

```
>>> @name_count
... def model():
...     for i in range(3):
...         with pyro.contrib.autoname.scope(prefix="a"):
...             pyro.sample("x", dist.Bernoulli(0.5))
...
>>> assert "a/x" in poutine.trace(model).get_trace()
>>> assert "a__1/x" in poutine.trace(model).get_trace()
>>> assert "a__2/x" in poutine.trace(model).get_trace()
```

Example:

```
>>> @name_count
... def model():
...     with pyro.contrib.autoname.scope(prefix="a"):
...         for i in range(3):
...             pyro.sample("x", dist.Bernoulli(0.5))
...
>>> assert "a/x" in poutine.trace(model).get_trace()
>>> assert "a/x__1" in poutine.trace(model).get_trace()
>>> assert "a/x__2" in poutine.trace(model).get_trace()
```


14.1 HiddenLayer

```
class HiddenLayer(X=None, A_mean=None, A_scale=None, non_linearity=<function
    relu>, KL_factor=1.0, A_prior_scale=1.0, include_hidden_bias=True,
    weight_space_sampling=False)
```

This distribution is a basic building block in a Bayesian neural network. It represents a single hidden layer, i.e. an affine transformation applied to a set of inputs X followed by a non-linearity. The uncertainty in the weights is encoded in a Normal variational distribution specified by the parameters A_scale and A_mean . The so-called ‘local reparameterization trick’ is used to reduce variance (see reference below). In effect, this means the weights are never sampled directly; instead one samples in pre-activation space (i.e. before the non-linearity is applied). Since the weights are never directly sampled, when this distribution is used within the context of variational inference, care must be taken to correctly scale the KL divergence term that corresponds to the weight matrix. This term is folded into the *log_prob* method of this distributions.

In effect, this distribution encodes the following generative process:

$A \sim \text{Normal}(A_mean, A_scale)$ output $\sim \text{non_linearity}(AX)$

Parameters

- **X** (*torch.Tensor*) – B x D dimensional mini-batch of inputs
- **A_mean** (*torch.Tensor*) – D x H dimensional specifying weight mean
- **A_scale** (*torch.Tensor*) – D x H dimensional (diagonal covariance matrix) specifying weight uncertainty
- **non_linearity** (*callable*) – a callable that specifies the non-linearity used. defaults to ReLU.
- **KL_factor** (*float*) – scaling factor for the KL divergence. prototypically this is equal to the size of the mini-batch divided by the size of the whole dataset. defaults to 1.0.
- **A_prior** (*float or torch.Tensor*) – the prior over the weights is assumed to be normal with mean zero and scale factor A_prior . default value is 1.0.

- **include_hidden_bias** (*bool*) – controls whether the activations should be augmented with a 1, which can be used to incorporate bias terms. defaults to *True*.
- **weight_space_sampling** (*bool*) – controls whether the local reparameterization trick is used. this is only intended to be used for internal testing. defaults to *False*.

Reference:

Kingma, Diederik P., Tim Salimans, and Max Welling. “Variational dropout and the local reparameterization trick.” *Advances in Neural Information Processing Systems*. 2015.

15.1 EasyGuide

class EasyGuide (*model*)

Bases: `object`

Base class for “easy guides”.

Derived classes should define a `guide()` method. This `guide()` method can combine ordinary guide statements (e.g. `pyro.sample` and `pyro.param`) with the following special statements:

- `group = self.group(...)` selects multiple `pyro.sample` sites in the model. See `Group` for subsequent methods.
- `with self.plate(...):` ... should be used instead of `pyro.plate`.
- `self.map_estimate(...)` uses a `Delta` guide for a single site.

Derived classes may also override the `init()` method to provide custom initialization for models sites.

Parameters `model` (*callable*) – A Pyro model.

guide (**args, **kwargs*)

Guide implementation, to be overridden by user.

init (*site*)

Model initialization method, may be overridden by user.

This should input a site and output a valid sample from that site. The default behavior is to draw a random sample:

```
return site["fn"]()
```

For other possible initialization functions see <http://docs.pyro.ai/en/stable/contrib.autoguide.html#module-pyro.contrib.autoguide.initialization>

__call__ (**args, **kwargs*)

Runs the guide. This is typically used by inference algorithms.

plate (*name*, *size=None*, *subsample_size=None*, *subsample=None*, **args*, ***kwargs*)

A wrapper around `pyro.plate` to allow *EasyGuide* to automatically construct plates. You should use this rather than `pyro.plate` inside your `guide()` implementation.

group (*match='.*'*)

Select a *Group* of model sites for joint guidance.

Parameters *match* (*str*) – A regex string matching names of model sample sites.

Returns A group of model sites.

Return type *Group*

map_estimate (*name*)

Construct a maximum a posteriori (MAP) guide using Delta distributions.

Parameters *name* (*str*) – The name of a model sample site.

Returns A sampled value.

Return type `torch.Tensor`

15.2 easy_guide

easy_guide (*model*)

Convenience decorator to create an *EasyGuide*. The following are equivalent:

```
# Version 1. Decorate a function.
@easy_guide(model)
def guide(self, foo, bar):
    return my_guide(foo, bar)

# Version 2. Create and instantiate a subclass of EasyGuide.
class Guide(EasyGuide):
    def guide(self, foo, bar):
        return my_guide(foo, bar)
guide = Guide(model)
```

Parameters *model* (*callable*) – a Pyro model.

15.3 Group

class Group (*guide*, *sites*)

Bases: `object`

An autoguide helper to match a group of model sites.

Variables

- **event_shape** (`torch.Size`) – The total flattened concatenated shape of all matching sample sites in the model.
- **prototype_sites** (`list`) – A list of all matching sample sites in a prototype trace of the model.

Parameters

- **guide** (`EasyGuide`) – An easyguide instance.

- **sites** (*list*) – A list of model sites.

guide

sample (*guide_name*, *fn*, *infer=None*)

Wrapper around `pyro.sample()` to create a single auxiliary sample site and then unpack to multiple sample sites for model replay.

Parameters

- **guide_name** (*str*) – The name of the auxiliary guide site.
- **fn** (*callable*) – A distribution with shape `self.event_shape`.
- **infer** (*dict*) – Optional inference configuration dict.

Returns A pair (*guide_z*, *model_zs*) where *guide_z* is the single concatenated blob and *model_zs* is a dict mapping site name to constrained model sample.

Return type *tuple*

map_estimate ()

Construct a maximum a posteriori (MAP) guide using Delta distributions.

Returns A dict mapping model site name to sampled value.

Return type *dict*

Generalised Linear Mixed Models

The `pyro.contrib.glmm` module provides models and guides for generalised linear mixed models (GLMM). It also includes the Normal-inverse-gamma family.

To create a classical Bayesian linear model, use:

```
from pyro.contrib.glmm import known_covariance_linear_model

# Note: coef is a p-vector, observation_sd is a scalar
# Here, p=1 (one feature)
model = known_covariance_linear_model(coef_mean=torch.tensor([0.]),
                                      coef_sd=torch.tensor([10.]),
                                      observation_sd=torch.tensor(2.))

# An n x p design tensor
# Here, n=2 (two observations)
design = torch.tensor(torch.tensor([[1.], [-1.]])

model(design)
```

A non-linear link function may be introduced, for instance:

```
from pyro.contrib.glmm import logistic_regression_model

# No observation_sd is needed for logistic models
model = logistic_regression_model(coef_mean=torch.tensor([0.]),
                                 coef_sd=torch.tensor([10.]))
```

Random effects may be incorporated as regular Bayesian regression coefficients. For random effects with a shared covariance matrix, see `pyro.contrib.glmm.lmer_model()`.

See the [Gaussian Processes tutorial](#) for an introduction.

17.1 Models

17.1.1 GPModel

class `GPModel` (*X*, *y*, *kernel*, *mean_function*=None, *jitter*=1e-06)

Bases: `pyro.contrib.gp.parameterized.Parameterized`

Base class for Gaussian Process models.

The core of a Gaussian Process is a covariance function k which governs the similarity between input points. Given k , we can establish a distribution over functions f by a multivariate normal distribution

$$p(f(X)) = \mathcal{N}(0, k(X, X)),$$

where X is any set of input points and $k(X, X)$ is a covariance matrix whose entries are outputs $k(x, z)$ of k over input pairs (x, z) . This distribution is usually denoted by

$$f \sim \mathcal{GP}(0, k).$$

Note: Generally, beside a covariance matrix k , a Gaussian Process can also be specified by a mean function m (which is a zero-value function by default). In that case, its distribution will be

$$p(f(X)) = \mathcal{N}(m(X), k(X, X)).$$

Gaussian Process models are `Parameterized` subclasses. So its parameters can be learned, set priors, or fixed by using corresponding methods from `Parameterized`. A typical way to define a Gaussian Process model is

```
>>> X = torch.tensor([[1., 5, 3], [4, 3, 7]])
>>> y = torch.tensor([2., 1])
>>> kernel = gp.kernels.RBF(input_dim=3)
>>> kernel.set_prior("variance", dist.Uniform(torch.tensor(0.5), torch.tensor(1.
↪5)))
>>> kernel.set_prior("lengthscale", dist.Uniform(torch.tensor(1.0), torch.
↪tensor(3.0)))
>>> gpr = gp.models.GPRegression(X, y, kernel)
```

There are two ways to train a Gaussian Process model:

- Using an MCMC algorithm (in module `pyro.infer.mcmc`) on `model()` to get posterior samples for the Gaussian Process's parameters. For example:

```
>>> hmc_kernel = HMC(gpr.model)
>>> mcmc = MCMC(hmc_kernel, num_samples=10)
>>> mcmc.run()
>>> ls_name = "GPR/RBF/lengthscale"
>>> posterior_ls = mcmc.get_samples()[ls_name]
```

- Using a variational inference on the pair `model()`, `guide()`:

```
>>> optimizer = torch.optim.Adam(gpr.parameters(), lr=0.01)
>>> loss_fn = pyro.infer.TraceMeanField_ELBO().differentiable_loss
>>>
>>> for i in range(1000):
...     svi.step() # doctest: +SKIP
...     optimizer.zero_grad()
...     loss = loss_fn(gpr.model, gpr.guide) # doctest: +SKIP
...     loss.backward() # doctest: +SKIP
...     optimizer.step()
```

To give a prediction on new dataset, simply use `forward()` like any PyTorch `torch.nn.Module`:

```
>>> Xnew = torch.tensor([[2., 3, 1]])
>>> f_loc, f_cov = gpr(Xnew, full_cov=True)
```

Reference:

- [1] *Gaussian Processes for Machine Learning*, Carl E. Rasmussen, Christopher K. I. Williams

Parameters

- **X** (`torch.Tensor`) – A input data for training. Its first dimension is the number of data points.
- **y** (`torch.Tensor`) – An output data for training. Its last dimension is the number of data points.
- **kernel** (`Kernel`) – A Pyro kernel object, which is the covariance function k .
- **mean_function** (`callable`) – An optional mean function m of this Gaussian process. By default, we use zero mean.
- **jitter** (`float`) – A small positive term which is added into the diagonal part of a covariance matrix to help stabilize its Cholesky decomposition.

`model()`

A “model” stochastic function. If `self.y` is `None`, this method returns mean and variance of the Gaussian Process prior.

guide()

A “guide” stochastic function to be used in variational inference methods. It also gives posterior information to the method `forward()` for prediction.

forward(*Xnew*, *full_cov=False*)

Computes the mean and covariance matrix (or variance) of Gaussian Process posterior on a test input data X_{new} :

$$p(f^* | X_{new}, X, y, k, \theta),$$

where θ are parameters of this model.

Note: Model’s parameters θ together with kernel’s parameters have been learned from a training procedure (MCMC or SVI).

Parameters

- **Xnew** (`torch.Tensor`) – A input data for testing. Note that `Xnew.shape[1:]` must be the same as `X.shape[1:]`.
- **full_cov** (`bool`) – A flag to decide if we want to predict full covariance matrix or just variance.

Returns loc and covariance matrix (or variance) of $p(f^*(X_{new}))$

Return type `tuple(torch.Tensor, torch.Tensor)`

set_data(*X*, *y=None*)

Sets data for Gaussian Process models.

Some examples to utilize this method are:

- Batch training on a sparse variational model:

```
>>> Xu = torch.tensor([[1., 0, 2]]) # inducing input
>>> likelihood = gp.likelihoods.Gaussian()
>>> vsgp = gp.models.VariationalSparseGP(X, y, kernel, Xu, likelihood)
>>> optimizer = torch.optim.Adam(vsgp.parameters(), lr=0.01)
>>> loss_fn = pyro.infer.TraceMeanField_ELBO().differentiable_loss
>>> batched_X, batched_y = X.split(split_size=10), y.split(split_size=10)
>>> for Xi, yi in zip(batched_X, batched_y):
...     optimizer.zero_grad()
...     vsgp.set_data(Xi, yi)
...     svi.step() # doctest: +SKIP
...     loss = loss_fn(vsgp.model, vsgp.guide) # doctest: +SKIP
...     loss.backward() # doctest: +SKIP
...     optimizer.step()
```

- Making a two-layer Gaussian Process stochastic function:

```
>>> gpr1 = gp.models.GPRegression(X, None, kernel)
>>> Z, _ = gpr1.model()
>>> gpr2 = gp.models.GPRegression(Z, y, kernel)
>>> def two_layer_model():
...     Z, _ = gpr1.model()
...     gpr2.set_data(Z, y)
...     return gpr2.model()
```

References:

- [1] *Scalable Variational Gaussian Process Classification*, James Hensman, Alexander G. de G. Matthews, Zoubin Ghahramani
- [2] *Deep Gaussian Processes*, Andreas C. Damianou, Neil D. Lawrence

Parameters

- **x** (`torch.Tensor`) – A input data for training. Its first dimension is the number of data points.
- **y** (`torch.Tensor`) – An output data for training. Its last dimension is the number of data points.

17.1.2 GPRegression

class `GPRegression` (*X*, *y*, *kernel*, *noise=None*, *mean_function=None*, *jitter=1e-06*)

Bases: `pyro.contrib.gp.models.model.GPModel`

Gaussian Process Regression model.

The core of a Gaussian Process is a covariance function k which governs the similarity between input points. Given k , we can establish a distribution over functions f by a multivariate normal distribution

$$p(f(X)) = \mathcal{N}(0, k(X, X)),$$

where X is any set of input points and $k(X, X)$ is a covariance matrix whose entries are outputs $k(x, z)$ of k over input pairs (x, z) . This distribution is usually denoted by

$$f \sim \mathcal{GP}(0, k).$$

Note: Generally, beside a covariance matrix k , a Gaussian Process can also be specified by a mean function m (which is a zero-value function by default). In that case, its distribution will be

$$p(f(X)) = \mathcal{N}(m(X), k(X, X)).$$

Given inputs X and their noisy observations y , the Gaussian Process Regression model takes the form

$$\begin{aligned} f &\sim \mathcal{GP}(0, k(X, X)), \\ y &\sim f + \epsilon, \end{aligned}$$

where ϵ is Gaussian noise.

Note: This model has $\mathcal{O}(N^3)$ complexity for training, $\mathcal{O}(N^3)$ complexity for testing. Here, N is the number of train inputs.

Reference:

- [1] *Gaussian Processes for Machine Learning*, Carl E. Rasmussen, Christopher K. I. Williams

Parameters

- **x** (`torch.Tensor`) – A input data for training. Its first dimension is the number of data points.

- **y** (`torch.Tensor`) – An output data for training. Its last dimension is the number of data points.
- **kernel** (`Kernel`) – A Pyro kernel object, which is the covariance function k .
- **noise** (`torch.Tensor`) – Variance of Gaussian noise of this model.
- **mean_function** (`callable`) – An optional mean function m of this Gaussian process. By default, we use zero mean.
- **jitter** (`float`) – A small positive term which is added into the diagonal part of a covariance matrix to help stabilize its Cholesky decomposition.

model (`**kwargs`)

guide (`**kwargs`)

forward (`Xnew`, `full_cov=False`, `noiseless=True`)

Computes the mean and covariance matrix (or variance) of Gaussian Process posterior on a test input data X_{new} :

$$p(f^* | X_{new}, X, y, k, \epsilon) = \mathcal{N}(loc, cov).$$

Note: The noise parameter `noise` (ϵ) together with kernel's parameters have been learned from a training procedure (MCMC or SVI).

Parameters

- **Xnew** (`torch.Tensor`) – A input data for testing. Note that `Xnew.shape[1:]` must be the same as `self.X.shape[1:]`.
- **full_cov** (`bool`) – A flag to decide if we want to predict full covariance matrix or just variance.
- **noiseless** (`bool`) – A flag to decide if we want to include noise in the prediction output or not.

Returns loc and covariance matrix (or variance) of $p(f^*(X_{new}))$

Return type `tuple(torch.Tensor, torch.Tensor)`

iter_sample (`noiseless=True`)

Iteratively constructs a sample from the Gaussian Process posterior.

Recall that at test input points X_{new} , the posterior is multivariate Gaussian distributed with mean and covariance matrix given by `forward()`.

This method samples lazily from this multivariate Gaussian. The advantage of this approach is that later query points can depend upon earlier ones. Particularly useful when the querying is to be done by an optimisation routine.

Note: The noise parameter `noise` (ϵ) together with kernel's parameters have been learned from a training procedure (MCMC or SVI).

Parameters **noiseless** (`bool`) – A flag to decide if we want to add sampling noise to the samples beyond the noise inherent in the GP posterior.

Returns sampler

Return type function

17.1.3 SparseGPRegression

class SparseGPRegression (*X*, *y*, *kernel*, *Xu*, *noise=None*, *mean_function=None*, *approx=None*, *jitter=1e-06*)

Bases: `pyro.contrib.gp.models.model.GPModel`

Sparse Gaussian Process Regression model.

In `GPRegression` model, when the number of input data X is large, the covariance matrix $k(X, X)$ will require a lot of computational steps to compute its inverse (for log likelihood and for prediction). By introducing an additional inducing-input parameter X_u , we can reduce computational cost by approximate $k(X, X)$ by a low-rank Nymström approximation Q (see reference [1]), where

$$Q = k(X, X_u)k(X, X)^{-1}k(X_u, X).$$

Given inputs X , their noisy observations y , and the inducing-input parameters X_u , the model takes the form:

$$\begin{aligned} u &\sim \mathcal{GP}(0, k(X_u, X_u)), \\ f &\sim q(f \mid X, X_u) = \mathbb{E}_{p(u)} q(f \mid X, X_u, u), \\ y &\sim f + \epsilon, \end{aligned}$$

where ϵ is Gaussian noise and the conditional distribution $q(f \mid X, X_u, u)$ is an approximation of

$$p(f \mid X, X_u, u) = \mathcal{N}(m, k(X, X) - Q),$$

whose terms m and $k(X, X) - Q$ is derived from the joint multivariate normal distribution:

$$[f, u] \sim \mathcal{GP}(0, k([X, X_u], [X, X_u])).$$

This class implements three approximation methods:

- Deterministic Training Conditional (DTC):

$$q(f \mid X, X_u, u) = \mathcal{N}(m, 0),$$

which in turns will imply

$$f \sim \mathcal{N}(0, Q).$$

- Fully Independent Training Conditional (FITC):

$$q(f \mid X, X_u, u) = \mathcal{N}(m, \text{diag}(k(X, X) - Q)),$$

which in turns will correct the diagonal part of the approximation in DTC:

$$f \sim \mathcal{N}(0, Q + \text{diag}(k(X, X) - Q)).$$

- Variational Free Energy (VFE), which is similar to DTC but has an additional *trace_term* in the model’s log likelihood. This additional term makes “VFE” equivalent to the variational approach in SparseVariationalGP (see reference [2]).

Note: This model has $\mathcal{O}(NM^2)$ complexity for training, $\mathcal{O}(NM^2)$ complexity for testing. Here, N is the number of train inputs, M is the number of inducing inputs.

References:

[1] *A Unifying View of Sparse Approximate Gaussian Process Regression*, Joaquin Quiñero-Candela, Carl E. Rasmussen

[2] *Variational learning of inducing variables in sparse Gaussian processes*, Michalis Titsias

Parameters

- **x** (*torch.Tensor*) – A input data for training. Its first dimension is the number of data points.
- **y** (*torch.Tensor*) – An output data for training. Its last dimension is the number of data points.
- **kernel** (*Kernel*) – A Pyro kernel object, which is the covariance function k .
- **Xu** (*torch.Tensor*) – Initial values for inducing points, which are parameters of our model.
- **noise** (*torch.Tensor*) – Variance of Gaussian noise of this model.
- **mean_function** (*callable*) – An optional mean function m of this Gaussian process. By default, we use zero mean.
- **approx** (*str*) – One of approximation methods: “DTC”, “FITC”, and “VFE” (default).
- **jitter** (*float*) – A small positive term which is added into the diagonal part of a covariance matrix to help stabilize its Cholesky decomposition.
- **name** (*str*) – Name of this model.

model (***kwargs*)

guide (***kwargs*)

forward (X_{new} , *full_cov=False*, *noiseless=True*)

Computes the mean and covariance matrix (or variance) of Gaussian Process posterior on a test input data X_{new} :

$$p(f^* \mid X_{new}, X, y, k, X_u, \epsilon) = \mathcal{N}(loc, cov).$$

Note: The noise parameter `noise` (ϵ), the inducing-point parameter `Xu`, together with `kernel`’s parameters have been learned from a training procedure (MCMC or SVI).

Parameters

- **Xnew** (*torch.Tensor*) – A input data for testing. Note that `Xnew.shape[1:]` must be the same as `self.X.shape[1:]`.
- **full_cov** (*bool*) – A flag to decide if we want to predict full covariance matrix or just variance.

- **noiseless** (*bool*) – A flag to decide if we want to include noise in the prediction output or not.

Returns loc and covariance matrix (or variance) of $p(f^*(X_{new}))$

Return type `tuple(torch.Tensor, torch.Tensor)`

17.1.4 VariationalGP

class VariationalGP (*X*, *y*, *kernel*, *likelihood*, *mean_function=None*, *latent_shape=None*, *whiten=False*, *jitter=1e-06*)

Bases: `pyro.contrib.gp.models.model.GPModel`

Variational Gaussian Process model.

This model deals with both Gaussian and non-Gaussian likelihoods. Given inputs X and their noisy observations y , the model takes the form

$$\begin{aligned} f &\sim \mathcal{GP}(0, k(X, X)), \\ y &\sim p(y) = p(y | f)p(f), \end{aligned}$$

where $p(y | f)$ is the likelihood.

We will use a variational approach in this model by approximating $q(f)$ to the posterior $p(f | y)$. Precisely, $q(f)$ will be a multivariate normal distribution with two parameters `f_loc` and `f_scale_tril`, which will be learned during a variational inference process.

Note: This model can be seen as a special version of `SparseVariationalGP` model with $X_u = X$.

Note: This model has $\mathcal{O}(N^3)$ complexity for training, $\mathcal{O}(N^3)$ complexity for testing. Here, N is the number of train inputs. Size of variational parameters is $\mathcal{O}(N^2)$.

Parameters

- **X** (*torch.Tensor*) – A input data for training. Its first dimension is the number of data points.
- **y** (*torch.Tensor*) – An output data for training. Its last dimension is the number of data points.
- **kernel** (*Kernel*) – A Pyro kernel object, which is the covariance function k .
- **Likelihood likelihood** (*likelihood*) – A likelihood object.
- **mean_function** (*callable*) – An optional mean function m of this Gaussian process. By default, we use zero mean.
- **latent_shape** (*torch.Size*) – Shape for latent processes (*batch_shape* of $q(f)$). By default, it equals to output batch shape `y.shape[: -1]`. For the multi-class classification problems, `latent_shape[-1]` should correspond to the number of classes.
- **whiten** (*bool*) – A flag to tell if variational parameters `f_loc` and `f_scale_tril` are transformed by the inverse of `Lff`, where `Lff` is the lower triangular decomposition of $\text{kernel}(X, X)$. Enable this flag will help optimization.
- **jitter** (*float*) – A small positive term which is added into the diagonal part of a covariance matrix to help stabilize its Cholesky decomposition.

model (***kwargs*)

guide (***kwargs*)

forward (X_{new} , $full_cov=False$)

Computes the mean and covariance matrix (or variance) of Gaussian Process posterior on a test input data X_{new} :

$$p(f^* \mid X_{new}, X, y, k, f_{loc}, f_{scale_tril}) = \mathcal{N}(loc, cov).$$

Note: Variational parameters `f_loc`, `f_scale_tril`, together with kernel's parameters have been learned from a training procedure (MCMC or SVI).

Parameters

- **Xnew** (*torch.Tensor*) – A input data for testing. Note that `Xnew.shape[1:]` must be the same as `self.X.shape[1:]`.
- **full_cov** (*bool*) – A flag to decide if we want to predict full covariance matrix or just variance.

Returns loc and covariance matrix (or variance) of $p(f^*(X_{new}))$

Return type `tuple(torch.Tensor, torch.Tensor)`

17.1.5 VariationalSparseGP

class VariationalSparseGP ($X, y, kernel, Xu, likelihood, mean_function=None, latent_shape=None, num_data=None, whiten=False, jitter=1e-06$)

Bases: `pyro.contrib.gp.models.model.GPModel`

Variational Sparse Gaussian Process model.

In *VariationalGP* model, when the number of input data X is large, the covariance matrix $k(X, X)$ will require a lot of computational steps to compute its inverse (for log likelihood and for prediction). This model introduces an additional inducing-input parameter X_u to solve that problem. Given inputs X , their noisy observations y , and the inducing-input parameters X_u , the model takes the form:

$$\begin{aligned} [f, u] &\sim \mathcal{GP}(0, k([X, X_u], [X, X_u])), \\ y &\sim p(y) = p(y \mid f)p(f), \end{aligned}$$

where $p(y \mid f)$ is the likelihood.

We will use a variational approach in this model by approximating $q(f, u)$ to the posterior $p(f, u \mid y)$. Precisely, $q(f) = p(f \mid u)q(u)$, where $q(u)$ is a multivariate normal distribution with two parameters `u_loc` and `u_scale_tril`, which will be learned during a variational inference process.

Note: This model can be learned using MCMC method as in reference [2]. See also `GPModel`.

Note: This model has $\mathcal{O}(NM^2)$ complexity for training, $\mathcal{O}(M^3)$ complexity for testing. Here, N is the number of train inputs, M is the number of inducing inputs. Size of variational parameters is $\mathcal{O}(M^2)$.

References:

[1] *Scalable variational Gaussian process classification*, James Hensman, Alexander G. de G. Matthews, Zoubin Ghahramani

[2] *MCMC for Variationally Sparse Gaussian Processes*, James Hensman, Alexander G. de G. Matthews, Maurizio Filippone, Zoubin Ghahramani

Parameters

- **x** (*torch.Tensor*) – A input data for training. Its first dimension is the number of data points.
- **y** (*torch.Tensor*) – An output data for training. Its last dimension is the number of data points.
- **kernel** (*Kernel*) – A Pyro kernel object, which is the covariance function k .
- **Xu** (*torch.Tensor*) – Initial values for inducing points, which are parameters of our model.
- **Likelihood likelihood** (*likelihood*) – A likelihood object.
- **mean_function** (*callable*) – An optional mean function m of this Gaussian process. By default, we use zero mean.
- **latent_shape** (*torch.Size*) – Shape for latent processes (*batch_shape* of $q(u)$). By default, it equals to output batch shape $y.shape[-1]$. For the multi-class classification problems, `latent_shape[-1]` should correspond to the number of classes.
- **num_data** (*int*) – The size of full training dataset. It is useful for training this model with mini-batch.
- **whiten** (*bool*) – A flag to tell if variational parameters `u_loc` and `u_scale_tril` are transformed by the inverse of `Luu`, where `Luu` is the lower triangular decomposition of $kernel(X_u, X_u)$. Enable this flag will help optimization.
- **jitter** (*float*) – A small positive term which is added into the diagonal part of a covariance matrix to help stabilize its Cholesky decomposition.

model (***kwargs*)

guide (***kwargs*)

forward (X_{new} , *full_cov=False*)

Computes the mean and covariance matrix (or variance) of Gaussian Process posterior on a test input data X_{new} :

$$p(f^* | X_{new}, X, y, k, X_u, u_{loc}, u_{scale_tril}) = \mathcal{N}(loc, cov).$$

Note: Variational parameters `u_loc`, `u_scale_tril`, the inducing-point parameter `Xu`, together with kernel's parameters have been learned from a training procedure (MCMC or SVI).

Parameters

- **Xnew** (*torch.Tensor*) – A input data for testing. Note that `Xnew.shape[1:]` must be the same as `self.X.shape[1:]`.
- **full_cov** (*bool*) – A flag to decide if we want to predict full covariance matrix or just variance.

Returns loc and covariance matrix (or variance) of $p(f^*(X_{new}))$

Return type `tuple(torch.Tensor, torch.Tensor)`

17.1.6 GPLVM

class GPLVM (*base_model*)

Bases: *pyro.contrib.gp.parameterized.Parameterized*

Gaussian Process Latent Variable Model (GPLVM) model.

GPLVM is a Gaussian Process model with its train input data is a latent variable. This model is useful for dimensional reduction of high dimensional data. Assume the mapping from low dimensional latent variable to is a Gaussian Process instance. Then the high dimensional data will play the role of train output y and our target is to learn latent inputs which best explain y . For the purpose of dimensional reduction, latent inputs should have lower dimensions than y .

We follows reference [1] to put a unit Gaussian prior to the input and approximate its posterior by a multivariate normal distribution with two variational parameters: X_{loc} and $X_{\text{scale_tril}}$.

For example, we can do dimensional reduction on Iris dataset as follows:

```
>>> # With y as the 2D Iris data of shape 150x4 and we want to reduce its_
↳dimension
>>> # to a tensor X of shape 150x2, we will use GPLVM.
```

```
>>> # First, define the initial values for X parameter:
>>> X_init = torch.zeros(150, 2)
>>> # Then, define a Gaussian Process model with input X_init and output_
↳y:
>>> kernel = gp.kernels.RBF(input_dim=2, lengthscale=torch.ones(2))
>>> Xu = torch.zeros(20, 2) # initial inducing inputs of sparse model
>>> gpmodule = gp.models.SparseGPRegression(X_init, y, kernel, Xu)
>>> # Finally, wrap gpmodule by GPLVM, optimize, and get the "learned"_
↳mean of X:
>>> gplvm = gp.models.GPLVM(gpmodule)
>>> gp.util.train(gplvm) # doctest: +SKIP
>>> X = gplvm.X
```

Reference:

[1] Bayesian Gaussian Process Latent Variable Model Michalis K. Titsias, Neil D. Lawrence

Parameters **base_model** (*GPModel*) – A Pyro Gaussian Process model object. Note that `base_model.X` will be the initial value for the variational parameter X_{loc} .

model ()

guide ()

forward (**kwargs)

Forward method has the same signal as its `base_model`. Note that the train input data of `base_model` is sampled from GPLVM.

17.2 Kernels

17.2.1 Kernel

class Kernel (*input_dim, active_dims=None*)

Bases: *pyro.contrib.gp.parameterized.Parameterized*

Base class for kernels used in this Gaussian Process module.

Every inherited class should implement a `forward()` pass which takes inputs X , Z and returns their covariance matrix.

To construct a new kernel from the old ones, we can use methods `add()`, `mul()`, `exp()`, `warp()`, `vertical_scale()`.

References:

[1] *Gaussian Processes for Machine Learning*, Carl E. Rasmussen, Christopher K. I. Williams

Parameters

- **input_dim** (*int*) – Number of feature dimensions of inputs.
- **variance** (*torch.Tensor*) – Variance parameter of this kernel.
- **active_dims** (*list*) – List of feature dimensions of the input which the kernel acts on.

forward (X , $Z=None$, $diag=False$)

Calculates covariance matrix of inputs on active dimensionals.

Parameters

- **x** (*torch.Tensor*) – A 2D tensor with shape $N \times \text{input_dim}$.
- **z** (*torch.Tensor*) – An (optional) 2D tensor with shape $M \times \text{input_dim}$.
- **diag** (*bool*) – A flag to decide if we want to return full covariance matrix or just its diagonal part.

Returns covariance matrix of X and Z with shape $N \times M$

Return type *torch.Tensor*

17.2.2 Brownian

class Brownian (*input_dim*, *variance=None*, *active_dims=None*)

Bases: `pyro.contrib.gp.kernels.kernel.Kernel`

This kernel corresponds to a two-sided Brownian motion (Wiener process):

$$k(x, z) = \begin{cases} \sigma^2 \min(|x|, |z|), & \text{if } x \cdot z \geq 0 \\ 0, & \text{otherwise.} \end{cases}$$

Note that the input dimension of this kernel must be 1.

Reference:

[1] *Theory and Statistical Applications of Stochastic Processes*, Yuliya Mishura, Georgiy Shevchenko

forward (X , $Z=None$, $diag=False$)

17.2.3 Combination

class Combination (*kern0*, *kern1*)

Bases: `pyro.contrib.gp.kernels.kernel.Kernel`

Base class for kernels derived from a combination of kernels.

Parameters

- **kern0** (*Kernel*) – First kernel to combine.
- **kern1** (*Kernel* or *numbers.Number*) – Second kernel to combine.

17.2.4 Constant

class Constant (*input_dim*, *variance=None*, *active_dims=None*)

Bases: `pyro.contrib.gp.kernels.kernel.Kernel`

Implementation of Constant kernel:

$$k(x, z) = \sigma^2.$$

forward (*X*, *Z=None*, *diag=False*)

17.2.5 Coregionalize

class Coregionalize (*input_dim*, *rank=None*, *components=None*, *diagonal=None*, *active_dims=None*)

Bases: `pyro.contrib.gp.kernels.kernel.Kernel`

A kernel for the linear model of coregionalization $k(x, z) = x^T(WW^T + D)z$ where W is an `input_dim`-by-rank matrix and typically `rank < input_dim`, and D is a diagonal matrix.

This generalizes the `Linear` kernel to multiple features with a low-rank-plus-diagonal weight matrix. The typical use case is for modeling correlations among outputs of a multi-output GP, where outputs are coded as distinct data points with one-hot coded features denoting which output each datapoint represents.

If only `rank` is specified, the kernel $(W W^T + D)$ will be randomly initialized to a matrix with expected value the identity matrix.

References:

[1] **Mauricio A. Alvarez, Lorenzo Rosasco, Neil D. Lawrence (2012)** [Kernels for Vector-Valued Functions: a Review](#)

Parameters

- **input_dim** (*int*) – Number of feature dimensions of inputs.
- **rank** (*int*) – Optional rank. This is only used if `components` is unspecified. If neither `rank` nor `components` is specified, then `rank` defaults to `input_dim`.
- **components** (*torch.Tensor*) – An optional (`input_dim`, `rank`) shaped matrix that maps features to `rank`-many components. If unspecified, this will be randomly initialized.
- **diagonal** (*torch.Tensor*) – An optional vector of length `input_dim`. If unspecified, this will be set to constant `0.5`.
- **active_dims** (*list*) – List of feature dimensions of the input which the kernel acts on.
- **name** (*str*) – Name of the kernel.

forward (*X*, *Z=None*, *diag=False*)

17.2.6 Cosine

class Cosine (*input_dim*, *variance=None*, *lengthscale=None*, *active_dims=None*)

Bases: `pyro.contrib.gp.kernels.isotropic.Isotropy`

Implementation of Cosine kernel:

$$k(x, z) = \sigma^2 \cos\left(\frac{|x-z|}{l}\right).$$

Parameters `lengthscale` (*torch.Tensor*) – Length-scale parameter of this kernel.

forward (*X*, *Z=None*, *diag=False*)

17.2.7 DotProduct

class `DotProduct` (*input_dim*, *variance=None*, *active_dims=None*)

Bases: `pyro.contrib.gp.kernels.kernel.Kernel`

Base class for kernels which are functions of $x \cdot z$.

17.2.8 Exponent

class `Exponent` (*kern*)

Bases: `pyro.contrib.gp.kernels.kernel.Transforming`

Creates a new kernel according to

$$k_{new}(x, z) = \exp(k(x, z)).$$

forward (*X*, *Z=None*, *diag=False*)

17.2.9 Exponential

class `Exponential` (*input_dim*, *variance=None*, *lengthscale=None*, *active_dims=None*)

Bases: `pyro.contrib.gp.kernels.isotropic.Isotropy`

Implementation of Exponential kernel:

$$k(x, z) = \sigma^2 \exp\left(-\frac{|x-z|}{l}\right).$$

forward (*X*, *Z=None*, *diag=False*)

17.2.10 Isotropy

class `Isotropy` (*input_dim*, *variance=None*, *lengthscale=None*, *active_dims=None*)

Bases: `pyro.contrib.gp.kernels.kernel.Kernel`

Base class for a family of isotropic covariance kernels which are functions of the distance $|x - z|/l$, where l is the length-scale parameter.

By default, the parameter `lengthscale` has size 1. To use the isotropic version (different `lengthscale` for each dimension), make sure that `lengthscale` has size equal to `input_dim`.

Parameters `lengthscale` (*torch.Tensor*) – Length-scale parameter of this kernel.

17.2.11 Linear

class `Linear` (*input_dim*, *variance=None*, *active_dims=None*)

Bases: `pyro.contrib.gp.kernels.dot_product.DotProduct`

Implementation of Linear kernel:

$$k(x, z) = \sigma^2 x \cdot z.$$

Doing Gaussian Process regression with linear kernel is equivalent to doing a linear regression.

Note: Here we implement the homogeneous version. To use the inhomogeneous version, consider using `Polynomial` kernel with `degree=1` or making a `Sum` with a `Constant` kernel.

forward ($X, Z=None, \text{diag}=False$)

17.2.12 Matern32

class `Matern32` ($\text{input_dim}, \text{variance}=None, \text{lengthscale}=None, \text{active_dims}=None$)

Bases: `pyro.contrib.gp.kernels.isotropic.Isotropy`

Implementation of Matern32 kernel:

$$k(x, z) = \sigma^2 \left(1 + \sqrt{3} \times \frac{|x-z|}{l} \right) \exp \left(-\sqrt{3} \times \frac{|x-z|}{l} \right).$$

forward ($X, Z=None, \text{diag}=False$)

17.2.13 Matern52

class `Matern52` ($\text{input_dim}, \text{variance}=None, \text{lengthscale}=None, \text{active_dims}=None$)

Bases: `pyro.contrib.gp.kernels.isotropic.Isotropy`

Implementation of Matern52 kernel:

$$k(x, z) = \sigma^2 \left(1 + \sqrt{5} \times \frac{|x-z|}{l} + \frac{5}{3} \times \frac{|x-z|^2}{l^2} \right) \exp \left(-\sqrt{5} \times \frac{|x-z|}{l} \right).$$

forward ($X, Z=None, \text{diag}=False$)

17.2.14 Periodic

class `Periodic` ($\text{input_dim}, \text{variance}=None, \text{lengthscale}=None, \text{period}=None, \text{active_dims}=None$)

Bases: `pyro.contrib.gp.kernels.kernel.Kernel`

Implementation of Periodic kernel:

$$k(x, z) = \sigma^2 \exp \left(-2 \times \frac{\sin^2(\pi(x-z)/p)}{l^2} \right),$$

where p is the period parameter.

References:

[1] *Introduction to Gaussian processes*, David J.C. MacKay

Parameters

- **lengthscale** (`torch.Tensor`) – Length scale parameter of this kernel.
- **period** (`torch.Tensor`) – Period parameter of this kernel.

forward ($X, Z=None, \text{diag}=False$)

17.2.15 Polynomial

class Polynomial (*input_dim, variance=None, bias=None, degree=1, active_dims=None*)

Bases: `pyro.contrib.gp.kernels.dot_product.DotProduct`

Implementation of Polynomial kernel:

$$k(x, z) = \sigma^2 (\text{bias} + x \cdot z)^d.$$

Parameters

- **bias** (`torch.Tensor`) – Bias parameter of this kernel. Should be positive.
- **degree** (`int`) – Degree d of the polynomial.

forward ($X, Z=None, \text{diag}=False$)

17.2.16 Product

class Product (*kern0, kern1*)

Bases: `pyro.contrib.gp.kernels.kernel.Combination`

Returns a new kernel which acts like a product/tensor product of two kernels. The second kernel can be a constant.

forward ($X, Z=None, \text{diag}=False$)

17.2.17 RBF

class RBF (*input_dim, variance=None, lengthscale=None, active_dims=None*)

Bases: `pyro.contrib.gp.kernels.isotropic.Isotropy`

Implementation of Radial Basis Function kernel:

$$k(x, z) = \sigma^2 \exp \left(-0.5 \times \frac{|x-z|^2}{l^2} \right).$$

Note: This kernel also has name *Squared Exponential* in literature.

forward ($X, Z=None, \text{diag}=False$)

17.2.18 RationalQuadratic

class RationalQuadratic (*input_dim, variance=None, lengthscale=None, scale_mixture=None, active_dims=None*)

Bases: `pyro.contrib.gp.kernels.isotropic.Isotropy`

Implementation of RationalQuadratic kernel:

$$k(x, z) = \sigma^2 \left(1 + 0.5 \times \frac{|x-z|^2}{\alpha l^2} \right)^{-\alpha}.$$

Parameters **scale_mixture** (`torch.Tensor`) – Scale mixture (α) parameter of this kernel. Should have size 1.

forward ($X, Z=None, \text{diag}=False$)

17.2.19 Sum

class `Sum(kern0, kern1)`

Bases: `pyro.contrib.gp.kernels.kernel.Combination`

Returns a new kernel which acts like a sum/direct sum of two kernels. The second kernel can be a constant.

forward (X , $Z=None$, $diag=False$)

17.2.20 Transforming

class `Transforming(kern)`

Bases: `pyro.contrib.gp.kernels.kernel.Kernel`

Base class for kernels derived from a kernel by some transforms such as warping, exponent, vertical scaling.

Parameters `kern` (`Kernel`) – The original kernel.

17.2.21 VerticalScaling

class `VerticalScaling(kern, vscaling_fn)`

Bases: `pyro.contrib.gp.kernels.kernel.Transforming`

Creates a new kernel according to

$$k_{new}(x, z) = f(x)k(x, z)f(z),$$

where f is a function.

Parameters `vscaling_fn` (*callable*) – A vertical scaling function f .

forward (X , $Z=None$, $diag=False$)

17.2.22 Warping

class `Warping(kern, iwarping_fn=None, owarping_coef=None)`

Bases: `pyro.contrib.gp.kernels.kernel.Transforming`

Creates a new kernel according to

$$k_{new}(x, z) = q(k(f(x), f(z))),$$

where f is an function and q is a polynomial with non-negative coefficients `owarping_coef`.

We can take advantage of f to combine a Gaussian Process kernel with a deep learning architecture. For example:

```
>>> linear = torch.nn.Linear(10, 3)
>>> # register its parameters to Pyro's ParamStore and wrap it by lambda
>>> # to call the primitive pyro.module each time we use the linear function
>>> pyro_linear_fn = lambda x: pyro.module("linear", linear)(x)
>>> kernel = gp.kernels.Matern52(input_dim=3, lengthscale=torch.ones(3))
>>> warped_kernel = gp.kernels.Warping(kernel, pyro_linear_fn)
```

Reference:

[1] *Deep Kernel Learning*, Andrew G. Wilson, Zhiting Hu, Ruslan Salakhutdinov, Eric P. Xing

Parameters

- **iwarping_fn** (*callable*) – An input warping function f .
- **owarping_coef** (*list*) – A list of coefficients of the output warping polynomial. These coefficients must be non-negative.

forward ($X, Z=None, diag=False$)

17.2.23 WhiteNoise

class WhiteNoise (*input_dim, variance=None, active_dims=None*)

Bases: `pyro.contrib.gp.kernels.kernel.Kernel`

Implementation of WhiteNoise kernel:

$$k(x, z) = \sigma^2 \delta(x, z),$$

where δ is a Dirac delta function.

forward ($X, Z=None, diag=False$)

17.3 Likelihoods

17.3.1 Likelihood

class Likelihood

Bases: `pyro.contrib.gp.parameterized.Parameterized`

Base class for likelihoods used in Gaussian Process.

Every inherited class should implement a forward pass which takes an input f and returns a sample y .

forward ($f_loc, f_var, y=None$)

Samples y given f_loc, f_var .

Parameters

- **f_loc** (`torch.Tensor`) – Mean of latent function output.
- **f_var** (`torch.Tensor`) – Variance of latent function output.
- **y** (`torch.Tensor`) – Training output tensor.

Returns a tensor sampled from likelihood

Return type `torch.Tensor`

17.3.2 Binary

class Binary (*response_function=None*)

Bases: `pyro.contrib.gp.likelihoods.likelihood.Likelihood`

Implementation of Binary likelihood, which is used for binary classification problems.

Binary likelihood uses `Bernoulli` distribution, so the output of `response_function` should be in range $(0, 1)$. By default, we use `sigmoid` function.

Parameters **response_function** (*callable*) – A mapping to correct domain for Binary likelihood.

forward (*f_loc*, *f_var*, *y=None*)
 Samples *y* given *f_loc*, *f_var* according to

$$f \sim \mathcal{N}(f_{loc}, f_{var}),$$

$$y \sim \mathcal{B}(\alpha, \beta)(f).$$

Note: The log likelihood is estimated using Monte Carlo with 1 sample of *f*.

Parameters

- **f_loc** (*torch.Tensor*) – Mean of latent function output.
- **f_var** (*torch.Tensor*) – Variance of latent function output.
- **y** (*torch.Tensor*) – Training output tensor.

Returns a tensor sampled from likelihood

Return type *torch.Tensor*

17.3.3 Gaussian

class Gaussian (*variance=None*)

Bases: *pyro.contrib.gp.likelihoods.likelihood.Likelihood*

Implementation of Gaussian likelihood, which is used for regression problems.

Gaussian likelihood uses *Normal* distribution.

Parameters **variance** (*torch.Tensor*) – A variance parameter, which plays the role of noise in regression problems.

forward (*f_loc*, *f_var*, *y=None*)
 Samples *y* given *f_loc*, *f_var* according to

$$y \sim \mathcal{N}(f_{loc}, f_{var} + \epsilon),$$

where ϵ is the variance parameter of this likelihood.

Parameters

- **f_loc** (*torch.Tensor*) – Mean of latent function output.
- **f_var** (*torch.Tensor*) – Variance of latent function output.
- **y** (*torch.Tensor*) – Training output tensor.

Returns a tensor sampled from likelihood

Return type *torch.Tensor*

17.3.4 MultiClass

class MultiClass (*num_classes*, *response_function=None*)

Bases: `pyro.contrib.gp.likelihoods.likelihood.Likelihood`

Implementation of MultiClass likelihood, which is used for multi-class classification problems.

MultiClass likelihood uses `Categorical` distribution, so `response_function` should normalize its input's rightmost axis. By default, we use `softmax` function.

Parameters

- **num_classes** (*int*) – Number of classes for prediction.
- **response_function** (*callable*) – A mapping to correct domain for MultiClass likelihood.

forward (*f_loc*, *f_var*, *y=None*)

Samples y given f_{loc} , f_{var} according to

$$\begin{aligned} f &\sim \mathbb{N}(\mu, \Sigma)(f_{loc}, f_{var}), \\ y &\sim \text{Categorical}(\pi)(f). \end{aligned}$$

Note: The log likelihood is estimated using Monte Carlo with 1 sample of f .

Parameters

- **f_loc** (`torch.Tensor`) – Mean of latent function output.
- **f_var** (`torch.Tensor`) – Variance of latent function output.
- **y** (`torch.Tensor`) – Training output tensor.

Returns a tensor sampled from likelihood

Return type `torch.Tensor`

17.3.5 Poisson

class Poisson (*response_function=None*)

Bases: `pyro.contrib.gp.likelihoods.likelihood.Likelihood`

Implementation of Poisson likelihood, which is used for count data.

Poisson likelihood uses the `Poisson` distribution, so the output of `response_function` should be positive. By default, we use `torch.exp()` as response function, corresponding to a log-Gaussian Cox process.

Parameters **response_function** (*callable*) – A mapping to positive real numbers.

forward (*f_loc*, *f_var*, *y=None*)

Samples y given f_{loc} , f_{var} according to

$$\begin{aligned} f &\sim \mathbb{N}(\mu, \Sigma)(f_{loc}, f_{var}), \\ y &\sim \text{Poisson}(\lambda)(\exp(f)). \end{aligned}$$

Note: The log likelihood is estimated using Monte Carlo with 1 sample of f .

Parameters

- `f_loc` (`torch.Tensor`) – Mean of latent function output.
- `f_var` (`torch.Tensor`) – Variance of latent function output.
- `y` (`torch.Tensor`) – Training output tensor.

Returns a tensor sampled from likelihood

Return type `torch.Tensor`

17.4 Parameterized

`class Parameterized`

Bases: `torch.nn.modules.module.Module`

A wrapper of `torch.nn.Module` whose parameters can be set constraints, set priors.

Under the hood, we move parameters to a buffer store and create “root” parameters which are used to generate that parameter’s value. For example, if we set a constraint to a parameter, an “unconstrained” parameter will be created, and the constrained value will be transformed from that “unconstrained” parameter.

By default, when we set a prior to a parameter, an auto Delta guide will be created. We can use the method `autoguide()` to setup other auto guides. To fix a parameter to a specific value, it is enough to turn off its “root” parameters’ `requires_grad` flags.

Example:

```
>>> class Linear(Parameterized):
...     def __init__(self, a, b):
...         super(Linear, self).__init__()
...         self.a = Parameter(a)
...         self.b = Parameter(b)
...
...     def forward(self, x):
...         return self.a * x + self.b
...
>>> linear = Linear(torch.tensor(1.), torch.tensor(0.))
>>> linear.set_constraint("a", constraints.positive)
>>> linear.set_prior("b", dist.Normal(0, 1))
>>> linear.autoguide("b", dist.Normal)
>>> assert "a_unconstrained" in dict(linear.named_parameters())
>>> assert "b_loc" in dict(linear.named_parameters())
>>> assert "b_scale_unconstrained" in dict(linear.named_parameters())
>>> assert "a" in dict(linear.named_buffers())
>>> assert "b" in dict(linear.named_buffers())
>>> assert "b_scale" in dict(linear.named_buffers())
```

Note that by default, data of a parameter is a float `torch.Tensor` (unless we use `torch.set_default_tensor_type()` to change default tensor type). To cast these parameters to a correct data type or GPU device, we can call methods such as `double()` or `cuda()`. See `torch.nn.Module` for more information.

set_constraint (*name*, *constraint*)

Sets the constraint of an existing parameter.

Parameters

- **name** (*str*) – Name of the parameter.
- **constraint** (*Constraint*) – A PyTorch constraint. See `torch.distributions.constraints` for a list of constraints.

set_prior (*name*, *prior*)

Sets the constraint of an existing parameter.

Parameters

- **name** (*str*) – Name of the parameter.
- **prior** (*Distribution*) – A Pyro prior distribution.

autoguide (*name*, *dist_constructor*)

Sets an autoguide for an existing parameter with name *name* (mimic the behavior of module `pyro.contrib.autoguide`).

Note: *dist_constructor* should be one of `Delta`, `Normal`, and `MultivariateNormal`. More distribution constructor will be supported in the future if needed.

Parameters

- **name** (*str*) – Name of the parameter.
- **dist_constructor** – A `Distribution` constructor.

set_mode (*mode*)

Sets *mode* of this object to be able to use its parameters in stochastic functions. If *mode*="model", a parameter will get its value from its prior. If *mode*="guide", the value will be drawn from its guide.

Note: This method automatically sets *mode* for submodules which belong to `Parameterized` class.

Parameters **mode** (*str*) – Either “model” or “guide”.

mode

17.5 Util

conditional (*Xnew*, *X*, *kernel*, *f_loc*, *f_scale_tril*=None, *Lff*=None, *full_cov*=False, *whiten*=False, *jitter*=1e-06)

Given X_{new} , predicts loc and covariance matrix of the conditional multivariate normal distribution

$$p(f^*(X_{new}) \mid X, k, f_{loc}, f_{scale_tril}).$$

Here *f_loc* and *f_scale_tril* are variation parameters of the variational distribution

$$q(f \mid f_{loc}, f_{scale_tril}) \sim p(f \mid X, y),$$

where f is the function value of the Gaussian Process given input X

$$p(f(X)) \sim \mathcal{N}(0, k(X, X))$$

and y is computed from f by some likelihood function $p(y|f)$.

In case `f_scale_tril=None`, we consider $f = f_{loc}$ and computes

$$p(f^*(X_{new}) \mid X, k, f).$$

In case `f_scale_tril` is not `None`, we follow the derivation from reference [1]. For the case `f_scale_tril=None`, we follow the popular reference [2].

References:

[1] [Sparse GPs: approximate the posterior, not the model](#)

[2] *Gaussian Processes for Machine Learning*, Carl E. Rasmussen, Christopher K. I. Williams

Parameters

- **Xnew** (`torch.Tensor`) – A new input data.
- **X** (`torch.Tensor`) – An input data to be conditioned on.
- **kernel** (`Kernel`) – A Pyro kernel object.
- **f_loc** (`torch.Tensor`) – Mean of $q(f)$. In case `f_scale_tril=None`, $f_{loc} = f$.
- **f_scale_tril** (`torch.Tensor`) – Lower triangular decomposition of covariance matrix of $q(f)$'s.
- **Lff** (`torch.Tensor`) – Lower triangular decomposition of $kernel(X, X)$ (optional).
- **full_cov** (`bool`) – A flag to decide if we want to return full covariance matrix or just variance.
- **whiten** (`bool`) – A flag to tell if `f_loc` and `f_scale_tril` are already transformed by the inverse of `Lff`.
- **jitter** (`float`) – A small positive term which is added into the diagonal part of a covariance matrix to help stabilize its Cholesky decomposition.

Returns loc and covariance matrix (or variance) of $p(f^*(X_{new}))$

Return type `tuple(torch.Tensor, torch.Tensor)`

train (`gpmodule`, `optimizer=None`, `loss_fn=None`, `retain_graph=None`, `num_steps=1000`)

A helper to optimize parameters for a GP module.

Parameters

- **gpmodule** (`GPMoDel`) – A GP module.
- **optimizer** (`Optimizer`) – A PyTorch optimizer instance. By default, we use Adam with `lr=0.01`.
- **loss_fn** (`callable`) – A loss function which takes inputs are `gpmodule`, `model`, `gpmodule.guide`, and returns ELBO loss. By default, `loss_fn=TraceMeanField_ELBO().differentiable_loss`.
- **retain_graph** (`bool`) – An optional flag of `torch.autograd.backward`.
- **num_steps** (`int`) – Number of steps to run SVI.

Returns a list of losses during the training procedure

Return type `list`

CHAPTER 18

Mini Pyro

This file contains a minimal implementation of the Pyro Probabilistic Programming Language. The API (method signatures, etc.) match that of the full implementation as closely as possible. This file is independent of the rest of Pyro, with the exception of the `pyro.distributions` module.

An accompanying example that makes use of this implementation can be found at `examples/minipyro.py`.

```
class Adam(optim_args)
    Bases: object

    __call__(params)

class JitTrace_ELBO(**kwargs)
    Bases: object

    __call__(model, guide, *args)

class Messenger(fn=None)
    Bases: object

    __call__(*args, **kwargs)

    postprocess_message(msg)

    process_message(msg)

class PlateMessenger(fn, size, dim)
    Bases: pyro.contrib.minipyro.Messenger

    process_message(msg)

class SVI(model, guide, optim, loss)
    Bases: object

    step(*args, **kwargs)

Trace_ELBO(**kwargs)

apply_stack(msg)
```

```
class block (fn=None, hide_fn=<function <lambda>>)
    Bases: pyro.contrib.minipyro.Messenger

    process_message (msg)

elbo (model, guide, *args, **kwargs)

get_param_store ()

param (name, init_value=None, constraint=Real())

plate (name, size, dim)

class replay (fn, guide_trace)
    Bases: pyro.contrib.minipyro.Messenger

    process_message (msg)

sample (name, fn, obs=None)

class trace (fn=None)
    Bases: pyro.contrib.minipyro.Messenger

    get_trace (*args, **kwargs)

    postprocess_message (msg)
```

Optimal Experiment Design

The `pyro.contrib.oed` module provides tools to create optimal experiment designs for pyro models. In particular, it provides estimators for the expected information gain (EIG) criterion.

To estimate the EIG for a particular design, use:

```
def model(design):
    ...

# Select an appropriate EIG estimator, such as
eig = vnmc_eig(model, design, ...)
```

EIG can then be maximised using existing optimisers in `pyro.optim`.

19.1 Expected Information Gain

`laplace_eig(model, design, observation_labels, target_labels, guide, loss, optim, num_steps, final_num_samples, y_dist=None, eig=True, **prior_entropy_kwargs)`

Estimates the expected information gain (EIG) by making repeated Laplace approximations to the posterior.

Parameters

- **model** (*function*) – Pyro stochastic function taking *design* as only argument.
- **design** (*torch.Tensor*) – Tensor of possible designs.
- **observation_labels** (*list*) – labels of sample sites to be regarded as observables.
- **target_labels** (*list*) – labels of sample sites to be regarded as latent variables of interest, i.e. the sites that we wish to gain information about.
- **guide** (*function*) – Pyro stochastic function corresponding to *model*.
- **loss** – a Pyro loss such as `pyro.infer.Trace_ELBO().differentiable_loss`.
- **optim** – optimizer for the loss

- **num_steps** (*int*) – Number of gradient steps to take per sampled pseudo-observation.
- **final_num_samples** (*int*) – Number of y samples (pseudo-observations) to take.
- **y_dist** – Distribution to sample y from- if *None* we use the Bayesian marginal distribution.
- **eig** (*bool*) – Whether to compute the EIG or the average posterior entropy (APE). The EIG is given by $EIG = \text{prior entropy} - APE$. If *True*, the prior entropy will be estimated analytically, or by Monte Carlo as appropriate for the *model*. If *False* the APE is returned.
- **prior_entropy_kwargs** (*dict*) – parameters for estimating the prior entropy: *num_prior_samples* indicating the number of samples for a MC estimate of prior entropy, and *mean_field* indicating if an analytic form for a mean-field prior should be tried.

Returns EIG estimate

Return type `torch.Tensor`

vi_eig (*model*, *design*, *observation_labels*, *target_labels*, *vi_parameters*, *is_parameters*, *y_dist=None*, *eig=True*, ***prior_entropy_kwargs*)
Estimates the expected information gain (EIG) using variational inference (VI).

The APE is defined as

$$APE(d) = E_{Y \sim p(y|\theta, d)}[H(p(\theta|Y, d))]$$

where $H[p(x)]$ is the [differential entropy](#). The APE is related to expected information gain (EIG) by the equation

$$EIG(d) = H[p(\theta)] - APE(d)$$

in particular, minimising the APE is equivalent to maximising EIG.

Parameters

- **model** (*function*) – A pyro model accepting *design* as only argument.
- **design** (`torch.Tensor`) – Tensor representation of design
- **observation_labels** (*list*) – A subset of the sample sites present in *model*. These sites are regarded as future observations and other sites are regarded as latent variables over which a posterior is to be inferred.
- **target_labels** (*list*) – A subset of the sample sites over which the posterior entropy is to be measured.
- **vi_parameters** (*dict*) – Variational inference parameters which should include: *optim*: an instance of `pyro.Optim`, *guide*: a guide function compatible with *model*, *num_steps*: the number of VI steps to make, and *loss*: the loss function to use for VI
- **is_parameters** (*dict*) – Importance sampling parameters for the marginal distribution of Y . May include *num_samples*: the number of samples to draw from the marginal.
- **y_dist** (`pyro.distributions.Distribution`) – (optional) the distribution assumed for the response variable Y
- **eig** (*bool*) – Whether to compute the EIG or the average posterior entropy (APE). The EIG is given by $EIG = \text{prior entropy} - APE$. If *True*, the prior entropy will be estimated analytically, or by Monte Carlo as appropriate for the *model*. If *False* the APE is returned.
- **prior_entropy_kwargs** (*dict*) – parameters for estimating the prior entropy: *num_prior_samples* indicating the number of samples for a MC estimate of prior entropy, and *mean_field* indicating if an analytic form for a mean-field prior should be tried.

Returns EIG estimate

Return type `torch.Tensor`

nmc_eig (*model*, *design*, *observation_labels*, *target_labels=None*, *N=100*, *M=10*, *M_prime=None*, *independent_priors=False*)

Nested Monte Carlo estimate of the expected information gain (EIG). The estimate is, when there are not any random effects,

$$\frac{1}{N} \sum_{n=1}^N \log p(y_n | \theta_n, d) - \frac{1}{N} \sum_{n=1}^N \log \left(\frac{1}{M} \sum_{m=1}^M p(y_n | \theta_m, d) \right)$$

The estimate is, in the presence of random effects,

$$\frac{1}{N} \sum_{n=1}^N \log \left(\frac{1}{M'} \sum_{m=1}^{M'} p(y_n | \theta_n, \tilde{\theta}_{nm}, d) \right) - \frac{1}{N} \sum_{n=1}^N \log \left(\frac{1}{M} \sum_{m=1}^M p(y_n | \theta_m, \tilde{\theta}_m, d) \right)$$

The latter form is used when *M_prime != None*.

param function model A pyro model accepting *design* as only argument.

param torch.Tensor design Tensor representation of design

param list observation_labels A subset of the sample sites present in *model*. These sites are regarded as future observations and other sites are regarded as latent variables over which a posterior is to be inferred.

param list target_labels A subset of the sample sites over which the posterior entropy is to be measured.

param int N Number of outer expectation samples.

param int M Number of inner expectation samples for $p(y|d)$.

param int M_prime Number of samples for $p(y | \theta, d)$ if required.

param bool independent_priors Only used when *M_prime* is not *None*. Indicates whether the prior distributions for the target variables and the nuisance variables are independent. In this case, it is not necessary to sample the targets conditional on the nuisance variables.

return EIG estimate

rtype *torch.Tensor*

donsker_varadhan_eig (*model*, *design*, *observation_labels*, *target_labels*, *num_samples*, *num_steps*, *T*, *optim*, *return_history=False*, *final_design=None*, *final_num_samples=None*)
Donsker-Varadhan estimate of the expected information gain (EIG).

The Donsker-Varadhan representation of EIG is

$$\sup_T E_{p(y, \theta | d)} [T(y, \theta)] - \log E_{p(y | d) p(\theta)} [\exp(T(\bar{y}, \bar{\theta}))]$$

where T is any (measurable) function.

This methods optimises the loss function over a pre-specified class of functions T .

Parameters

- **model** (*function*) – A pyro model accepting *design* as only argument.
- **design** (*torch.Tensor*) – Tensor representation of design
- **observation_labels** (*list*) – A subset of the sample sites present in *model*. These sites are regarded as future observations and other sites are regarded as latent variables over which a posterior is to be inferred.

- **target_labels** (*list*) – A subset of the sample sites over which the posterior entropy is to be measured.
- **num_samples** (*int*) – Number of samples per iteration.
- **num_steps** (*int*) – Number of optimisation steps.
- or **torch.nn.Module T** (*function*) – optimisable function T for use in the Donsker-Varadhan loss function.
- **optim** (*pyro.optim.Optim*) – Optimiser to use.
- **return_history** (*bool*) – If *True*, also returns a tensor giving the loss function at each step of the optimisation.
- **final_design** (*torch.Tensor*) – The final design tensor to evaluate at. If *None*, uses *design*.
- **final_num_samples** (*int*) – The number of samples to use at the final evaluation, If *None*, uses 'num_samples'.

Returns EIG estimate, optionally includes full optimisation history

Return type *torch.Tensor* or *tuple*

posterior_eig (*model*, *design*, *observation_labels*, *target_labels*, *num_samples*, *num_steps*, *guide*, *optim*, *return_history=False*, *final_design=None*, *final_num_samples=None*, *eig=True*, *prior_entropy_kwargs={}*, **args*, ***kwargs*)

Posterior estimate of expected information gain (EIG) computed from the average posterior entropy (APE) using $EIG = \text{prior entropy} - \text{APE}$. See [1] for full details.

The posterior representation of APE is

$$\sup_q E_{p(y, \theta|d)} [\log q(\theta|y, d)]$$

where q is any distribution on θ .

This method optimises the loss over a given guide family *guide* representing q .

[1] Foster, Adam, et al. “Variational Bayesian Optimal Experimental Design.” arXiv preprint arXiv:1903.05480 (2019).

Parameters

- **model** (*function*) – A pyro model accepting *design* as only argument.
- **design** (*torch.Tensor*) – Tensor representation of design
- **observation_labels** (*list*) – A subset of the sample sites present in *model*. These sites are regarded as future observations and other sites are regarded as latent variables over which a posterior is to be inferred.
- **target_labels** (*list*) – A subset of the sample sites over which the posterior entropy is to be measured.
- **num_samples** (*int*) – Number of samples per iteration.
- **num_steps** (*int*) – Number of optimisation steps.
- **guide** (*function*) – guide family for use in the (implicit) posterior estimation. The parameters of *guide* are optimised to maximise the posterior objective.
- **optim** (*pyro.optim.Optim*) – Optimiser to use.
- **return_history** (*bool*) – If *True*, also returns a tensor giving the loss function at each step of the optimisation.

- **final_design** (*torch.Tensor*) – The final design tensor to evaluate at. If *None*, uses *design*.
- **final_num_samples** (*int*) – The number of samples to use at the final evaluation, If *None*, uses *num_samples*.
- **eig** (*bool*) – Whether to compute the EIG or the average posterior entropy (APE). The EIG is given by $EIG = \text{prior entropy} - APE$. If *True*, the prior entropy will be estimated analytically, or by Monte Carlo as appropriate for the *model*. If *False* the APE is returned.
- **prior_entropy_kwargs** (*dict*) – parameters for estimating the prior entropy: *num_prior_samples* indicating the number of samples for a MC estimate of prior entropy, and *mean_field* indicating if an analytic form for a mean-field prior should be tried.

Returns EIG estimate, optionally includes full optimisation history

Return type *torch.Tensor* or *tuple*

marginal_eig (*model*, *design*, *observation_labels*, *target_labels*, *num_samples*, *num_steps*, *guide*, *optim*, *return_history=False*, *final_design=None*, *final_num_samples=None*)

Estimate EIG by estimating the marginal entropy $p(y|d)$. See [1] for full details.

The marginal representation of EIG is

$$\inf_q E_{p(y, \theta|d)} \left[\log \frac{p(y|\theta, d)}{q(y|d)} \right]$$

where q is any distribution on y .

Warning: this method does **not** estimate the correct quantity in the presence of random effects.

[1] Foster, Adam, et al. “Variational Bayesian Optimal Experimental Design.” arXiv preprint arXiv:1903.05480 (2019).

Parameters

- **model** (*function*) – A pyro model accepting *design* as only argument.
- **design** (*torch.Tensor*) – Tensor representation of design
- **observation_labels** (*list*) – A subset of the sample sites present in *model*. These sites are regarded as future observations and other sites are regarded as latent variables over which a posterior is to be inferred.
- **target_labels** (*list*) – A subset of the sample sites over which the posterior entropy is to be measured.
- **num_samples** (*int*) – Number of samples per iteration.
- **num_steps** (*int*) – Number of optimisation steps.
- **guide** (*function*) – guide family for use in the marginal estimation. The parameters of *guide* are optimised to maximise the log-likelihood objective.
- **optim** (*pyro.optim.Optim*) – Optimiser to use.
- **return_history** (*bool*) – If *True*, also returns a tensor giving the loss function at each step of the optimisation.
- **final_design** (*torch.Tensor*) – The final design tensor to evaluate at. If *None*, uses *design*.
- **final_num_samples** (*int*) – The number of samples to use at the final evaluation, If *None*, uses *num_samples*.

Returns EIG estimate, optionally includes full optimisation history

Return type *torch.Tensor* or *tuple*

lfire_eig(*model*, *design*, *observation_labels*, *target_labels*, *num_y_samples*, *num_theta_samples*, *num_steps*, *classifier*, *optim*, *return_history=False*, *final_design=None*, *final_num_samples=None*)

Estimates the EIG using the method of Likelihood-Free Inference by Ratio Estimation (LFIRE) as in [1]. LFIRE is run separately for several samples of θ .

[1] Kleinegesse, Steven, and Michael Gutmann. “Efficient Bayesian Experimental Design for Implicit Models.” arXiv preprint arXiv:1810.09912 (2018).

Parameters

- **model** (*function*) – A pyro model accepting *design* as only argument.
- **design** (*torch.Tensor*) – Tensor representation of design
- **observation_labels** (*list*) – A subset of the sample sites present in *model*. These sites are regarded as future observations and other sites are regarded as latent variables over which a posterior is to be inferred.
- **target_labels** (*list*) – A subset of the sample sites over which the posterior entropy is to be measured.
- **num_y_samples** (*int*) – Number of samples to take in y for each θ .
- **num_steps** (*int*) – Number of optimisation steps.
- **classifier** (*function*) – a Pytorch or Pyro classifier used to distinguish between samples of y under $p(y|d)$ and samples under $p(y|\theta, d)$ for some θ .
- **optim** (*pyro.optim.Optim*) – Optimiser to use.
- **return_history** (*bool*) – If *True*, also returns a tensor giving the loss function at each step of the optimisation.
- **final_design** (*torch.Tensor*) – The final design tensor to evaluate at. If *None*, uses *design*.
- **final_num_samples** (*int*) – The number of samples to use at the final evaluation, If *None*, uses ‘*num_samples*’.

Param int num_theta_samples: Number of initial samples in θ to take. The likelihood ratio is estimated by LFIRE for each sample.

Returns EIG estimate, optionally includes full optimisation history

Return type *torch.Tensor* or *tuple*

vnmc_eig(*model*, *design*, *observation_labels*, *target_labels*, *num_samples*, *num_steps*, *guide*, *optim*, *return_history=False*, *final_design=None*, *final_num_samples=None*)

Estimates the EIG using Variational Nested Monte Carlo (VNMC). The VNMC estimate [1] is

$$\frac{1}{N} \sum_{n=1}^N \left[\log p(y_n | \theta_n, d) - \log \left(\frac{1}{M} \sum_{m=1}^M \frac{p(\theta_{mn}) p(y_n | \theta_{mn}, d)}{q(\theta_{mn} | y_n)} \right) \right]$$

where $q(\theta|y)$ is the learned variational posterior approximation and $\theta_n, y_n \sim p(\theta, y|d)$, $\theta_{mn} \sim q(\theta|y = y_n)$.

As $N \rightarrow \infty$ this is an upper bound on EIG. We minimise this upper bound by stochastic gradient descent.

[1] Foster, Adam, et al. “Variational Bayesian Optimal Experimental Design.” arXiv preprint arXiv:1903.05480 (2019).

Parameters

- **model** (*function*) – A pyro model accepting *design* as only argument.
- **design** (*torch.Tensor*) – Tensor representation of design
- **observation_labels** (*list*) – A subset of the sample sites present in *model*. These sites are regarded as future observations and other sites are regarded as latent variables over which a posterior is to be inferred.
- **target_labels** (*list*) – A subset of the sample sites over which the posterior entropy is to be measured.
- **num_samples** (*tuple*) – Number of (N, M) samples per iteration.
- **num_steps** (*int*) – Number of optimisation steps.
- **guide** (*function*) – guide family for use in the posterior estimation. The parameters of *guide* are optimised to minimise the VNMC upper bound.
- **optim** (*pyro.optim.Optim*) – Optimiser to use.
- **return_history** (*bool*) – If *True*, also returns a tensor giving the loss function at each step of the optimisation.
- **final_design** (*torch.Tensor*) – The final design tensor to evaluate at. If *None*, uses *design*.
- **final_num_samples** (*tuple*) – The number of (N, M) samples to use at the final evaluation, If *None*, uses 'num_samples'.

Returns EIG estimate, optionally includes full optimisation history

Return type *torch.Tensor* or *tuple*

20.1 Data Association

class `MarginalAssignment` (*exists_logits*, *assign_logits*, *bp_iters=None*)

Computes marginal data associations between objects and detections.

This assumes that each detection corresponds to zero or one object, and each object corresponds to zero or more detections. Specifically this does not assume detections have been partitioned into frames of mutual exclusion as is common in 2-D assignment problems.

Parameters

- **exists_logits** (*torch.Tensor*) – a tensor of shape `[num_objects]` representing per-object factors for existence of each potential object.
- **assign_logits** (*torch.Tensor*) – a tensor of shape `[num_detections, num_objects]` representing per-edge factors of assignment probability, where each edge denotes that a given detection associates with a single object.
- **bp_iters** (*int*) – optional number of belief propagation iterations. If unspecified or `None` an expensive exact algorithm will be used.

Variables

- **num_detections** (*int*) – the number of detections
- **num_objects** (*int*) – the number of (potentially existing) objects
- **exists_dist** (*pyro.distributions.Bernoulli*) – a mean field posterior distribution over object existence.
- **assign_dist** (*pyro.distributions.Categorical*) – a mean field posterior distribution over the object (or `None`) to which each detection associates. This has `.event_shape == (num_objects + 1,)` where the final element denotes spurious detection, and `.batch_shape == (num_frames, num_detections)`.

```
class MarginalAssignmentSparse (num_objects, num_detections, edges, exists_logits, assign_logits,
                                bp_iters)
```

A cheap sparse version of *MarginalAssignment*.

Parameters

- **num_detections** (*int*) – the number of detections
- **num_objects** (*int*) – the number of (potentially existing) objects
- **edges** (*torch.LongTensor*) – a $[2, \text{num_edges}]$ -shaped tensor of (detection, object) index pairs specifying feasible associations.
- **exists_logits** (*torch.Tensor*) – a tensor of shape $[\text{num_objects}]$ representing per-object factors for existence of each potential object.
- **assign_logits** (*torch.Tensor*) – a tensor of shape $[\text{num_edges}]$ representing per-edge factors of assignment probability, where each edge denotes that a given detection associates with a single object.
- **bp_iters** (*int*) – optional number of belief propagation iterations. If unspecified or None an expensive exact algorithm will be used.

Variables

- **num_detections** (*int*) – the number of detections
- **num_objects** (*int*) – the number of (potentially existing) objects
- **exists_dist** (*pyro.distributions.Bernoulli*) – a mean field posterior distribution over object existence.
- **assign_dist** (*pyro.distributions.Categorical*) – a mean field posterior distribution over the object (or None) to which each detection associates. This has `.event_shape == (num_objects + 1,)` where the final element denotes spurious detection, and `.batch_shape == (num_frames, num_detections)`.

```
class MarginalAssignmentPersistent (exists_logits,      assign_logits,      bp_iters=None,
                                    bp_momentum=0.5)
```

This computes marginal distributions of a multi-frame multi-object data association problem with an unknown number of persistent objects.

The inputs are factors in a factor graph (existence probabilities for each potential object and assignment probabilities for each object-detection pair), and the outputs are marginal distributions of posterior existence probability of each potential object and posterior assignment probabilities of each object-detection pair.

This assumes a shared (maximum) number of detections per frame; to handle variable number of detections, simply set corresponding elements of `assign_logits` to `-float('inf')`.

Parameters

- **exists_logits** (*torch.Tensor*) – a tensor of shape $[\text{num_objects}]$ representing per-object factors for existence of each potential object.
- **assign_logits** (*torch.Tensor*) – a tensor of shape $[\text{num_frames}, \text{num_detections}, \text{num_objects}]$ representing per-edge factors of assignment probability, where each edge denotes that at a given time frame a given detection associates with a single object.
- **bp_iters** (*int*) – optional number of belief propagation iterations. If unspecified or None an expensive exact algorithm will be used.
- **bp_momentum** (*float*) – optional momentum to use for belief propagation. Should be in the interval $[0, 1)$.

Variables

- **num_frames** (*int*) – the number of time frames
- **num_detections** (*int*) – the (maximum) number of detections per frame
- **num_objects** (*int*) – the number of (potentially existing) objects
- **exists_dist** (`pyro.distributions.Bernoulli`) – a mean field posterior distribution over object existence.
- **assign_dist** (`pyro.distributions.Categorical`) – a mean field posterior distribution over the object (or None) to which each detection associates. This has `.event_shape == (num_objects + 1,)` where the final element denotes spurious detection, and `.batch_shape == (num_frames, num_detections)`.

compute_marginals (*exists_logits, assign_logits*)

This implements exact inference of pairwise marginals via enumeration. This is very expensive and is only useful for testing.

See *MarginalAssignment* for args and problem description.

compute_marginals_bp (*exists_logits, assign_logits, bp_iters*)

This implements approximate inference of pairwise marginals via loopy belief propagation, adapting the approach of [1].

See *MarginalAssignment* for args and problem description.

[1] Jason L. Williams, Roslyn A. Lau (2014) Approximate evaluation of marginal association probabilities with belief propagation <https://arxiv.org/abs/1209.6299>

compute_marginals_sparse_bp (*num_objects, num_detections, edges, exists_logits, assign_logits, bp_iters*)

This implements approximate inference of pairwise marginals via loopy belief propagation, adapting the approach of [1].

See *MarginalAssignmentSparse* for args and problem description.

[1] Jason L. Williams, Roslyn A. Lau (2014) Approximate evaluation of marginal association probabilities with belief propagation <https://arxiv.org/abs/1209.6299>

compute_marginals_persistent (*exists_logits, assign_logits*)

This implements exact inference of pairwise marginals via enumeration. This is very expensive and is only useful for testing.

See *MarginalAssignmentPersistent* for args and problem description.

compute_marginals_persistent_bp (*exists_logits, assign_logits, bp_iters, bp_momentum=0.5*)

This implements approximate inference of pairwise marginals via loopy belief propagation, adapting the approach of [1], [2].

See *MarginalAssignmentPersistent* for args and problem description.

[1] Jason L. Williams, Roslyn A. Lau (2014) Approximate evaluation of marginal association probabilities with belief propagation <https://arxiv.org/abs/1209.6299>

[2] Ryan Turner, Steven Bottone, Bhargav Avasarala (2014) A Complete Variational Tracker <https://papers.nips.cc/paper/5572-a-complete-variational-tracker.pdf>

20.2 Distributions

class `EKFDistribution` (*x0*, *P0*, *dynamic_model*, *measurement_cov*, *time_steps=1*, *dt=1.0*, *validate_args=None*)

Distribution over EKF states. See [EKState](#). Currently only supports *log_prob*.

Parameters

- **x0** (*torch.Tensor*) – PV tensor (mean)
- **P0** (*torch.Tensor*) – covariance
- **dynamic_model** – *DynamicModel* object
- **measurement_cov** (*torch.Tensor*) – measurement covariance
- **time_steps** (*int*) – number time step
- **dt** (*torch.Tensor*) – time step

filter_states (*value*)

Returns the ekf states given measurements

Parameters **value** (*torch.Tensor*) – measurement means of shape (*time_steps*, *event_shape*)

log_prob (*value*)

Returns the joint log probability of the innovations of a tensor of measurements

Parameters **value** (*torch.Tensor*) – measurement means of shape (*time_steps*, *event_shape*)

20.3 Dynamic Models

class `DynamicModel` (*dimension*, *dimension_pv*, *num_process_noise_parameters=None*)

Dynamic model interface.

Parameters

- **dimension** – native state dimension.
- **dimension_pv** – PV state dimension.
- **num_process_noise_parameters** – process noise parameter space dimension. This for UKF applications. Can be left as `None` for EKF and most other filters.

dimension

Native state dimension access.

dimension_pv

PV state dimension access.

num_process_noise_parameters

Process noise parameters space dimension access.

forward (*x*, *dt*, *do_normalization=True*)

Integrate native state *x* over time interval *dt*.

Parameters

- **x** – current native state. If the *DynamicModel* is non-differentiable, be sure to handle the case of *x* being augmented with process noise parameters.

- **dt** – time interval to integrate over.
- **do_normalization** – whether to perform normalization on output, e.g., mod'ing angles into an interval.

Returns Native state x integrated dt into the future.

geodesic_difference ($x1, x0$)

Compute and return the geodesic difference between 2 native states. This is a generalization of the Euclidean operation $x1 - x0$.

Parameters

- **x1** – native state.
- **x0** – native state.

Returns Geodesic difference between native states $x1$ and $x2$.

mean2pv (x)

Compute and return PV state from native state. Useful for combining state estimates of different types in IMM (Interacting Multiple Model) filtering.

Parameters **x** – native state estimate mean.

Returns PV state estimate mean.

cov2pv (P)

Compute and return PV covariance from native covariance. Useful for combining state estimates of different types in IMM (Interacting Multiple Model) filtering.

Parameters **P** – native state estimate covariance.

Returns PV state estimate covariance.

process_noise_cov ($dt=0.0$)

Compute and return process noise covariance (Q).

Parameters **dt** – time interval to integrate over.

Returns Read-only covariance (Q). For a `DifferentiableDynamicModel`, this is the covariance of the native state x resulting from stochastic integration (for use with EKF). Otherwise, it is the covariance directly of the process noise parameters (for use with UKF).

process_noise_dist ($dt=0.0$)

Return a distribution object of state displacement from the process noise distribution over a time interval.

Parameters **dt** – time interval that process noise accumulates over.

Returns `MultivariateNormal`.

class DifferentiableDynamicModel (*dimension, dimension_pv, num_process_noise_parameters=None*)

DynamicModel for which state transition Jacobians can be efficiently calculated, usu. analytically or by automatic differentiation.

jacobian (dt)

Compute and return native state transition Jacobian (F) over time interval dt .

Parameters **dt** – time interval to integrate over.

Returns Read-only Jacobian (F) of integration map (f).

class Ncp (*dimension, sv2*)

NCP (Nearly-Constant Position) dynamic model. May be subclassed, e.g., with CWNV (Continuous White Noise Velocity) or DWNV (Discrete White Noise Velocity).

Parameters

- **dimension** – native state dimension.
- **sv2** – variance of velocity. Usually chosen so that the standard deviation is roughly half of the max velocity one would ever expect to observe.

forward (*x*, *dt*, *do_normalization=True*)Integrate native state *x* over time interval *dt*.**Parameters**

- **x** – current native state. If the DynamicModel is non-differentiable, be sure to handle the case of *x* being augmented with process noise parameters.
- **dt** – time interval to integrate over. *do_normalization*: whether to perform normalization on output, e.g., mod'ing angles into an interval. Has no effect for this subclass.

Returns Native state *x* integrated *dt* into the future.**mean2pv** (*x*)

Compute and return PV state from native state. Useful for combining state estimates of different types in IMM (Interacting Multiple Model) filtering.

Parameters **x** – native state estimate mean.**Returns** PV state estimate mean.**cov2pv** (*P*)

Compute and return PV covariance from native covariance. Useful for combining state estimates of different types in IMM (Interacting Multiple Model) filtering.

Parameters **P** – native state estimate covariance.**Returns** PV state estimate covariance.**jacobian** (*dt*)Compute and return cached native state transition Jacobian (F) over time interval *dt*.**Parameters** **dt** – time interval to integrate over.**Returns** Read-only Jacobian (F) of integration map (f).**process_noise_cov** (*dt=0.0*)

Compute and return cached process noise covariance (Q).

Parameters **dt** – time interval to integrate over.**Returns** Read-only covariance (Q) of the native state *x* resulting from stochastic integration (for use with EKF).**class Ncv** (*dimension*, *sa2*)

NCV (Nearly-Constant Velocity) dynamic model. May be subclassed, e.g., with CWNA (Continuous White Noise Acceleration) or DWNA (Discrete White Noise Acceleration).

Parameters

- **dimension** – native state dimension.
- **sa2** – variance of acceleration. Usually chosen so that the standard deviation is roughly half of the max acceleration one would ever expect to observe.

forward (*x*, *dt*, *do_normalization=True*)Integrate native state *x* over time interval *dt*.**Parameters**

- **x** – current native state. If the DynamicModel is non-differentiable, be sure to handle the case of **x** being augmented with process noise parameters.
- **dt** – time interval to integrate over.
- **do_normalization** – whether to perform normalization on output, e.g., mod'ing angles into an interval. Has no effect for this subclass.

Returns Native state **x** integrated **dt** into the future.

mean2pv (*x*)

Compute and return PV state from native state. Useful for combining state estimates of different types in IMM (Interacting Multiple Model) filtering.

Parameters **x** – native state estimate mean.

Returns PV state estimate mean.

cov2pv (*P*)

Compute and return PV covariance from native covariance. Useful for combining state estimates of different types in IMM (Interacting Multiple Model) filtering.

Parameters **P** – native state estimate covariance.

Returns PV state estimate covariance.

jacobian (*dt*)

Compute and return cached native state transition Jacobian (**F**) over time interval **dt**.

Parameters **dt** – time interval to integrate over.

Returns Read-only Jacobian (**F**) of integration map (**f**).

process_noise_cov (*dt=0.0*)

Compute and return cached process noise covariance (**Q**).

Parameters **dt** – time interval to integrate over.

Returns Read-only covariance (**Q**) of the native state **x** resulting from stochastic integration (for use with EKF).

class NcpContinuous (*dimension, sv2*)

NCP (Nearly-Constant Position) dynamic model with CWNV (Continuous White Noise Velocity).

References: “Estimation with Applications to Tracking and Navigation” by Y. Bar- Shalom et al, 2001, p.269.

Parameters

- **dimension** – native state dimension.
- **sv2** – variance of velocity. Usually chosen so that the standard deviation is roughly half of the max velocity one would ever expect to observe.

process_noise_cov (*dt=0.0*)

Compute and return cached process noise covariance (**Q**).

Parameters **dt** – time interval to integrate over.

Returns Read-only covariance (**Q**) of the native state **x** resulting from stochastic integration (for use with EKF).

class NcvContinuous (*dimension, sa2*)

NCV (Nearly-Constant Velocity) dynamic model with CWNA (Continuous White Noise Acceleration).

References: “Estimation with Applications to Tracking and Navigation” by Y. Bar- Shalom et al, 2001, p.269.

Parameters

- **dimension** – native state dimension.
- **sa2** – variance of acceleration. Usually chosen so that the standard deviation is roughly half of the max acceleration one would ever expect to observe.

process_noise_cov (*dt=0.0*)

Compute and return cached process noise covariance (Q).

Parameters **dt** – time interval to integrate over.

Returns Read-only covariance (Q) of the native state x resulting from stochastic integration (for use with EKF).

class NcpDiscrete (*dimension, sv2*)

NCP (Nearly-Constant Position) dynamic model with DWNV (Discrete White Noise Velocity).

Parameters

- **dimension** – native state dimension.
- **sv2** – variance of velocity. Usually chosen so that the standard deviation is roughly half of the max velocity one would ever expect to observe.

References: “Estimation with Applications to Tracking and Navigation” by Y. Bar- Shalom et al, 2001, p.273.

process_noise_cov (*dt=0.0*)

Compute and return cached process noise covariance (Q).

Parameters **dt** – time interval to integrate over.

Returns Read-only covariance (Q) of the native state x resulting from stochastic integration (for use with EKF).

class NcvDiscrete (*dimension, sa2*)

NCV (Nearly-Constant Velocity) dynamic model with DWNA (Discrete White Noise Acceleration).

Parameters

- **dimension** – native state dimension.
- **sa2** – variance of acceleration. Usually chosen so that the standard deviation is roughly half of the max acceleration one would ever expect to observe.

References: “Estimation with Applications to Tracking and Navigation” by Y. Bar- Shalom et al, 2001, p.273.

process_noise_cov (*dt=0.0*)

Compute and return cached process noise covariance (Q).

Parameters **dt** – time interval to integrate over.

Returns Read-only covariance (Q) of the native state x resulting from stochastic integration (for use with EKF). (Note that this Q, modulo numerical error, has rank $\text{dimension}/2$. So, it is only positive semi-definite.)

20.4 Extended Kalman Filter

class EKFState (*dynamic_model, mean, cov, time=None, frame_num=None*)

State-Centric EKF (Extended Kalman Filter) for use with either an NCP (Nearly-Constant Position) or NCV

(Nearly-Constant Velocity) target dynamic model. Stores a target dynamic model, state estimate, and state time. Incoming `Measurement` provide sensor information for updates.

Warning: For efficiency, the dynamic model is only shallow-copied. Make a deep copy outside as necessary to protect against unexpected changes.

Parameters

- **dynamic_model** – target dynamic model.
- **mean** – mean of target state estimate.
- **cov** – covariance of target state estimate.
- **time** – time of state estimate.

dynamic_model

Dynamic model access.

dimension

Native state dimension access.

mean

Native state estimate mean access.

cov

Native state estimate covariance access.

dimension_pv

PV state dimension access.

mean_pv

Compute and return cached PV state estimate mean.

cov_pv

Compute and return cached PV state estimate covariance.

time

Continuous State time access.

frame_num

Discrete State time access.

predict (*dt=None, destination_time=None, destination_frame_num=None*)

Use dynamic model to predict (aka propagate aka integrate) state estimate in-place.

Parameters

- **dt** – time to integrate over. The state time will be automatically incremented this amount unless you provide `destination_time`. Using `destination_time` may be preferable for prevention of roundoff error accumulation.
- **destination_time** – optional value to set continuous state time to after integration. If this is not provided, then `destination_frame_num` must be.
- **destination_frame_num** – optional value to set discrete state time to after integration. If this is not provided, then `destination_frame_num` must be.

innovation (*measurement*)

Compute and return the innovation that a measurement would induce if it were used for an update, but

don't actually perform the update. Assumes state and measurement are time-aligned. Useful for computing Chi^2 stats and likelihoods.

Parameters `measurement` – measurement

Returns Innovation mean and covariance of hypothetical update.

Return type `tuple(torch.Tensor, torch.Tensor)`

`log_likelihood_of_update` (*measurement*)

Compute and return the likelihood of a potential update, but don't actually perform the update. Assumes state and measurement are time-aligned. Useful for gating and calculating costs in assignment problems for data association.

Param `measurement`.

Returns Likelihood of hypothetical update.

`update` (*measurement*)

Use measurement to update state estimate in-place and return innovation. The innovation is useful, e.g., for evaluating filter consistency or updating model likelihoods when the `EKFState` is part of an `IMMState`.

Param `measurement`.

Returns EKF State, Innovation mean and covariance.

20.5 Hashing

class `LSH` (*radius*)

Implements locality-sensitive hashing for low-dimensional euclidean space.

Allows to efficiently find neighbours of a point. Provides 2 guarantees:

- Difference between coordinates of points not returned by `nearby()` and input point is larger than `radius`.
- Difference between coordinates of points returned by `nearby()` and input point is smaller than `2 * radius`.

Example:

```
>>> radius = 1
>>> lsh = LSH(radius)
>>> a = torch.tensor([-0.51, -0.51]) # hash(a)=(-1,-1)
>>> b = torch.tensor([-0.49, -0.49]) # hash(a)=(0,0)
>>> c = torch.tensor([1.0, 1.0]) # hash(b)=(1,1)
>>> lsh.add('a', a)
>>> lsh.add('b', b)
>>> lsh.add('c', c)
>>> # even though c is within 2radius of a
>>> lsh.nearby('a') # doctest: +SKIP
{'b'}
>>> lsh.nearby('b') # doctest: +SKIP
{'a', 'c'}
>>> lsh.remove('b')
>>> lsh.nearby('a') # doctest: +SKIP
set()
```


Parameters **radius** (*float*) – Scaling parameter used in hash function. Determines the size of the neighbourhood.

add (*key*, *point*)

Adds (*key*, *point*) pair to the hash.

Parameters

- **key** – Key used identify *point*.
- **point** (*torch.Tensor*) – data, should be detached and on cpu.

remove (*key*)

Removes *key* and corresponding point from the hash.

Raises `KeyError` if *key* is not in hash.

Parameters **key** – key used to identify point.

nearby (*key*)

Returns a set of keys which are neighbours of the point identified by *key*.

Two points are nearby if difference of each element of their hashes is smaller than 2. In euclidean space, this corresponds to all points \mathbf{p} where $|\mathbf{p}_k - (\mathbf{p}_{\text{key}})_k| < r$, and some points (all points not guaranteed) where $|\mathbf{p}_k - (\mathbf{p}_{\text{key}})_k| < 2r$.

Parameters **key** – key used to identify input point.

Returns a set of keys identifying neighbours of the input point.

Return type `set`

class **ApproxSet** (*radius*)

Queries low-dimensional euclidean space for approximate occupancy.

Parameters **radius** (*float*) – scaling parameter used in hash function. Determines the size of the bin. See [LSH](#) for details.

try_add (*point*)

Attempts to add *point* to set. Only adds there are no points in the *point*'s bin.

Parameters **point** (*torch.Tensor*) – Point to be queried, should be detached and on cpu.

Returns `True` if point is successfully added, `False` if there is already a point in *point*'s bin.

Return type `bool`

merge_points (*points*, *radius*)

Greedy merge points that are closer than given radius.

This uses [LSH](#) to achieve complexity that is linear in the number of merged clusters and quadratic in the size of the largest merged cluster.

Parameters

- **points** (*torch.Tensor*) – A tensor of shape (K, D) where K is the number of points and D is the number of dimensions.
- **radius** (*float*) – The minimum distance nearer than which points will be merged.

Returns A tuple (*merged_points*, *groups*) where *merged_points* is a tensor of shape (J, D) where $J \leq K$, and *groups* is a list of tuples of indices mapping merged points to original points. Note that $\text{len}(\text{groups}) == J$ and $\text{sum}(\text{len}(\text{group}) \text{ for group in groups}) == K$.

Return type `tuple`

20.6 Measurements

class Measurement (*mean, cov, time=None, frame_num=None*)

Gaussian measurement interface.

Parameters

- **mean** – mean of measurement distribution.
- **cov** – covariance of measurement distribution.
- **time** – continuous time of measurement. If this is not provided, *frame_num* must be.
- **frame_num** – discrete time of measurement. If this is not provided, *time* must be.

dimension

Measurement space dimension access.

mean

Measurement mean (z in most Kalman Filtering literature).

cov

Noise covariance (R in most Kalman Filtering literature).

time

Continuous time of measurement.

frame_num

Discrete time of measurement.

geodesic_difference (*z1, z0*)

Compute and return the geodesic difference between 2 measurements. This is a generalization of the Euclidean operation $z1 - z0$.

Parameters

- **z1** – measurement.
- **z0** – measurement.

Returns Geodesic difference between $z1$ and $z2$.

class DifferentiableMeasurement (*mean, cov, time=None, frame_num=None*)

Interface for Gaussian measurement for which Jacobians can be efficiently calculated, usu. analytically or by automatic differentiation.

jacobian (*x=None*)

Compute and return Jacobian (H) of measurement map (h) at target PV state x .

Parameters **x** – PV state. Use default argument `None` when the Jacobian is not state-dependent.

Returns Read-only Jacobian (H) of measurement map (h).

class PositionMeasurement (*mean, cov, time=None, frame_num=None*)

Full-rank Gaussian position measurement in Euclidean space.

Parameters

- **mean** – mean of measurement distribution.
- **cov** – covariance of measurement distribution.
- **time** – time of measurement.

jacobian (*x=None*)

Compute and return Jacobian (H) of measurement map (h) at target PV state x .

Parameters \mathbf{x} – PV state. The default argument `None` may be used in this subclass since the Jacobian is not state-dependent.

Returns Read-only Jacobian (H) of measurement map (h).

CHAPTER 21

Indices and tables

- `genindex`
- `search`

p

pyro.contrib.autoguide, 103
pyro.contrib.autoguide.initialization, 110
pyro.contrib.autoname, 111
pyro.contrib.autoname.named, 113
pyro.contrib.autoname.scoping, 115
pyro.contrib.bnn, 119
pyro.contrib.bnn.hidden_layer, 119
pyro.contrib.easyguide, 121
pyro.contrib.glm, 125
pyro.contrib.gp, 127
pyro.contrib.gp.kernels, 137
pyro.contrib.gp.likelihoods, 144
pyro.contrib.gp.models.gplvm, 137
pyro.contrib.gp.models.gpr, 130
pyro.contrib.gp.models.model, 127
pyro.contrib.gp.models.sgpr, 132
pyro.contrib.gp.models.vgp, 134
pyro.contrib.gp.models.vsgp, 135
pyro.contrib.gp.parameterized, 147
pyro.contrib.gp.util, 148
pyro.contrib.minipyro, 149
pyro.contrib.oed, 153
pyro.contrib.oed.eig, 153
pyro.contrib.tracking, 161
pyro.contrib.tracking.assignment, 161
pyro.contrib.tracking.distributions, 164
pyro.contrib.tracking.dynamic_models, 164
pyro.contrib.tracking.extended_kalman_filter, 168
pyro.contrib.tracking.hashing, 170
pyro.contrib.tracking.measurements, 172
pyro.distributions.torch, 27
pyro.generic, 101
pyro.infer.abstract_infer, 18
pyro.infer.discrete, 17
pyro.infer.elbo, 10
pyro.infer.importance, 16
pyro.infer.renyi_elbo, 15
pyro.infer.svi, 9
pyro.infer.trace_elbo, 10
pyro.infer.trace_mean_field_elbo, 14
pyro.infer.trace_tail_adaptive_elbo, 15
pyro.infer.traceenum_elbo, 12
pyro.infer.tracegraph_elbo, 11
pyro.nn.auto_reg_nn, 65
pyro.ops.dual_averaging, 91
pyro.ops.einsum, 96
pyro.ops.indexing, 95
pyro.ops.integrator, 92
pyro.ops.newton, 93
pyro.ops.stats, 98
pyro.ops.welford, 92
pyro.optim.adagrad_rmsprop, 70
pyro.optim.clipped_adam, 71
pyro.optim.lr_scheduler, 70
pyro.optim.multi, 72
pyro.optim.optim, 69
pyro.optim.pytorch_optimizers, 71
pyro.params.param_store, 61
pyro.poutine.block_messenger, 85
pyro.poutine.broadcast_messenger, 86
pyro.poutine.condition_messenger, 86
pyro.poutine.escape_messenger, 86
pyro.poutine.handlers, 75
pyro.poutine.indep_messenger, 86
pyro.poutine.lift_messenger, 87
pyro.poutine.messenger, 84
pyro.poutine.replay_messenger, 87
pyro.poutine.runtime, 88
pyro.poutine.scale_messenger, 87
pyro.poutine.trace_messenger, 88
pyro.poutine.util, 89

Symbols

[__call__\(\)](#) (*Adam method*), 151
[__call__\(\)](#) (*AutoCallable method*), 105
[__call__\(\)](#) (*AutoContinuous method*), 106
[__call__\(\)](#) (*AutoDelta method*), 105
[__call__\(\)](#) (*AutoDiscreteParallel method*), 109
[__call__\(\)](#) (*AutoGuide method*), 103
[__call__\(\)](#) (*AutoGuideList method*), 104
[__call__\(\)](#) (*Distribution method*), 31
[__call__\(\)](#) (*EasyGuide method*), 121
[__call__\(\)](#) (*JitTrace_ELBO method*), 151
[__call__\(\)](#) (*Messenger method*), 151
[__call__\(\)](#) (*PyroLRScheduler method*), 70
[__call__\(\)](#) (*PyroOptim method*), 69
[__call__\(\)](#) (*TorchDistributionMixin method*), 32

A

[Adadelta\(\)](#) (*in module pyro.optim.pytorch_optimizers*), 72
[Adagrad\(\)](#) (*in module pyro.optim.pytorch_optimizers*), 71
[AdagradRMSPROP](#) (*class in pyro.optim.adagrad_rmprop*), 70
[AdagradRMSPROP\(\)](#) (*in module pyro.optim.optim*), 70
[Adam](#) (*class in pyro.contrib.minipyro*), 151
[Adam\(\)](#) (*in module pyro.optim.pytorch_optimizers*), 71
[Adamax\(\)](#) (*in module pyro.optim.pytorch_optimizers*), 71
[add\(\)](#) (*AutoGuideList method*), 104
[add\(\)](#) (*List method*), 115
[add\(\)](#) (*LSH method*), 171
[add_edge\(\)](#) (*Trace method*), 83
[add_node\(\)](#) (*Trace method*), 83
[AffineCoupling](#) (*class in pyro.distributions.transforms*), 47
[all_escape\(\)](#) (*in module pyro.poutine.util*), 89
[am_i_wrapped\(\)](#) (*in module pyro.poutine.runtime*), 88
[apply_stack\(\)](#) (*in module pyro.contrib.minipyro*), 151
[apply_stack\(\)](#) (*in module pyro.poutine.runtime*), 88
[ApproxSet](#) (*class in pyro.contrib.tracking.hashing*), 171
[arg_constraints](#) (*AVFMultivariateNormal attribute*), 35
[arg_constraints](#) (*BetaBinomial attribute*), 36
[arg_constraints](#) (*Delta attribute*), 36
[arg_constraints](#) (*DirichletMultinomial attribute*), 37
[arg_constraints](#) (*Empirical attribute*), 38
[arg_constraints](#) (*GammaPoisson attribute*), 39
[arg_constraints](#) (*GaussianScaleMixture attribute*), 40
[arg_constraints](#) (*InverseGamma attribute*), 40
[arg_constraints](#) (*LKJCorrCholesky attribute*), 41
[arg_constraints](#) (*MaskedMixture attribute*), 41
[arg_constraints](#) (*MixtureOfDiagNormals attribute*), 42
[arg_constraints](#) (*MixtureOfDiagNormalsSharedCovariance attribute*), 43
[arg_constraints](#) (*OMTMultivariateNormal attribute*), 43
[arg_constraints](#) (*SpanningTree attribute*), 45
[arg_constraints](#) (*VonMises attribute*), 46
[arg_constraints](#) (*VonMises3D attribute*), 47
[ASGD\(\)](#) (*in module pyro.optim.pytorch_optimizers*), 71
[AutoCallable](#) (*class in pyro.contrib.autoguide*), 104
[AutoContinuous](#) (*class in pyro.contrib.autoguide*), 106
[autocorrelation\(\)](#) (*in module pyro.ops.stats*), 98
[autocovariance\(\)](#) (*in module pyro.ops.stats*), 98
[AutoDelta](#) (*class in pyro.contrib.autoguide*), 105
[AutoDiagonalNormal](#) (*class in pyro.contrib.autoguide*), 107
[AutoDiscreteParallel](#) (*class in pyro.contrib.autoguide*), 109
[AutoGuide](#) (*class in pyro.contrib.autoguide*), 103
[autoguide\(\)](#) (*Parameterized method*), 148

AutoGuideList (class in *pyro.contrib.autoguide*), 104
 AutoIAFNormal (class in *pyro.contrib.autoguide*), 108
 AutoLaplaceApproximation (class in *pyro.contrib.autoguide*), 109
 AutoLowRankMultivariateNormal (class in *pyro.contrib.autoguide*), 108
 AutoMultivariateNormal (class in *pyro.contrib.autoguide*), 107
 autoregressive (*BlockAutoregressive* attribute), 50
 autoregressive (*InverseAutoregressiveFlow* attribute), 54
 autoregressive (*PolynomialFlow* attribute), 58
 AutoRegressiveNN (class in *pyro.nn.auto_reg_nn*), 65
 AVFMultivariateNormal (class in *pyro.distributions*), 35

B

BackwardSampleMessenger (class in *pyro.infer.traceenum_elbo*), 12
 BatchNormTransform (class in *pyro.distributions.transforms*), 48
 Bernoulli (class in *pyro.distributions*), 27
 Beta (class in *pyro.distributions*), 27
 BetaBinomial (class in *pyro.distributions*), 35
 bijective (*AffineCoupling* attribute), 48
 bijective (*BatchNormTransform* attribute), 49
 bijective (*BlockAutoregressive* attribute), 50
 bijective (*HouseholderFlow* attribute), 53
 bijective (*InverseAutoregressiveFlow* attribute), 54
 bijective (*InverseAutoregressiveFlowStable* attribute), 55
 bijective (*PermuteTransform* attribute), 56
 bijective (*PlanarFlow* attribute), 56
 bijective (*PolynomialFlow* attribute), 58
 bijective (*RadialFlow* attribute), 58
 bijective (*SylvesterFlow* attribute), 59
 Binary (class in *pyro.contrib.gp.likelihoods*), 144
 Binomial (class in *pyro.distributions*), 27
 block (class in *pyro.contrib.minipyro*), 151
 block() (in module *pyro.poutine*), 76
 BlockAutoregressive (class in *pyro.distributions.transforms*), 49
 BlockMessenger (class in *pyro.poutine.block_messenger*), 85
 broadcast() (in module *pyro.poutine*), 77
 BroadcastMessenger (class in *pyro.poutine.broadcast_messenger*), 86
 Brownian (class in *pyro.contrib.gp.kernels*), 138

C

Categorical (class in *pyro.distributions*), 27
 Cauchy (class in *pyro.distributions*), 28
 Chi2 (class in *pyro.distributions*), 28

cleanup() (*HMC* method), 23
 clear() (*ParamStoreDict* method), 61
 clear_cache() (*HMC* method), 23
 clear_param_store() (in module *pyro*), 8
 ClippedAdam (class in *pyro.optim.clipped_adam*), 71
 ClippedAdam() (in module *pyro.optim.optim*), 70
 codomain (*AffineCoupling* attribute), 48
 codomain (*BatchNormTransform* attribute), 49
 codomain (*BlockAutoregressive* attribute), 50
 codomain (*HouseholderFlow* attribute), 53
 codomain (*InverseAutoregressiveFlow* attribute), 54
 codomain (*InverseAutoregressiveFlowStable* attribute), 55
 codomain (*PermuteTransform* attribute), 56
 codomain (*PlanarFlow* attribute), 57
 codomain (*PolynomialFlow* attribute), 58
 codomain (*RadialFlow* attribute), 58
 codomain (*SylvesterFlow* attribute), 59
 Combination (class in *pyro.contrib.gp.kernels*), 138
 compute_log_prob() (*Trace* method), 83
 compute_marginals() (in module *pyro.contrib.tracking.assignment*), 163
 compute_marginals() (*TraceEnum_ELBO* method), 13
 compute_marginals_bp() (in module *pyro.contrib.tracking.assignment*), 163
 compute_marginals_persistent() (in module *pyro.contrib.tracking.assignment*), 163
 compute_marginals_persistent_bp() (in module *pyro.contrib.tracking.assignment*), 163
 compute_marginals_sparse_bp() (in module *pyro.contrib.tracking.assignment*), 163
 compute_score_parts() (*Trace* method), 83
 concentration (*DirichletMultinomial* attribute), 37
 concentration (*GammaPoisson* attribute), 39
 concentration (*InverseGamma* attribute), 40
 concentration0 (*BetaBinomial* attribute), 36
 concentration1 (*BetaBinomial* attribute), 36
 CondIndepStackFrame (class in *pyro.poutine.indep_messenger*), 86
 condition() (in module *pyro.poutine*), 77
 conditional() (in module *pyro.contrib.gp.util*), 148
 ConditionMessenger (class in *pyro.poutine.condition_messenger*), 86
 config_enumerate() (in module *pyro.infer.enum*), 81
 Constant (class in *pyro.contrib.gp.kernels*), 139
 constrained_gamma (*BatchNormTransform* attribute), 49
 contract() (in module *pyro.ops.einsum*), 96
 contract_expression() (in module *pyro.ops.einsum*), 96
 copy() (*Trace* method), 83

- Coregionalize (class in `pyro.contrib.gp.kernels`), 139
- Cosine (class in `pyro.contrib.gp.kernels`), 139
- CosineAnnealingLR() (in module `pyro.optim.pytorch_optimizers`), 72
- CosineAnnealingWarmRestarts() (in module `pyro.optim.pytorch_optimizers`), 72
- cov (EKFState attribute), 169
- cov (Measurement attribute), 172
- cov2pv() (DynamicModel method), 165
- cov2pv() (Ncp method), 166
- cov2pv() (Ncv method), 167
- cov_pv (EKFState attribute), 169
- create_mask() (in module `pyro.nn.auto_reg_nn`), 66
- current_backend (GenericModule attribute), 101
- CyclicLR() (in module `pyro.optim.pytorch_optimizers`), 72
- ## D
- DeepELUFlow (class in `pyro.distributions.transforms`), 50
- DeepLeakyReLUFlow (class in `pyro.distributions.transforms`), 51
- DeepSigmoidalFlow (class in `pyro.distributions.transforms`), 52
- default_process_message() (in module `pyro.poutine.runtime`), 89
- Delta (class in `pyro.distributions`), 36
- diagnostics() (HMC method), 23
- diagnostics() (in module `pyro.infer.mcmc.util`), 25
- diagnostics() (MCMC method), 21
- Dict (class in `pyro.contrib.autoname.named`), 115
- differentiable_loss() (JitTrace_ELBO method), 11
- differentiable_loss() (JitTraceEnum_ELBO method), 14
- differentiable_loss() (JitTraceMean-Field_ELBO method), 14
- differentiable_loss() (Trace_ELBO method), 11
- differentiable_loss() (TraceEnum_ELBO method), 13
- DifferentiableDynamicModel (class in `pyro.contrib.tracking.dynamic_models`), 165
- DifferentiableMeasurement (class in `pyro.contrib.tracking.measurements`), 172
- dimension (DynamicModel attribute), 164
- dimension (EKFState attribute), 169
- dimension (Measurement attribute), 172
- dimension_pv (DynamicModel attribute), 164
- dimension_pv (EKFState attribute), 169
- Dirichlet (class in `pyro.distributions`), 28
- DirichletMultinomial (class in `pyro.distributions`), 37
- discrete_escape() (in module `pyro.poutine.util`), 89
- Distribution (class in `pyro.distributions`), 31
- do() (in module `pyro.poutine`), 78
- domain (AffineCoupling attribute), 48
- domain (BatchNormTransform attribute), 49
- domain (BlockAutoregressive attribute), 50
- domain (HouseholderFlow attribute), 53
- domain (InverseAutoregressiveFlow attribute), 54
- domain (InverseAutoregressiveFlowStable attribute), 55
- domain (PlanarFlow attribute), 57
- domain (PolynomialFlow attribute), 58
- domain (RadialFlow attribute), 58
- domain (SylvesterFlow attribute), 59
- donsker_varadhan_eig() (in module `pyro.contrib.oed.eig`), 155
- DotProduct (class in `pyro.contrib.gp.kernels`), 140
- dtanh_dx() (SylvesterFlow method), 59
- DualAveraging (class in `pyro.ops.dual_averaging`), 91
- dynamic_model (EKFState attribute), 169
- DynamicModel (class in `pyro.contrib.tracking.dynamic_models`), 164
- ## E
- easy_guide() (in module `pyro.contrib.easyguide`), 122
- EasyGuide (class in `pyro.contrib.easyguide`), 121
- edges (Trace attribute), 83
- effective() (in module `pyro.poutine.runtime`), 89
- effective_sample_size() (in module `pyro.ops.stats`), 98
- einsum() (in module `pyro.ops.contract`), 96
- EKFDistribution (class in `pyro.contrib.tracking.distributions`), 164
- EKFState (class in `pyro.contrib.tracking.extended_kalman_filter`), 168
- ELBO (class in `pyro.infer.elbo`), 10
- elbo() (in module `pyro.contrib.minipyro`), 152
- Empirical (class in `pyro.distributions`), 37
- empirical (Marginals attribute), 19
- EmpiricalMarginal (class in `pyro.infer.abstract_infer`), 18
- enable_validation() (in module `pyro`), 8
- enable_validation() (in module `pyro.poutine.util`), 89
- enum() (in module `pyro.poutine`), 78
- enum_extend() (in module `pyro.poutine.util`), 89
- enumerate_support() (BetaBinomial method), 36
- enumerate_support() (Distribution method), 31
- enumerate_support() (Empirical method), 38
- enumerate_support() (SpanningTree method), 45
- escape() (in module `pyro.poutine`), 78

EscapeMessenger (class *pyro.poutine.escape_messenger*), 86
 evaluate_loss() (SVI method), 9
 event_dim (AffineCoupling attribute), 48
 event_dim (BatchNormTransform attribute), 49
 event_dim (BlockAutoregressive attribute), 50
 event_dim (HouseholderFlow attribute), 53
 event_dim (InverseAutoregressiveFlow attribute), 54
 event_dim (InverseAutoregressiveFlowStable attribute), 55
 event_dim (PermuteTransform attribute), 56
 event_dim (PlanarFlow attribute), 57
 event_dim (PolynomialFlow attribute), 58
 event_dim (RadialFlow attribute), 58
 event_dim (SylvesterFlow attribute), 59
 event_dim (TorchDistributionMixin attribute), 33
 event_shape (Empirical attribute), 38
 expand() (BetaBinomial method), 36
 expand() (Delta method), 36
 expand() (DirichletMultinomial method), 37
 expand() (GammaPoisson method), 39
 expand() (InverseGamma method), 40
 expand() (LKJCorrCholesky method), 41
 expand() (MaskedMixture method), 41
 expand() (MixtureOfDiagNormals method), 42
 expand() (MixtureOfDiagNormalsSharedCovariance method), 43
 expand() (VonMises method), 46
 expand() (VonMises3D method), 47
 expand_by() (TorchDistributionMixin method), 33
 Exponent (class in *pyro.contrib.gp.kernels*), 140
 Exponential (class in *pyro.contrib.gp.kernels*), 140
 Exponential (class in *pyro.distributions*), 28
 ExponentialFamily (class in *pyro.distributions*), 28
 ExponentialLR() (in module *pyro.optim.pytorch_optimizers*), 72

F

filter_states() (EKFDistribution method), 164
 FisherSnedecor (class in *pyro.distributions*), 28
 fit_generalized_pareto() (in module *pyro.ops.stats*), 100
 format_shapes() (Trace method), 83
 forward() (AutoRegressiveNN method), 66
 forward() (Binary method), 144
 forward() (Brownian method), 138
 forward() (Constant method), 139
 forward() (Coregionalize method), 139
 forward() (Cosine method), 140
 forward() (DynamicModel method), 164
 forward() (Exponent method), 140
 forward() (Exponential method), 140
 forward() (Gaussian method), 145
 forward() (GPLVM method), 137

in forward() (GPModel method), 129
 forward() (GPRegression method), 131
 forward() (Kernel method), 138
 forward() (Likelihood method), 144
 forward() (Linear method), 141
 forward() (MaskedLinear method), 66
 forward() (Matern32 method), 141
 forward() (Matern52 method), 141
 forward() (MultiClass method), 146
 forward() (Ncp method), 166
 forward() (Ncv method), 166
 forward() (Periodic method), 141
 forward() (Poisson method), 146
 forward() (Polynomial method), 142
 forward() (Product method), 142
 forward() (RationalQuadratic method), 142
 forward() (RBF method), 142
 forward() (SparseGPRegression method), 133
 forward() (Sum method), 143
 forward() (VariationalGP method), 135
 forward() (VariationalSparseGP method), 136
 forward() (VerticalScaling method), 143
 forward() (Warping method), 144
 forward() (WhiteNoise method), 144
 frame_num (EKFState attribute), 169
 frame_num (Measurement attribute), 172

G

Gamma (class in *pyro.distributions*), 28
 GammaPoisson (class in *pyro.distributions*), 38
 Gaussian (class in *pyro.contrib.gp.likelihoods*), 145
 GaussianScaleMixture (class in *pyro.distributions*), 39
 gelman_rubin() (in module *pyro.ops.stats*), 98
 GenericModule (class in *pyro.generic*), 101
 geodesic_difference() (DynamicModel method), 165
 geodesic_difference() (Measurement method), 172
 Geometric (class in *pyro.distributions*), 28
 get_all_param_names() (ParamStoreDict method), 62
 get_covariance() (WelfordCovariance method), 93
 get_ESS() (Importance method), 16
 get_log_normalizer() (Importance method), 16
 get_normalized_weights() (Importance method), 16
 get_param() (ParamStoreDict method), 62
 get_param_store() (in module *pyro*), 8
 get_param_store() (in module *pyro.contrib.minipyro*), 152
 get_permutation() (AutoRegressiveNN method), 66
 get_posterior() (AutoContinuous method), 106

[get_posterior\(\)](#) (*AutoDiagonalNormal method*), 107
[get_posterior\(\)](#) (*AutoIAFNormal method*), 109
[get_posterior\(\)](#) (*AutoLaplaceApproximation method*), 109
[get_posterior\(\)](#) (*AutoLowRankMultivariateNormal method*), 108
[get_posterior\(\)](#) (*AutoMultivariateNormal method*), 107
[get_samples\(\)](#) (*MCMC method*), 21
[get_state\(\)](#) (*DualAveraging method*), 92
[get_state\(\)](#) (*ParamStoreDict method*), 63
[get_state\(\)](#) (*PyroOptim method*), 69
[get_step\(\)](#) (*MixedMultiOptimizer method*), 73
[get_step\(\)](#) (*MultiOptimizer method*), 73
[get_step\(\)](#) (*Newton method*), 73
[get_trace\(\)](#) (*trace method*), 152
[get_trace\(\)](#) (*TraceHandler method*), 88
[get_trace\(\)](#) (*TraceMessenger method*), 88
[GPLVM](#) (*class in pyro.contrib.gp.models.gplvm*), 137
[GPModel](#) (*class in pyro.contrib.gp.models.model*), 127
[GPRRegression](#) (*class in pyro.contrib.gp.models.gpr*), 130
[Group](#) (*class in pyro.contrib.easyguide.easyguide*), 122
[group\(\)](#) (*EasyGuide method*), 122
[guide](#) (*Group attribute*), 123
[guide\(\)](#) (*EasyGuide method*), 121
[guide\(\)](#) (*GPLVM method*), 137
[guide\(\)](#) (*GPModel method*), 128
[guide\(\)](#) (*GPRRegression method*), 131
[guide\(\)](#) (*SparseGPRRegression method*), 133
[guide\(\)](#) (*VariationalGP method*), 135
[guide\(\)](#) (*VariationalSparseGP method*), 136
[Gumbel](#) (*class in pyro.distributions*), 28

H

[HalfCauchy](#) (*class in pyro.distributions*), 29
[HalfNormal](#) (*class in pyro.distributions*), 29
[has_enumerate_support](#) (*BetaBinomial attribute*), 36
[has_enumerate_support](#) (*Distribution attribute*), 32
[has_enumerate_support](#) (*Empirical attribute*), 38
[has_enumerate_support](#) (*SpanningTree attribute*), 45
[has_rsample](#) (*Delta attribute*), 36
[has_rsample](#) (*Distribution attribute*), 32
[has_rsample](#) (*GaussianScaleMixture attribute*), 40
[has_rsample](#) (*InverseGamma attribute*), 40
[has_rsample](#) (*LKJCorrCholesky attribute*), 41
[has_rsample](#) (*MaskedMixture attribute*), 41
[has_rsample](#) (*MixtureOfDiagNormals attribute*), 42
[has_rsample](#) (*MixtureOfDiagNormalsSharedCovariance attribute*), 43

[has_rsample](#) (*Rejector attribute*), 44
[has_rsample](#) (*VonMises attribute*), 46
[HiddenLayer](#) (*class in pyro.contrib.bnn.hidden_layer*), 119
[HMC](#) (*class in pyro.infer.mcmc*), 21
[HouseholderFlow](#) (*class in pyro.distributions.transforms*), 52
[hpdi\(\)](#) (*in module pyro.ops.stats*), 99

I

[identify_dense_edges\(\)](#) (*in module pyro.poutine.trace_messenger*), 88
[Importance](#) (*class in pyro.infer.importance*), 16
[Independent](#) (*class in pyro.distributions*), 29
[independent\(\)](#) (*TorchDistributionMixin method*), 34
[IndepMessenger](#) (*class in pyro.poutine.indep_messenger*), 87
[indices](#) (*IndepMessenger attribute*), 87
[infer_config\(\)](#) (*in module pyro.poutine*), 78
[infer_discrete\(\)](#) (*in module pyro.infer.discrete*), 17
[information_criterion\(\)](#) (*TracePosterior method*), 19
[init\(\)](#) (*EasyGuide method*), 121
[init_to_feasible\(\)](#) (*in module pyro.contrib.autoguide.initialization*), 110
[init_to_mean\(\)](#) (*in module pyro.contrib.autoguide.initialization*), 110
[init_to_median\(\)](#) (*in module pyro.contrib.autoguide.initialization*), 110
[init_to_sample\(\)](#) (*in module pyro.contrib.autoguide.initialization*), 110
[initial_params](#) (*HMC attribute*), 23
[initialize_model\(\)](#) (*in module pyro.infer.mcmc.util*), 25
[InitMessenger](#) (*class in pyro.contrib.autoguide.initialization*), 110
[innovation\(\)](#) (*EKFState method*), 169
[inv_permutation](#) (*PermuteTransform attribute*), 56
[inverse_mass_matrix](#) (*HMC attribute*), 23
[InverseAutoregressiveFlow](#) (*class in pyro.distributions.transforms*), 53
[InverseAutoregressiveFlowStable](#) (*class in pyro.distributions.transforms*), 54
[InverseGamma](#) (*class in pyro.distributions*), 40
[is_validation_enabled\(\)](#) (*in module pyro.poutine.util*), 90
[Isotropy](#) (*class in pyro.contrib.gp.kernels*), 140
[items\(\)](#) (*ParamStoreDict method*), 61
[iter_sample\(\)](#) (*GPRRegression method*), 131
[iter_stochastic_nodes\(\)](#) (*Trace method*), 83

J

[jacobian\(\)](#) (*DifferentiableDynamicModel method*),

- 165
- `jacobian()` (*DifferentiableMeasurement method*), 172
- `jacobian()` (*Ncp method*), 166
- `jacobian()` (*Ncv method*), 167
- `jacobian()` (*PositionMeasurement method*), 172
- `JitTrace_ELBO` (*class in pyro.contrib.minipyro*), 151
- `JitTrace_ELBO` (*class in pyro.infer.trace_elbo*), 11
- `JitTraceEnum_ELBO` (*class in pyro.infer.traceenum_elbo*), 13
- `JitTraceGraph_ELBO` (*class in pyro.infer.tracegraph_elbo*), 12
- `JitTraceMeanField_ELBO` (*class in pyro.infer.trace_mean_field_elbo*), 14
- ## K
- `Kernel` (*class in pyro.contrib.gp.kernels*), 137
- `keys()` (*ParamStoreDict method*), 62

L

`LambdaLR()` (*in module pyro.optim.pytorch_optimizers*), 72

`Laplace` (*class in pyro.distributions*), 29

`laplace_approximation()` (*AutoLaplaceApproximation method*), 109

`laplace_eig()` (*in module pyro.contrib.oed.eig*), 153

`lfire_eig()` (*in module pyro.contrib.oed.eig*), 158

`lift()` (*in module pyro.poutine*), 78

`LiftMessenger` (*class in pyro.poutine.lift_messenger*), 87

`Likelihood` (*class in pyro.contrib.gp.likelihoods*), 144

`Linear` (*class in pyro.contrib.gp.kernels*), 140

`List` (*class in pyro.contrib.autoname.named*), 114

`lkj_constant()` (*LKJCorrCholesky method*), 41

`LKJCorrCholesky` (*class in pyro.distributions*), 40

`load()` (*ParamStoreDict method*), 63

`load()` (*PyroOptim method*), 70

`log_abs_det_jacobian()` (*AffineCoupling method*), 48

`log_abs_det_jacobian()` (*BatchNormTransform method*), 49

`log_abs_det_jacobian()` (*BlockAutoregressive method*), 50

`log_abs_det_jacobian()` (*HouseholderFlow method*), 53

`log_abs_det_jacobian()` (*InverseAutoregressiveFlow method*), 54

`log_abs_det_jacobian()` (*InverseAutoregressiveFlowStable method*), 55

`log_abs_det_jacobian()` (*PermuteTransform method*), 56

`log_abs_det_jacobian()` (*PlanarFlow method*), 57

`log_abs_det_jacobian()` (*PolynomialFlow method*), 58

`log_abs_det_jacobian()` (*RadialFlow method*), 58

`log_abs_det_jacobian()` (*SylvesterFlow method*), 59

`log_likelihood_of_update()` (*EKFState method*), 170

`log_partition_function` (*SpanningTree attribute*), 45

`log_prob()` (*BetaBinomial method*), 36

`log_prob()` (*Delta method*), 36

`log_prob()` (*DirichletMultinomial method*), 37

`log_prob()` (*Distribution method*), 32

`log_prob()` (*EKFDistribution method*), 164

`log_prob()` (*Empirical method*), 38

`log_prob()` (*GammaPoisson method*), 39

`log_prob()` (*GaussianScaleMixture method*), 40

`log_prob()` (*LKJCorrCholesky method*), 41

`log_prob()` (*MaskedMixture method*), 41

`log_prob()` (*MixtureOfDiagNormals method*), 42

`log_prob()` (*MixtureOfDiagNormalsSharedCovariance method*), 43

`log_prob()` (*Rejector method*), 44

`log_prob()` (*RelaxedBernoulliStraightThrough method*), 43

`log_prob()` (*RelaxedOneHotCategoricalStraightThrough method*), 44

`log_prob()` (*SpanningTree method*), 45

`log_prob()` (*VonMises method*), 46

`log_prob()` (*VonMises3D method*), 47

`log_prob_sum()` (*Trace method*), 83

`log_weights` (*Empirical attribute*), 38

`logging()` (*HMC method*), 23

`LogisticNormal` (*class in pyro.distributions*), 29

`LogNormal` (*class in pyro.distributions*), 29

`loss()` (*RenyiELBO method*), 16

`loss()` (*Trace_ELBO method*), 11

`loss()` (*TraceEnum_ELBO method*), 13

`loss()` (*TraceGraph_ELBO method*), 12

`loss()` (*TraceMeanField_ELBO method*), 14

`loss()` (*TraceTailAdaptive_ELBO method*), 15

`loss_and_grads()` (*JitTrace_ELBO method*), 11

`loss_and_grads()` (*JitTraceEnum_ELBO method*), 14

`loss_and_grads()` (*JitTraceGraph_ELBO method*), 12

`loss_and_grads()` (*JitTraceMeanField_ELBO method*), 15

`loss_and_grads()` (*RenyiELBO method*), 16

`loss_and_grads()` (*Trace_ELBO method*), 11

`loss_and_grads()` (*TraceEnum_ELBO method*), 13

`loss_and_grads()` (*TraceGraph_ELBO method*), 12

`loss_and_surrogate_loss()` (*JitTrace_ELBO method*), 11

LowRankMultivariateNormal (class in *pyro.distributions*), 29
 LSH (class in *pyro.contrib.tracking.hashing*), 170

M

map_estimate() (EasyGuide method), 122
 map_estimate() (Group method), 123
 marginal() (TracePosterior method), 19
 marginal() (TracePredictive method), 20
 marginal_eig() (in module *pyro.contrib.oed.eig*), 157
 MarginalAssignment (class in *pyro.contrib.tracking.assignment*), 161
 MarginalAssignmentPersistent (class in *pyro.contrib.tracking.assignment*), 162
 MarginalAssignmentSparse (class in *pyro.contrib.tracking.assignment*), 161
 Marginals (class in *pyro.infer.abstract_infer*), 19
 markov() (in module *pyro.poutine*), 79
 mask() (in module *pyro.poutine*), 79
 mask() (TorchDistributionMixin method), 34
 MaskedLinear (class in *pyro.nn.auto_reg_nn*), 66
 MaskedMixture (class in *pyro.distributions*), 41
 match() (ParamStoreDict method), 62
 Matern32 (class in *pyro.contrib.gp.kernels*), 141
 Matern52 (class in *pyro.contrib.gp.kernels*), 141
 mc_extend() (in module *pyro.poutine.util*), 90
 MCMC (class in *pyro.infer.mcmc.api*), 20
 mean (BetaBinomial attribute), 36
 mean (Delta attribute), 36
 mean (DirichletMultinomial attribute), 37
 mean (EKFFState attribute), 169
 mean (Empirical attribute), 38
 mean (GammaPoisson attribute), 39
 mean (MaskedMixture attribute), 41
 mean (Measurement attribute), 172
 mean (VonMises attribute), 46
 mean2pv() (DynamicModel method), 165
 mean2pv() (Ncp method), 166
 mean2pv() (Ncv method), 167
 mean_pv (EKFFState attribute), 169
 Measurement (class in *pyro.contrib.tracking.measurements*), 172
 median() (AutoContinuous method), 106
 median() (AutoDelta method), 106
 median() (AutoGuide method), 103
 median() (AutoGuideList method), 104
 merge_points() (in module *pyro.contrib.tracking.hashing*), 171
 Messenger (class in *pyro.contrib.minipyro*), 151
 Messenger (class in *pyro.poutine.messenger*), 84
 MixedMultiOptimizer (class in *pyro.optim.multi*), 73

in MixtureOfDiagNormals (class in *pyro.distributions*), 42
 MixtureOfDiagNormalsSharedCovariance (class in *pyro.distributions*), 42
 mode (Parameterized attribute), 148
 model() (GPLVM method), 137
 model() (GPModel method), 128
 model() (GPRegression method), 131
 model() (SparseGPRegression method), 133
 model() (VariationalGP method), 135
 model() (VariationalSparseGP method), 136
 in module() (in module *pyro*), 6
 module_from_param_with_module_name() (in module *pyro.params.param_store*), 63
 MultiClass (class in *pyro.contrib.gp.likelihoods*), 146
 Multinomial (class in *pyro.distributions*), 29
 MultiOptimizer (class in *pyro.optim.multi*), 72
 MultiStepLR() (in module *pyro.optim.pytorch_optimizers*), 72
 MultivariateNormal (class in *pyro.distributions*), 29

N

name_count() (in module *pyro.contrib.autoname*), 112
 name_count() (in module *pyro.contrib.autoname.scoping*), 116
 NameCountMessenger (class in *pyro.contrib.autoname.scoping*), 115
 named_parameters() (ParamStoreDict method), 62
 Ncp (class in *pyro.contrib.tracking.dynamic_models*), 165
 NcpContinuous (class in *pyro.contrib.tracking.dynamic_models*), 167
 NcpDiscrete (class in *pyro.contrib.tracking.dynamic_models*), 168
 Ncv (class in *pyro.contrib.tracking.dynamic_models*), 166
 NcvContinuous (class in *pyro.contrib.tracking.dynamic_models*), 167
 NcvDiscrete (class in *pyro.contrib.tracking.dynamic_models*), 168
 nearby() (LSH method), 171
 NegativeBinomial (class in *pyro.distributions*), 30
 Newton (class in *pyro.optim.multi*), 73
 newton_step() (in module *pyro.ops.newton*), 93
 newton_step_1d() (in module *pyro.ops.newton*), 93
 newton_step_2d() (in module *pyro.ops.newton*), 94
 newton_step_3d() (in module *pyro.ops.newton*), 94
 next_context() (IndepMessenger method), 87
 nmc_eig() (in module *pyro.contrib.oed.eig*), 154
 NonlocalExit, 88
 nonreparam_stochastic_nodes (Trace attribute), 83

- Normal (class in `pyro.distributions`), 30
- `num_process_noise_parameters` (Dynamic-Model attribute), 164
- `num_steps` (HMC attribute), 23
- NUTS (class in `pyro.infer.mcmc`), 23
- ## O
- Object (class in `pyro.contrib.autoname.named`), 113
- `observation_nodes` (Trace attribute), 83
- OMTMultivariateNormal (class in `pyro.distributions`), 43
- OneHotCategorical (class in `pyro.distributions`), 30
- ## P
- `pack_tensors()` (Trace method), 83
- `param()` (in module `pyro`), 5
- `param()` (in module `pyro.contrib.minipyro`), 152
- `param_()` (Object method), 114
- `param_name()` (ParamStoreDict method), 63
- `param_nodes` (Trace attribute), 83
- `param_with_module_name()` (in module `pyro.params.param_store`), 63
- Parameterized (class in `pyro.contrib.gp.parameterized`), 147
- ParamStoreDict (class in `pyro.params.param_store`), 61
- Pareto (class in `pyro.distributions`), 30
- Periodic (class in `pyro.contrib.gp.kernels`), 141
- PermuteTransform (class in `pyro.distributions.transforms`), 55
- `pi()` (in module `pyro.ops.stats`), 99
- PlanarFlow (class in `pyro.distributions.transforms`), 56
- `plate` (class in `pyro`), 6
- `plate()` (EasyGuide method), 121
- `plate()` (in module `pyro.contrib.minipyro`), 152
- PlateMessenger (class in `pyro.contrib.minipyro`), 151
- Poisson (class in `pyro.contrib.gp.likelihoods`), 146
- Poisson (class in `pyro.distributions`), 30
- Polynomial (class in `pyro.contrib.gp.kernels`), 142
- PolynomialFlow (class in `pyro.distributions.transforms`), 57
- PositionMeasurement (class in `pyro.contrib.tracking.measurements`), 172
- `posterior_eig()` (in module `pyro.contrib.oed.eig`), 156
- `postprocess_message()` (Messenger method), 151
- `postprocess_message()` (trace method), 152
- `potential_grad()` (in module `pyro.ops.integrator`), 92
- `predecessors()` (Trace method), 84
- `predict()` (EKFSState method), 169
- `predictive()` (in module `pyro.infer.mcmc.util`), 26
- `process_message()` (block method), 152
- `process_message()` (Messenger method), 151
- `process_message()` (PlateMessenger method), 151
- `process_message()` (replay method), 152
- `process_noise_cov()` (DynamicModel method), 165
- `process_noise_cov()` (Ncp method), 166
- `process_noise_cov()` (NcpContinuous method), 167
- `process_noise_cov()` (NcpDiscrete method), 168
- `process_noise_cov()` (Ncv method), 167
- `process_noise_cov()` (NcvContinuous method), 168
- `process_noise_cov()` (NcvDiscrete method), 168
- `process_noise_dist()` (DynamicModel method), 165
- Product (class in `pyro.contrib.gp.kernels`), 142
- `prune_subsample_sites()` (in module `pyro.poutine.util`), 90
- `psis_diagnostic()` (in module `pyro.infer.importance`), 16
- `pyro.contrib.autoguide` (module), 103
- `pyro.contrib.autoguide.initialization` (module), 110
- `pyro.contrib.autoname` (module), 111
- `pyro.contrib.autoname.named` (module), 113
- `pyro.contrib.autoname.scoping` (module), 115
- `pyro.contrib.bnn` (module), 119
- `pyro.contrib.bnn.hidden_layer` (module), 119
- `pyro.contrib.easyguide` (module), 121
- `pyro.contrib.glm` (module), 125
- `pyro.contrib.gp` (module), 127
- `pyro.contrib.gp.kernels` (module), 137
- `pyro.contrib.gp.likelihoods` (module), 144
- `pyro.contrib.gp.models.gplvm` (module), 137
- `pyro.contrib.gp.models.gpr` (module), 130
- `pyro.contrib.gp.models.model` (module), 127
- `pyro.contrib.gp.models.sgpr` (module), 132
- `pyro.contrib.gp.models.vgp` (module), 134
- `pyro.contrib.gp.models.vsgp` (module), 135
- `pyro.contrib.gp.parameterized` (module), 147
- `pyro.contrib.gp.util` (module), 148
- `pyro.contrib.minipyro` (module), 149
- `pyro.contrib.oed` (module), 153
- `pyro.contrib.oed.eig` (module), 153
- `pyro.contrib.tracking` (module), 161
- `pyro.contrib.tracking.assignment` (module), 161
- `pyro.contrib.tracking.distributions` (module), 164

pyro.contrib.tracking.dynamic_models (module), 164
 pyro.contrib.tracking.extended_kalman_filter (module), 168
 pyro.contrib.tracking.hashing (module), 170
 pyro.contrib.tracking.measurements (module), 172
 pyro.distributions.torch (module), 27
 pyro.generic (module), 101
 pyro.infer.abstract_infer (module), 18
 pyro.infer.discrete (module), 17
 pyro.infer.elbo (module), 10
 pyro.infer.importance (module), 16
 pyro.infer.renyi_elbo (module), 15
 pyro.infer.svi (module), 9
 pyro.infer.trace_elbo (module), 10
 pyro.infer.trace_mean_field_elbo (module), 14
 pyro.infer.trace_tail_adaptive_elbo (module), 15
 pyro.infer.traceenum_elbo (module), 12
 pyro.infer.tracegraph_elbo (module), 11
 pyro.nn.auto_reg_nn (module), 65
 pyro.ops.dual_averaging (module), 91
 pyro.ops.einsum (module), 96
 pyro.ops.indexing (module), 95
 pyro.ops.integrator (module), 92
 pyro.ops.newton (module), 93
 pyro.ops.stats (module), 98
 pyro.ops.welford (module), 92
 pyro.optim.adagrad_rmsprop (module), 70
 pyro.optim.clipped_adam (module), 71
 pyro.optim.lr_scheduler (module), 70
 pyro.optim.multi (module), 72
 pyro.optim.optim (module), 69
 pyro.optim.pytorch_optimizers (module), 71
 pyro.params.param_store (module), 61
 pyro.poutine.block_messenger (module), 85
 pyro.poutine.broadcast_messenger (module), 86
 pyro.poutine.condition_messenger (module), 86
 pyro.poutine.escape_messenger (module), 86
 pyro.poutine.handlers (module), 75
 pyro.poutine.indep_messenger (module), 86
 pyro.poutine.lift_messenger (module), 87
 pyro.poutine.messenger (module), 84
 pyro.poutine.replay_messenger (module), 87
 pyro.poutine.runtime (module), 88
 pyro.poutine.scale_messenger (module), 87
 pyro.poutine.trace_messenger (module), 88
 pyro.poutine.util (module), 89
 pyro_backend() (in module pyro.generic), 101
 PyroLRScheduler (class in pyro.optim.lr_scheduler), 70
 PyroMultiOptimizer (class in pyro.optim.multi), 73
 PyroOptim (class in pyro.optim.optim), 69

Q

Q() (SylvesterFlow method), 59
 quantile() (in module pyro.ops.stats), 99
 quantiles() (AutoContinuous method), 106
 queue() (in module pyro.poutine), 79

R

R() (SylvesterFlow method), 59
 RadialFlow (class in pyro.distributions.transforms), 58
 random_module() (in module pyro), 6
 rate (GammaPoisson attribute), 39
 rate (InverseGamma attribute), 40
 RationalQuadratic (class in pyro.contrib.gp.kernels), 142
 RBF (class in pyro.contrib.gp.kernels), 142
 ReduceLROnPlateau() (in module pyro.optim.pytorch_optimizers), 72
 register() (pyro.poutine.messenger.Messenger class method), 84
 Rejection (class in pyro.distributions), 44
 RelaxedBernoulli (class in pyro.distributions), 30
 RelaxedBernoulliStraightThrough (class in pyro.distributions), 43
 RelaxedOneHotCategorical (class in pyro.distributions), 30
 RelaxedOneHotCategoricalStraightThrough (class in pyro.distributions), 44
 remove() (LSH method), 171
 remove_node() (Trace method), 84
 RenyiELBO (class in pyro.infer.renyi_elbo), 15
 reparameterized_nodes (Trace attribute), 84
 replace_param() (ParamStoreDict method), 62
 replay (class in pyro.contrib.minipyro), 152
 replay() (in module pyro.poutine), 80
 ReplayMessenger (class in pyro.poutine.replay_messenger), 87
 resample() (in module pyro.ops.stats), 99
 reset() (DualAveraging method), 92
 reset() (WelfordCovariance method), 92
 reset_parameters() (HouseholderFlow method), 53
 reset_parameters() (PlanarFlow method), 57
 reset_parameters() (PolynomialFlow method), 58
 reset_parameters() (RadialFlow method), 58
 reset_parameters2() (SylvesterFlow method), 59
 reset_stack() (NonlocalExit method), 88
 reshape() (TorchDistributionMixin method), 33

- RMSprop() (in module `pyro.optim.pytorch_optimizers`), 71
 Rprop() (in module `pyro.optim.pytorch_optimizers`), 71
 rsample() (AVFMultivariateNormal method), 35
 rsample() (Delta method), 36
 rsample() (GaussianScaleMixture method), 40
 rsample() (MaskedMixture method), 41
 rsample() (MixtureOfDiagNormals method), 42
 rsample() (MixtureOfDiagNormalsSharedCovariance method), 43
 rsample() (OMTMultivariateNormal method), 43
 rsample() (Rejector method), 44
 rsample() (RelaxedBernoulliStraightThrough method), 44
 rsample() (RelaxedOneHotCategoricalStraightThrough method), 44
 run() (MCMC method), 21
 run() (SVI method), 10
 run() (TracePosterior method), 20
- ## S
- S() (SylvesterFlow method), 59
 sample() (BetaBinomial method), 36
 sample() (DirichletMultinomial method), 37
 sample() (Distribution method), 32
 sample() (Empirical method), 38
 sample() (GammaPoisson method), 39
 sample() (Group method), 123
 sample() (HMC method), 23
 sample() (in module `pyro`), 5
 sample() (in module `pyro.contrib.minipyro`), 152
 sample() (LKJCorrCholesky method), 41
 sample() (MaskedMixture method), 41
 sample() (NUTS method), 25
 sample() (SpanningTree method), 45
 sample() (VonMises method), 46
 sample_() (Object method), 114
 sample_latent() (AutoContinuous method), 107
 sample_latent() (AutoGuide method), 104
 sample_mask_indices() (in module `pyro.nn.auto_reg_nn`), 66
 sample_posterior() (TraceEnum_ELBO method), 13
 sample_saved() (TraceEnumSample_ELBO method), 18
 sample_size (Empirical attribute), 38
 save() (ParamStoreDict method), 63
 save() (PyroOptim method), 69
 scale() (in module `pyro.poutine`), 80
 ScaleMessenger (class in `pyro.poutine.scale_messenger`), 87
 scope() (in module `pyro.contrib.autoname`), 111
 scope() (in module `pyro.contrib.autoname.scoping`), 115
 ScopeMessenger (class in `pyro.contrib.autoname.scoping`), 115
 score_parts() (Distribution method), 32
 score_parts() (Rejector method), 44
 set_constraint() (Parameterized method), 147
 set_data() (GPModel method), 129
 set_mode() (Parameterized method), 148
 set_prior() (Parameterized method), 148
 set_state() (ParamStoreDict method), 63
 set_state() (PyroOptim method), 69
 setdefault() (ParamStoreDict method), 62
 setup() (HMC method), 23
 SGD() (in module `pyro.optim.pytorch_optimizers`), 71
 shape() (TorchDistributionMixin method), 33
 share_memory() (AdagradRMSProp method), 71
 site_is_subsample() (in module `pyro.poutine.util`), 90
 SpanningTree (class in `pyro.distributions`), 45
 SparseAdam() (in module `pyro.optim.pytorch_optimizers`), 72
 SparseGPRegression (class in `pyro.contrib.gp.models.sgpr`), 132
 split_gelman_rubin() (in module `pyro.ops.stats`), 98
 step() (AdagradRMSProp method), 71
 step() (ClippedAdam method), 71
 step() (DualAveraging method), 92
 step() (MixedMultiOptimizer method), 73
 step() (MultiOptimizer method), 72
 step() (PyroLRScheduler method), 70
 step() (PyroMultiOptimizer method), 73
 step() (SVI method), 10, 151
 step_size (HMC attribute), 23
 StepLR() (in module `pyro.optim.pytorch_optimizers`), 72
 stochastic_nodes (Trace attribute), 84
 StudentT (class in `pyro.distributions`), 30
 successors() (Trace method), 84
 Sum (class in `pyro.contrib.gp.kernels`), 143
 summary() (MCMC method), 21
 support (BetaBinomial attribute), 36
 support (Delta attribute), 36
 support (DirichletMultinomial attribute), 37
 support (Empirical attribute), 38
 support (GammaPoisson attribute), 39
 support (InverseGamma attribute), 40
 support (LKJCorrCholesky attribute), 41
 support (MaskedMixture attribute), 41
 support (SpanningTree attribute), 46
 support (VonMises attribute), 46
 support (VonMises3D attribute), 47
 support() (Marginals method), 19
 SVI (class in `pyro.contrib.minipyro`), 151
 SVI (class in `pyro.infer.svi`), 9

SylvesterFlow (class
pyro.distributions.transforms), 59
symbolize_dims() (Trace method), 84

T

time (EKFState attribute), 169
time (Measurement attribute), 172
to_event() (TorchDistributionMixin method), 33
topological_sort() (Trace method), 84
TorchDistribution (class in pyro.distributions), 34
TorchDistributionMixin (class in
pyro.distributions.torch_distribution), 32
TorchMultiOptimizer (class in pyro.optim.multi),
73
trace (class in pyro.contrib.minipyro), 152
Trace (class in pyro.poutine), 82
trace (TraceHandler attribute), 88
trace() (in module pyro.ops.jit), 8
trace() (in module pyro.poutine), 81
Trace_ELBO (class in pyro.infer.trace_elbo), 10
Trace_ELBO() (in module pyro.contrib.minipyro), 151
TraceEnum_ELBO (class in
pyro.infer.traceenum_elbo), 12
TraceEnumSample_ELBO (class
pyro.infer.discrete), 18
TraceGraph_ELBO (class
pyro.infer.tracegraph_elbo), 11
TraceHandler (class
pyro.poutine.trace_messenger), 88
TraceMeanField_ELBO (class
pyro.infer.trace_mean_field_elbo), 14
TraceMessenger (class
pyro.poutine.trace_messenger), 88
TracePosterior (class in pyro.infer.abstract_infer),
19
TracePredictive (class
pyro.infer.abstract_infer), 20
TraceTailAdaptive_ELBO (class
pyro.infer.trace_tail_adaptive_elbo), 15
train() (in module pyro.contrib.gp.util), 149
TransformedDistribution (class in
pyro.distributions), 31
Transforming (class in pyro.contrib.gp.kernels), 143
TransformModule (class in pyro.distributions), 60
try_add() (ApproxSet method), 171

U

u() (HouseholderFlow method), 53
u_hat() (PlanarFlow method), 57
ubersum() (in module pyro.ops.contract), 98
Uniform (class in pyro.distributions), 31
unregister() (pyro.poutine.messenger.Messenger
class method), 85
update() (EKFState method), 170

in update() (WelfordCovariance method), 93
user_param_name() (in module
pyro.params.param_store), 63

V

validate_edges() (SpanningTree method), 46
validation_enabled() (in module pyro), 8
values() (ParamStoreDict method), 62
variance (BetaBinomial attribute), 36
variance (Delta attribute), 36
variance (DirichletMultinomial attribute), 37
variance (Empirical attribute), 38
variance (GammaPoisson attribute), 39
variance (MaskedMixture attribute), 41
variance (VonMises attribute), 46
VariationalGP (class in pyro.contrib.gp.models.vgp),
134
VariationalSparseGP (class in
pyro.contrib.gp.models.vsgp), 135
vectorized (CondIndepStackFrame attribute), 86
vectorized_importance_weights() (in mod-
ule pyro.infer.importance), 17
velocity_verlet() (in module
pyro.ops.integrator), 92
VerticalScaling (class in pyro.contrib.gp.kernels),
143
vi_eig() (in module pyro.contrib.oed.eig), 154
Vindex (class in pyro.ops.indexing), 96
vindex() (in module pyro.ops.indexing), 95
vnmc_eig() (in module pyro.contrib.oed.eig), 158
volume_preserving (HouseholderFlow attribute),
53
volume_preserving (PermuteTransform attribute),
56
VonMises (class in pyro.distributions), 46
in VonMises3D (class in pyro.distributions), 47

W

waic() (in module pyro.ops.stats), 99
Warping (class in pyro.contrib.gp.kernels), 143
Weibull (class in pyro.distributions), 31
WelfordCovariance (class in pyro.ops.welford), 92
WhiteNoise (class in pyro.contrib.gp.kernels), 144